

rueegg_wissiak_optimized

March 3, 2023

1 Optimized Model

This model does not show under- or overfitting and performs well on both, training and testing data. Afterwards, a brief description on how to tackle the challenges of an optimal model complexity.

To address underfitting, one approach is to increase the complexity of the model by adding more layers or increasing the number of filters in each layer. To address overfitting, we can try several approaches. One approach is to simplify the model by removing some layers or decreasing the number of filters in each layer. Another approach is to use less epochs for example.

Adding dropout or weight decay can help to address both of the above mentioned issues. We can also try adjusting the hyperparameters such as learning rate, batch size, or number of epochs.

```
[24]: import numpy as np
import matplotlib.pyplot as plt
```

```
[25]: from tensorflow.keras.preprocessing.image import ImageDataGenerator

img_size = 150
batch_size = 32

# TODO: add data augmentation for optimized model
train_datagen = ImageDataGenerator(
    rescale=1/255.,
    rotation_range=1,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.005,
    zoom_range=0.05,
    horizontal_flip=True,
    fill_mode='nearest'
)
test_datagen = ImageDataGenerator(rescale=1/255.)

train_generator = train_datagen.flow_from_directory(
    './dataset/seg_train/seg_train',
    target_size=(img_size, img_size),
    batch_size=batch_size,
    shuffle=True,
```

```

        class_mode='sparse'
    )
    test_generator = test_datagen.flow_from_directory(
        './dataset/seg_test/seg_test',
        target_size=(img_size, img_size),
        batch_size=batch_size,
        shuffle=True,
        class_mode='sparse'
    )
    labels = ['buildings', 'forest', 'glacier', 'mountain', 'sea']

    samples = train_generator.__next__()
    images = samples[0]
    target = samples[1]

    plt.figure(figsize = (10,10))
    for i in range(15):
        plt.subplot(5,5,i+1)
        plt.subplots_adjust(hspace=0.3,wspace=0.3)
        plt.imshow(images[i])
        plt.title(f"Class: {labels[int(target[i])]}")
        plt.axis('off')

```

Found 2666 images belonging to 5 classes.

Found 2499 images belonging to 5 classes.



1.1 Building the Model

Here we use the same model as the overfitting one but we add some extras to reduce the overfitting behavior.

Regularization: reduce the impact of the weights. The weights then end up having less impact on the loss function which determines the error between the actual label and predicted label. This reduces complexity of the model and therefore reduces overfitting. We are adding regularization only to those layers which have the largest number of parameters according to the model summary.

Dropout Layers: The benefit of using dropout is no node in the network will be assigned with high parameter values, as a result the parameter values will be dispersed and the output of the current layer will not depend on a single node. E.g. Dropout(0.2) drops the input layers at a probability of 0.2.

To **improve generalization** of the model, data augmentation is a useful tool. With data augmentation we can add artificial effects to the images such as shearing, stretching, flipping, rotating and translating. Through these effects, the images always appear differently each time they appear in the training step and therefore the CNN doesn't adapt to the exact images but rather learns about the relative features inside of an image.

```
[26]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
      ↪Dropout
      from tensorflow.keras.regularizers import l2

      model = Sequential()

      model.add(Conv2D(32, (3,3), input_shape= (img_size,img_size,3), activation =
      ↪'relu', padding = 'same')) #padding = same size output
      model.add(MaxPooling2D())

      model.add(Conv2D(64, (3,3), activation = 'relu', padding = 'same'))
      model.add(MaxPooling2D())

      model.add(Conv2D(128, (3,3), activation = 'relu', padding = 'same'))
      model.add(MaxPooling2D())

      model.add(Conv2D(256, (3,3), activation = 'relu', padding = 'same'))
      model.add(MaxPooling2D())

      model.add(Conv2D(512, (3,3), activation = 'relu', padding = 'same',
      ↪kernel_regularizer=l2(1=0.0001)))
      model.add(MaxPooling2D())
      model.add(Dropout(0.05))
```

```

model.add(Conv2D(1024, (3,3), activation = 'relu', padding = 'same',
    ↳kernel_regularizer=l2(l=0.001)))
model.add(MaxPooling2D())
model.add(Dropout(0.15))

model.add(Conv2D(512, (3,3), activation = 'relu', padding = 'same',
    ↳kernel_regularizer=l2(l=0.001)))
model.add(MaxPooling2D())
model.add(Dropout(0.15))

model.add(Flatten())

model.add(Dense(256, activation='relu'))
model.add(Dense(5, activation = 'softmax'))

model.compile(optimizer = 'adam', loss='sparse_categorical_crossentropy',
    ↳metrics=['accuracy'])

```

[27]: `model.summary()`

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_14 (Conv2D)	(None, 150, 150, 32)	896
max_pooling2d_14 (MaxPooling2D)	(None, 75, 75, 32)	0
conv2d_15 (Conv2D)	(None, 75, 75, 64)	18496
max_pooling2d_15 (MaxPooling2D)	(None, 37, 37, 64)	0
conv2d_16 (Conv2D)	(None, 37, 37, 128)	73856
max_pooling2d_16 (MaxPooling2D)	(None, 18, 18, 128)	0
conv2d_17 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_17 (MaxPooling2D)	(None, 9, 9, 256)	0
conv2d_18 (Conv2D)	(None, 9, 9, 512)	1180160
max_pooling2d_18 (MaxPooling2D)	(None, 4, 4, 512)	0

```

g2D)

dropout_6 (Dropout)          (None, 4, 4, 512)          0

conv2d_19 (Conv2D)           (None, 4, 4, 1024)         4719616

max_pooling2d_19 (MaxPoolin (None, 2, 2, 1024)         0
g2D)

dropout_7 (Dropout)          (None, 2, 2, 1024)         0

conv2d_20 (Conv2D)           (None, 2, 2, 512)          4719104

max_pooling2d_20 (MaxPoolin (None, 1, 1, 512)         0
g2D)

dropout_8 (Dropout)          (None, 1, 1, 512)         0

flatten_2 (Flatten)          (None, 512)                0

dense_4 (Dense)              (None, 256)                131328

dense_5 (Dense)              (None, 5)                  1285

=====
Total params: 11,139,909
Trainable params: 11,139,909
Non-trainable params: 0
-----

```

1.2 Training the Model

```
[28]: history = model.fit(train_generator, validation_data=test_generator, epochs=20)
```



```

Epoch 1/20
84/84 [=====] - 103s 1s/step - loss: 1.9049 - accuracy:
0.2427 - val_loss: 1.6389 - val_accuracy: 0.2185
Epoch 2/20
84/84 [=====] - 104s 1s/step - loss: 1.2831 - accuracy:
0.4771 - val_loss: 1.1710 - val_accuracy: 0.5166
Epoch 3/20
84/84 [=====] - 102s 1s/step - loss: 1.1166 - accuracy:
0.5416 - val_loss: 1.0446 - val_accuracy: 0.5674
Epoch 4/20
84/84 [=====] - 103s 1s/step - loss: 1.0358 - accuracy:
0.5750 - val_loss: 1.1101 - val_accuracy: 0.5554
Epoch 5/20
84/84 [=====] - 103s 1s/step - loss: 0.9611 - accuracy:

```

0.6249 - val_loss: 0.9900 - val_accuracy: 0.5714
 Epoch 6/20
 84/84 [=====] - 102s 1s/step - loss: 0.8924 - accuracy:
 0.6489 - val_loss: 0.8405 - val_accuracy: 0.6755
 Epoch 7/20
 84/84 [=====] - 102s 1s/step - loss: 0.8536 - accuracy:
 0.6733 - val_loss: 0.8021 - val_accuracy: 0.6975
 Epoch 8/20
 84/84 [=====] - 103s 1s/step - loss: 0.7983 - accuracy:
 0.7089 - val_loss: 1.0597 - val_accuracy: 0.5918
 Epoch 9/20
 84/84 [=====] - 103s 1s/step - loss: 0.7701 - accuracy:
 0.7161 - val_loss: 0.7549 - val_accuracy: 0.7103
 Epoch 10/20
 84/84 [=====] - 102s 1s/step - loss: 0.7852 - accuracy:
 0.7086 - val_loss: 0.7423 - val_accuracy: 0.7235
 Epoch 11/20
 84/84 [=====] - 103s 1s/step - loss: 0.7113 - accuracy:
 0.7581 - val_loss: 0.7122 - val_accuracy: 0.7227
 Epoch 12/20
 84/84 [=====] - 102s 1s/step - loss: 0.7008 - accuracy:
 0.7584 - val_loss: 0.6669 - val_accuracy: 0.7627
 Epoch 13/20
 84/84 [=====] - 102s 1s/step - loss: 0.6609 - accuracy:
 0.7738 - val_loss: 0.6151 - val_accuracy: 0.7975
 Epoch 14/20
 84/84 [=====] - 103s 1s/step - loss: 0.6307 - accuracy:
 0.7824 - val_loss: 0.6130 - val_accuracy: 0.7979
 Epoch 15/20
 84/84 [=====] - 103s 1s/step - loss: 0.6139 - accuracy:
 0.7941 - val_loss: 0.6364 - val_accuracy: 0.7791
 Epoch 16/20
 84/84 [=====] - 103s 1s/step - loss: 0.5878 - accuracy:
 0.8053 - val_loss: 0.6348 - val_accuracy: 0.7895
 Epoch 17/20
 84/84 [=====] - 103s 1s/step - loss: 0.6011 - accuracy:
 0.8008 - val_loss: 0.6113 - val_accuracy: 0.8043
 Epoch 18/20
 84/84 [=====] - 103s 1s/step - loss: 0.5473 - accuracy:
 0.8248 - val_loss: 0.5588 - val_accuracy: 0.8107
 Epoch 19/20
 84/84 [=====] - 104s 1s/step - loss: 0.5554 - accuracy:
 0.8278 - val_loss: 0.5630 - val_accuracy: 0.8127
 Epoch 20/20
 84/84 [=====] - 104s 1s/step - loss: 0.5331 - accuracy:
 0.8282 - val_loss: 0.5094 - val_accuracy: 0.8399

```
[29]: %run rueegg_wissiak_model_visualization.ipynb
```



```
[30]: %run rueegg_wissiak_model_evaluation.ipynb
```

```
79/79 [=====] - 13s 169ms/step  
Predicted classes: [3 3 0 ... 1 3 4]  
True labels: [0 0 0 ... 4 4 4]  
Accuracy:  
0.1880752300920368
```