

March 17, 2023

Made with  by Thomas Rüegg and Patrick Wissiak

1 Introduction

In this Jupyter Notebook, we'll introduce a convolutional neural network that is underfitting, then, create a new model that will overfit to the training data. Since that will not be the desired behavior, we'll showcase various methods to defeat overfitting, resulting in a model that'll be optimal. We'll discuss learnings and wrap everything up with a summary and an outlook.

2 General Information

2.1 Environment Setup

Here we're providing our environment yaml file to use the same environment as we did to run the project.

```
name: aiap
channels:
  - apple
  - conda-forge
dependencies:
  - anyio=3.6.2=pyhd8ed1ab_0
  - appnope=0.1.3=pyhd8ed1ab_0
  - argon2-cffi=21.3.0=pyhd8ed1ab_0
  - argon2-cffi-bindings=21.2.0=py310h8e9501a_3
  - asttokens=2.2.1=pyhd8ed1ab_0
  - attrs=22.2.0=pyh71513ae_0
  - backcall=0.2.0=pyh9f0ad1d_0
  - backports=1.0=pyhd8ed1ab_3
  - backports.functools_lru_cache=1.6.4=pyhd8ed1ab_0
  - beautifulsoup4=4.11.2=pyha770c72_0
  - bleach=6.0.0=pyhd8ed1ab_0
  - brotli=1.0.9=h1a8c8d9_8
  - brotli-bin=1.0.9=h1a8c8d9_8
  - brotlipy=0.7.0=py310h8e9501a_1005
  - bzip2=1.0.8=h3422bc3_4
  - c-ares=1.18.1=h3422bc3_0
  - ca-certificates=2022.12.7=h4653dfc_0
```

- cached-property=1.5.2=hd8ed1ab_1
- cached_property=1.5.2=pyha770c72_1
- cffi=1.15.1=py310h2399d43_3
- comm=0.1.2=pyhd8ed1ab_0
- cryptography=39.0.2=py310hfc83b78_0
- cycler=0.11.0=pyhd8ed1ab_0
- debugpy=1.6.6=py310h0f1eb42_0
- decorator=5.1.1=pyhd8ed1ab_0
- defusedxml=0.7.1=pyhd8ed1ab_0
- entrypoints=0.4=pyhd8ed1ab_0
- executing=1.2.0=pyhd8ed1ab_0
- flit-core=3.8.0=pyhd8ed1ab_0
- fonttools=4.39.0=py310h8e9501a_0
- freetype=2.12.1=hd633e50_1
- grpcio=1.46.3=py310ha26ec5d_0
- h5py=3.6.0=nompi_py310hb8bbf05_100
- hdf5=1.12.1=nompi_hd9dbc9e_104
- idna=3.4=pyhd8ed1ab_0
- importlib-metadata=6.0.0=pyha770c72_0
- importlib_resources=5.12.0=pyhd8ed1ab_0
- ipykernel=6.21.2=pyh736e0ef_0
- ipython=8.11.0=pyhd1c38e8_0
- ipython_genutils=0.2.0=py_1
- jedi=0.18.2=pyhd8ed1ab_0
- jinja2=3.1.2=pyhd8ed1ab_1
- joblib=1.2.0=pyhd8ed1ab_0
- jsonschema=4.17.3=pyhd8ed1ab_0
- jupyter_client=7.4.9=pyhd8ed1ab_0
- jupyter_core=5.2.0=py310hbe9552e_0
- jupyter_events=0.6.3=pyhd8ed1ab_0
- jupyter_server=2.4.0=pyhd8ed1ab_0
- jupyter_server_terminals=0.4.4=pyhd8ed1ab_1
- jupyterlab_pygments=0.2.2=pyhd8ed1ab_0
- kiwisolver=1.4.4=py310h2887b22_1
- krb5=1.20.1=h69eda48_0
- lcms2=2.15=hd835a16_1
- lerc=4.0.0=h9a09cb3_0
- libblas=3.9.0=16_osxarm64_openblas
- libbrotlicommon=1.0.9=h1a8c8d9_8
- libbrotlidec=1.0.9=h1a8c8d9_8
- libbrotlienc=1.0.9=h1a8c8d9_8
- libcblas=3.9.0=16_osxarm64_openblas
- libcurl=7.88.1=h9049daf_0
- libcxx=15.0.7=h75e25f2_0
- libdeflate=1.17=h1a8c8d9_0
- libedit=3.1.20191231=hc8eb9b7_2
- libev=4.33=h642e427_1
- libffi=3.4.2=h3422bc3_5

- libgfortran=5.0.0=12_2_0_hd922786_31
- libgfortran5=12.2.0=h0eea778_31
- libjpeg-turbo=2.1.5.1=h1a8c8d9_0
- liblapack=3.9.0=16_osxarm64_openblas
- libnnghttp2=1.52.0=hae82a92_0
- libopenblas=0.3.21=openmp_hc731615_3
- libpng=1.6.39=h76d750c_0
- libprotobuf=3.19.4=hccf11d3_0
- libsodium=1.0.18=h27ca646_1
- libsqlite=3.40.0=h76d750c_0
- libssh2=1.10.0=h7a5bd25_3
- libtiff=4.5.0=hdc14d85_5
- libwebp-base=1.3.0=h1a8c8d9_0
- libxcb=1.13=h9b22ae9_1004
- libzlib=1.2.13=h03a7124_4
- llvm-openmp=15.0.7=h7cfbb63_0
- markupsafe=2.1.2=py310h8e9501a_0
- matplotlib=3.5.0=py310hb6292c7_0
- matplotlib-base=3.5.0=py310hacb9267_0
- matplotlib-inline=0.1.6=pyhd8ed1ab_0
- mistune=2.0.5=pyhd8ed1ab_0
- munkres=1.1.4=pyh9f0ad1d_0
- nbclassic=0.5.3=pyhb4ecaf3_3
- nbclient=0.7.2=pyhd8ed1ab_0
- nbconvert=7.2.9=pyhd8ed1ab_0
- nbconvert-core=7.2.9=pyhd8ed1ab_0
- nbconvert-pandoc=7.2.9=pyhd8ed1ab_0
- nbformat=5.7.3=pyhd8ed1ab_0
- ncurses=6.3=h07bb92c_1
- nest-asyncio=1.5.6=pyhd8ed1ab_0
- notebook=6.5.3=pyha770c72_0
- notebook-shim=0.2.2=pyhd8ed1ab_0
- numpy=1.23.2=py310h127c7cf_0
- openjpeg=2.5.0=hbc2ba62_2
- openssl=3.0.8=h03a7124_0
- packaging=23.0=pyhd8ed1ab_0
- pandoc=2.19.2=hce30654_2
- pandocfilters=1.5.0=pyhd8ed1ab_0
- parso=0.8.3=pyhd8ed1ab_0
- pexpect=4.8.0=pyh1a96a4e_2
- pickleshare=0.7.5=py_1003
- pillow=9.4.0=py310h07496d3_2
- pip=23.0.1=pyhd8ed1ab_0
- pkgutil-resolve-name=1.3.10=pyhd8ed1ab_0
- platformdirs=3.1.1=pyhd8ed1ab_0
- pooch=1.7.0=pyhd8ed1ab_0
- prometheus_client=0.16.0=pyhd8ed1ab_0
- prompt-toolkit=3.0.38=pyha770c72_0

- prompt_toolkit=3.0.38=hd8ed1ab_0
- protobuf=3.19.4=py310h1b49c16_0
- psutil=5.9.4=py310h8e9501a_0
- pthread-stubs=0.4=h27ca646_1001
- ptyprocess=0.7.0=pyhd3deb0d_0
- pure_eval=0.2.2=pyhd8ed1ab_0
- pycparser=2.21=pyhd8ed1ab_0
- pygments=2.14.0=pyhd8ed1ab_0
- pyopenssl=23.0.0=pyhd8ed1ab_0
- pyparsing=3.0.9=pyhd8ed1ab_0
- pyrsistent=0.19.3=py310h8e9501a_0
- pysocks=1.7.1=pyha2e5f31_6
- python=3.10.9=h3ba56d0_0_cpython
- python-dateutil=2.8.2=pyhd8ed1ab_0
- python-fastjsonschema=2.16.3=pyhd8ed1ab_0
- python-json-logger=2.0.7=pyhd8ed1ab_0
- python_abi=3.10=3_cp310
- pyyaml=6.0=py310h8e9501a_5
- pyzmq=24.0.1=py310hc407298_1
- readline=8.1.2=h46ed386_0
- requests=2.28.2=pyhd8ed1ab_0
- rfc3339-validator=0.1.4=pyhd8ed1ab_0
- rfc3986-validator=0.1.1=pyh9f0ad1d_0
- scikit-learn=1.2.1=py310ha00a7cd_0
- scipy=1.10.1=py310ha0d8a01_0
- send2trash=1.8.0=pyhd8ed1ab_0
- setuptools=67.6.0=pyhd8ed1ab_0
- six=1.16.0=pyh6c4a22f_0
- sniffio=1.3.0=pyhd8ed1ab_0
- soupsieve=2.3.2.post1=pyhd8ed1ab_0
- stack_data=0.6.2=pyhd8ed1ab_0
- tensorflow-deps=2.10.0=0
- terminado=0.17.1=pyhd1c38e8_0
- threadpoolctl=3.1.0=pyh8a188c0_0
- tinycc2=1.2.1=pyhd8ed1ab_0
- tk=8.6.12=he1e0b03_0
- tornado=6.2=py310h8e9501a_1
- traitlets=5.9.0=pyhd8ed1ab_0
- typing-extensions=4.5.0=hd8ed1ab_0
- typing_extensions=4.5.0=pyha770c72_0
- tzdata=2022g=h191b570_0
- unicodedata2=15.0.0=py310h8e9501a_0
- urllib3=1.26.15=pyhd8ed1ab_0
- wcwidth=0.2.6=pyhd8ed1ab_0
- webencodings=0.5.1=py_1
- websocket-client=1.5.1=pyhd8ed1ab_0
- wheel=0.38.4=pyhd8ed1ab_0
- xorg-libxau=1.0.9=h27ca646_0

- xorg-libxdmcp=1.1.3=h27ca646_0
- xz=5.2.6=h57fd34a_0
- yaml=0.2.5=h3422bc3_2
- zeromq=4.3.4=hbdaafb3b_1
- zipp=3.15.0=pyhd8ed1ab_0
- zlib=1.2.13=h03a7124_4
- zstd=1.5.2=hf913c23_6
- pip:
 - absl-py==1.4.0
 - arrow==1.2.3
 - astunparse==1.6.3
 - cachetools==5.3.0
 - certifi==2022.12.7
 - charset-normalizer==3.1.0
 - extra-keras-datasets==1.2.0
 - flatbuffers==1.12
 - fqdn==1.5.1
 - gast==0.4.0
 - google-auth==2.16.2
 - google-auth-oauthlib==0.4.6
 - google-pasta==0.2.0
 - ipython-genutils==0.2.0
 - isoduration==20.11.0
 - jsonpointer==2.3
 - jupyter-server==2.3.0
 - keras==2.9.0
 - keras-preprocessing==1.1.2
 - libclang==15.0.6.1
 - markdown==3.4.1
 - oauthlib==3.2.2
 - opt-einsum==3.3.0
 - pandas==1.5.3
 - pyasn1==0.4.8
 - pyasn1-modules==0.2.8
 - pytz==2022.7.1
 - qtconsole==5.4.0
 - qtpy==2.3.0
 - requests-oauthlib==1.3.1
 - rsa==4.9
 - scikeras==0.10.0
 - tensorboard==2.9.1
 - tensorboard-data-server==0.6.1
 - tensorboard-plugin-wit==1.8.1
 - tensorcross==0.4.4
 - tensorflow-estimator==2.9.0
 - tensorflow-macos==2.9.0
 - tensorflow-metal==0.5.1
 - termcolor==2.2.0

- `uri-template==1.2.0`
- `webcolors==1.12`
- `werkzeug==2.2.3`
- `wrapt==1.15.0`

2.2 The Dataset

We used [this dataset from kaggle](#) for our miniproject.

The dataset contains 6 classes of images: `building`, `forest`, `glacier`, `mountain`, `sea`, and `street`. We decided to delete the class `street` and its corresponding pictures to simplify the learning process. The images are divided into folders with their respective labels. We used the `image_dataset_from_directory` method of Keras to convert the images into a TensorFlow dataset object for training.

The dataset fulfills all of the requirements posed in the assignment:

- from 4 to 15 classes (had 6, now 5)
- minimum of 500 samples per class (has ~2300)
- less than 150k samples in total (has 24.3k)

2.3 Importing the Data

```
[1]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.python.client import device_lib
from sklearn import metrics

print(device_lib.list_local_devices())

[name: "/device:CPU:0"
device_type: "CPU"
memory_limit: 268435456
locality {
}
incarnation: 17957140826399196183
xla_global_id: -1
, name: "/device:GPU:0"
device_type: "GPU"
locality {
    bus_id: 1
}
incarnation: 5849238088427255724
physical_device_desc: "device: 0, name: METAL, pci bus id: <undefined>"
xla_global_id: -1
]Metal device set to: Apple M2 Pro

systemMemory: 16.00 GB
maxCacheSize: 5.33 GB
```

```
2023-03-17 12:05:10.195545: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2023-03-17 12:05:10.195653: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/device:GPU:0 with 0 MB memory) -> physical
PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)
```

```
[2]: from tensorflow.keras.utils import image_dataset_from_directory
training_img_path = "./dataset/seg_train/seg_train/"

img_size = 150
seed = 42    # the seed will make sure the two datasets are not overlapping

def import_models(batch_size=16):
    train_ds = image_dataset_from_directory(
        './dataset/seg_train/seg_train',
        validation_split=0.2,
        subset="training",
        labels="inferred",
        seed=seed,
        image_size=(img_size, img_size),
        batch_size=batch_size
    )
    val_ds = image_dataset_from_directory(
        './dataset/seg_train/seg_train',
        validation_split=0.2,
        subset="validation",
        labels="inferred",
        seed=seed,
        image_size=(img_size, img_size),
        batch_size=batch_size
    )
    test_ds = image_dataset_from_directory(
        './dataset/seg_test/seg_test',
        labels="inferred",
        seed=seed,
        image_size=(img_size, img_size),
        batch_size=batch_size
    )
    return train_ds, val_ds, test_ds
train_ds, val_ds, test_ds = import_models()

ds = test_ds.unbatch()
ds = list(ds.as_numpy_iterator())
test_images = [x for x, y in ds]
```

```

test_labels = [y for x, y in ds]

full_train_dataset = image_dataset_from_directory(
    './dataset/seg_train/seg_train',
    seed=seed,
    image_size=(img_size, img_size)
)

ds = full_train_dataset.unbatch()
ds = list(ds.as_numpy_iterator())
train_images = np.array([x for x, y in ds])
train_labels = np.array([y for x, y in ds])

```

Found 2666 files belonging to 5 classes.

Using 2133 files for training.

2023-03-17 12:05:10.339362: I

tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.

2023-03-17 12:05:10.339387: I

tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
<undefined>)

Found 2666 files belonging to 5 classes.

Using 533 files for validation.

Found 2499 files belonging to 5 classes.

2023-03-17 12:05:10.604495: W

tensorflow/core/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU
frequency: 0 Hz

Found 2666 files belonging to 5 classes.

2.4 Common functionalities

Here we define functions that we will need in later sections of the notebook.

This code handles the data augmentation which can be added to the beginning of a model. As data augmentation we use flipping as well as zooming and rotating both with a percentage of 0.1.

```
[3]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import RandomRotation, RandomZoom, RandomFlip

data_augmentation = Sequential(
    [
        RandomFlip("horizontal", input_shape=(img_size, img_size, 3)),

```

```

        RandomRotation(0.1),
        RandomZoom(0.1),
    ]
)

```

The below function shows sample images from the dataset and can apply data augmentation if activated per parameter.

```
[4]: def show_images(title, augment = False):
    plt.figure(figsize=(5, 5))
    plt.suptitle(title)

    for i in range(9):
        if augment:
            # Convert to shape with batch for prediction
            img = np.reshape(test_images[i], (1, img_size, img_size, 3))
            img = data_augmentation(img).numpy()
            # Reshape back to normal image
            img = np.reshape(img, (img_size, img_size, 3))
        else:
            img = test_images[i]
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(img.astype("uint8"))
        plt.title(train_ds.class_names[test_labels[i]])
        plt.axis("off")
```

2.4.1 Visualize the Results

The beneath function if used to visualize the results of the accuracy and loss curves after a model has been fit. The x-axis shows the epochs and the y-axis shows the loss or accuracy, respectively. In one diagram the curve for training and validation is shown.

```
[5]: def visualize_results(history):
    # Plot model loss
    plt.figure(figsize=(10,4))
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('Model Loss')
    plt.ylabel('loss')
    plt.xlabel('epochs')
    plt.legend(['Train', 'Val'], loc= 'upper left')
    plt.show()

    # Plot model accuracy
    plt.figure(figsize=(10,3))
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
```

```

plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['Train', 'Val'], loc='upper left')
plt.show()

```

Next comes an easy-to-use method that shows 12 sample images along with the predictions of the model on the same images and the corresponding labels.

```
[6]: def show_images_with_predictions(model):
    plt.figure(figsize=(15, 15))
    for i in range(12):
        img = np.reshape(test_images[i], (1, img_size, img_size, 3))
        pred = model.predict(img).flatten()
        pred_idx = np.argmax(pred)
        pred_perc = "{:.2f}".format(pred[pred_idx])
        true_label = test_ds.class_names[test_labels[i]]
        pred_label = test_ds.class_names[pred_idx]
        true_label_perc = "{:.2f}".format(pred[test_labels[i]])
        ax = plt.subplot(3, 4, i + 1)
        plt.imshow(test_images[i].astype("uint8"))
        plt.title(f"True: {true_label}, {true_label_perc}\nPred: {pred_label}, {pred_perc}")
        plt.axis("off")
```

Underneath method shows a confusion matrix based on a given model's predictions on the test dataset.

```
[7]: def show_confusion_matrix(model):
    true_labels = np.array([], dtype=np.uint8)
    pred_labels = np.array([], dtype=np.uint8)
    for images_batch, true_labels_batch in test_ds:
        pred_labels_batch = np.argmax(model.predict(images_batch), axis=-1)
        true_labels = np.concatenate([true_labels, true_labels_batch])
        pred_labels = np.concatenate([pred_labels, pred_labels_batch])

    confusion_matrix = metrics.confusion_matrix(y_true=true_labels,
                                                y_pred=pred_labels)
    display = metrics.ConfusionMatrixDisplay(confusion_matrix=confusion_matrix,
                                              display_labels=test_ds.class_names)
    display.plot()
    plt.show()
    return true_labels, pred_labels
```

3 Underfitting Model

We will create a model that is too simple and won't describe the data accurately enough.

3.1 Building the Model

We start with a convolutional layer that'll import the images and use the ReLU activation function (we use that activation function for all the convolution and dense layers). Using a saturating activation function (like e.g. hyperbolic tangent tanh) would not be as efficient, but it would not contribute to underfitting. It would only make the model slower. That's why we use the more efficient non-saturating activation function ReLU.

The next layer will run a kernel of 3×3 over each image, 6 times. `padding=same` is referring to the padding of the image (needed because of the kernel size) being filled with zeros. We have set the stride of the first convolutional layer to $(5, 5)$. This was needed (for the underfitting model at least) to ensure there are no more than 5000 parameters. It's similar to downscaling the images from $150 \times 150px$ to $30 \times 30px$. We did not rescale the images in the beginning, because this way, we are able to preserve the quality of the images within the plots. Furthermore, we did not notice any increase in runtime, each epoch took the same amount of time to calculate (thanks to the GPU acceleration with tensorflow-metal, more on that later).

After that there's a `MaxPooling2D` layer which goes through the feature map (output of the `Conv2D` layer) and only takes the most relevant information.

Following that is another convolution layer and max pooling layer. Next, the image will be flattened into a vector, ready to be fed to the following dense layer.

The last layer is a dense layer which represents the output layer, having a shared softmax activation layer to determine the probabilities of the 5 different classes.

```
[27]: from tensorflow.keras.models import Sequential
       from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

def create_model_underfitting():
    model = Sequential([
        Conv2D(6, (3,3), input_shape=(img_size,img_size,3), activation='relu', ↴
               padding='same', strides=(5,5)),
        MaxPooling2D(),
        Conv2D(6, (3,3), activation='relu', padding='same'),
        MaxPooling2D(),
        Flatten(),
        Dense(5, activation='softmax')
    ])
    model._name='underfitting_model'

    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', ↴
                  metrics=['accuracy'])

    return model
```

```
[28]: model_underfitting = create_model_underfitting()
model_underfitting.summary()
```

```
Model: "underfitting_model"
-----
Layer (type)          Output Shape       Param #
=====
conv2d_7 (Conv2D)     (None, 50, 50, 6)    168
max_pooling2d_7 (MaxPooling 2D) (None, 25, 25, 6) 0
conv2d_8 (Conv2D)     (None, 25, 25, 6)    330
max_pooling2d_8 (MaxPooling 2D) (None, 12, 12, 6) 0
flatten_2 (Flatten)   (None, 864)         0
dense_3 (Dense)      (None, 5)           4325
=====
Total params: 4,823
Trainable params: 4,823
Non-trainable params: 0
```

3.1.1 Dropout?

At first we had some more dense layers in our underfitting model, and we were thinking if we should apply dropout to some of them, because dropout will reduce overfitting, and we thought that should therefore make our underfitting model even more underfitting. This is now no longer relevant (because the current model only has one dense layer and two convolution layers), but we wanted to share our findings.

We ended up halving the amount of nodes of those layers for the underfitting model. Why? We learned that the difference between applying `dropout(0.5)` and halving the nodes of the layer is that dropout randomly drops out nodes (more precisely, it reduces the weights to 0) during each training iteration, which means that **different** nodes will be dropped out in each iteration. This allows the network to learn more robust and generalizable representations of the data, as it is forced to rely on a subset of nodes in each iteration, which prevents overfitting. Each node needs to supply relevant information to the next layer on its own, as the next layer's node can no longer depend on the combination of information from the previous nodes (because some nodes are not there anymore). Also, dropout is present during training, but not during inference.

While, on the other hand, halving the nodes reduces the number of nodes in the layer permanently, which means that the network has less capacity to learn and represent complex patterns in the data. This then leads to underfitting, where the model is not able to capture the important features in the data. Which is exactly what we want.

3.2 Training the Model

```
[29]: history = model_underfitting.fit(  
    train_ds,  
    validation_data=val_ds,  
    epochs=30  
)  
  
Epoch 1/30  
6/67 [=>...] - ETA: 0s - loss: 22.4902 - accuracy:  
0.1771  
2023-03-16 15:36:09.319446: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.  
67/67 [=====] - 1s 12ms/step - loss: 6.4289 - accuracy:  
0.2386 - val_loss: 2.0746 - val_accuracy: 0.2570  
Epoch 2/30  
1/67 [...] - ETA: 1s - loss: 2.0553 - accuracy:  
0.2500  
2023-03-16 15:36:10.040110: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.  
67/67 [=====] - 1s 10ms/step - loss: 1.8675 - accuracy:  
0.2818 - val_loss: 1.7222 - val_accuracy: 0.2946  
Epoch 3/30  
67/67 [=====] - 1s 9ms/step - loss: 1.6406 - accuracy:  
0.3221 - val_loss: 1.6075 - val_accuracy: 0.3021  
Epoch 4/30  
67/67 [=====] - 1s 9ms/step - loss: 1.5222 - accuracy:  
0.3638 - val_loss: 1.5152 - val_accuracy: 0.3771  
Epoch 5/30  
67/67 [=====] - 1s 10ms/step - loss: 1.4324 - accuracy:  
0.4069 - val_loss: 1.4698 - val_accuracy: 0.3790  
Epoch 6/30  
67/67 [=====] - 1s 9ms/step - loss: 1.3759 - accuracy:  
0.4187 - val_loss: 1.4463 - val_accuracy: 0.3940  
Epoch 7/30  
67/67 [=====] - 1s 10ms/step - loss: 1.3354 - accuracy:  
0.4299 - val_loss: 1.4311 - val_accuracy: 0.3959  
Epoch 8/30  
67/67 [=====] - 1s 10ms/step - loss: 1.3036 - accuracy:  
0.4379 - val_loss: 1.4406 - val_accuracy: 0.3977  
Epoch 9/30  
67/67 [=====] - 1s 10ms/step - loss: 1.2746 - accuracy:  
0.4487 - val_loss: 1.4403 - val_accuracy: 0.3977  
Epoch 10/30
```

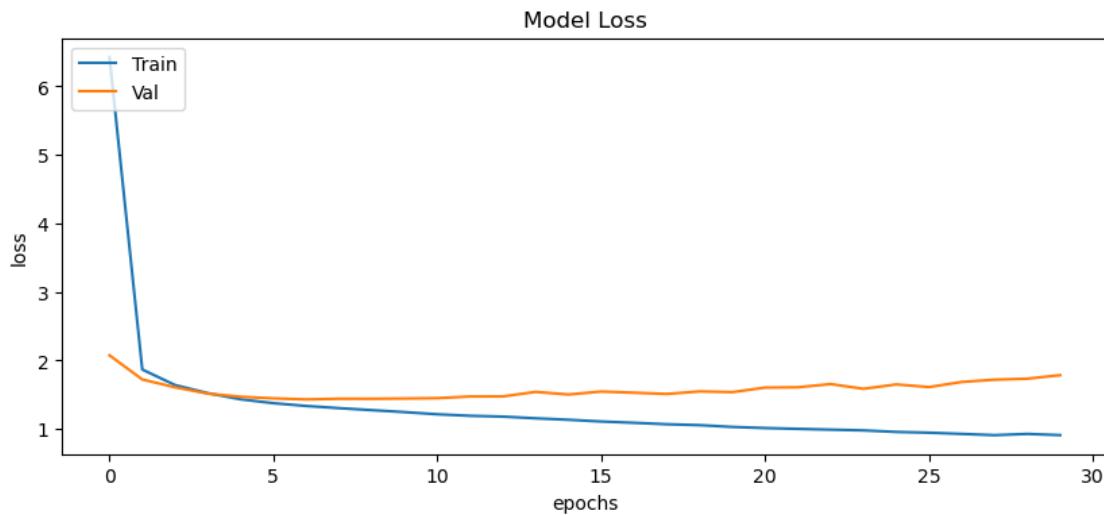
```
67/67 [=====] - 1s 10ms/step - loss: 1.2459 - accuracy: 0.4609 - val_loss: 1.4444 - val_accuracy: 0.4071
Epoch 11/30
67/67 [=====] - 1s 10ms/step - loss: 1.2140 - accuracy: 0.4674 - val_loss: 1.4495 - val_accuracy: 0.4015
Epoch 12/30
67/67 [=====] - 1s 10ms/step - loss: 1.1920 - accuracy: 0.4698 - val_loss: 1.4737 - val_accuracy: 0.4090
Epoch 13/30
67/67 [=====] - 1s 10ms/step - loss: 1.1800 - accuracy: 0.4759 - val_loss: 1.4744 - val_accuracy: 0.4109
Epoch 14/30
67/67 [=====] - 1s 10ms/step - loss: 1.1553 - accuracy: 0.4852 - val_loss: 1.5407 - val_accuracy: 0.3996
Epoch 15/30
67/67 [=====] - 1s 10ms/step - loss: 1.1340 - accuracy: 0.5007 - val_loss: 1.5027 - val_accuracy: 0.4034
Epoch 16/30
67/67 [=====] - 1s 10ms/step - loss: 1.1087 - accuracy: 0.4998 - val_loss: 1.5458 - val_accuracy: 0.4109
Epoch 17/30
67/67 [=====] - 1s 10ms/step - loss: 1.0897 - accuracy: 0.5199 - val_loss: 1.5289 - val_accuracy: 0.4090
Epoch 18/30
67/67 [=====] - 1s 9ms/step - loss: 1.0686 - accuracy: 0.5298 - val_loss: 1.5106 - val_accuracy: 0.4146
Epoch 19/30
67/67 [=====] - 1s 10ms/step - loss: 1.0555 - accuracy: 0.5321 - val_loss: 1.5478 - val_accuracy: 0.4221
Epoch 20/30
67/67 [=====] - 1s 10ms/step - loss: 1.0294 - accuracy: 0.5466 - val_loss: 1.5372 - val_accuracy: 0.4221
Epoch 21/30
67/67 [=====] - 1s 9ms/step - loss: 1.0131 - accuracy: 0.5485 - val_loss: 1.6034 - val_accuracy: 0.4221
Epoch 22/30
67/67 [=====] - 1s 10ms/step - loss: 1.0013 - accuracy: 0.5551 - val_loss: 1.6085 - val_accuracy: 0.4278
Epoch 23/30
67/67 [=====] - 1s 10ms/step - loss: 0.9897 - accuracy: 0.5607 - val_loss: 1.6548 - val_accuracy: 0.4221
Epoch 24/30
67/67 [=====] - 1s 10ms/step - loss: 0.9788 - accuracy: 0.5551 - val_loss: 1.5858 - val_accuracy: 0.4259
Epoch 25/30
67/67 [=====] - 1s 10ms/step - loss: 0.9565 - accuracy: 0.5795 - val_loss: 1.6495 - val_accuracy: 0.4203
Epoch 26/30
```

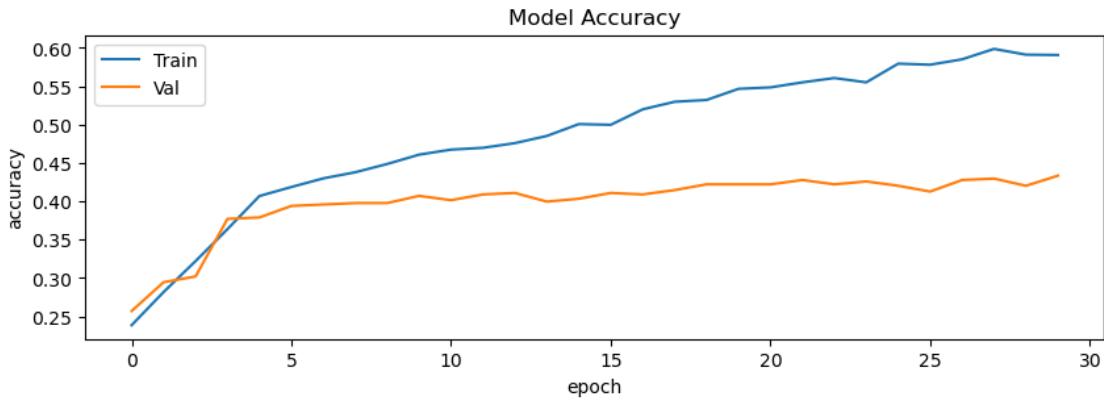
```
67/67 [=====] - 1s 10ms/step - loss: 0.9453 - accuracy: 0.5781 - val_loss: 1.6114 - val_accuracy: 0.4128
Epoch 27/30
67/67 [=====] - 1s 9ms/step - loss: 0.9274 - accuracy: 0.5851 - val_loss: 1.6858 - val_accuracy: 0.4278
Epoch 28/30
67/67 [=====] - 1s 10ms/step - loss: 0.9094 - accuracy: 0.5987 - val_loss: 1.7191 - val_accuracy: 0.4296
Epoch 29/30
67/67 [=====] - 1s 10ms/step - loss: 0.9274 - accuracy: 0.5912 - val_loss: 1.7337 - val_accuracy: 0.4203
Epoch 30/30
67/67 [=====] - 1s 9ms/step - loss: 0.9099 - accuracy: 0.5907 - val_loss: 1.7845 - val_accuracy: 0.4334
```

3.3 Visualize and analyze the model

Now, we'll have a look at how well the model performs.

```
[30]: visualize_results(history)
```





We can see from the learning curves that the model slowly converges and that the model therefore is underfitting. Both the training and validation accuracy curves plateau at a low value, indicating that the model is not learning the patterns in the data well enough.

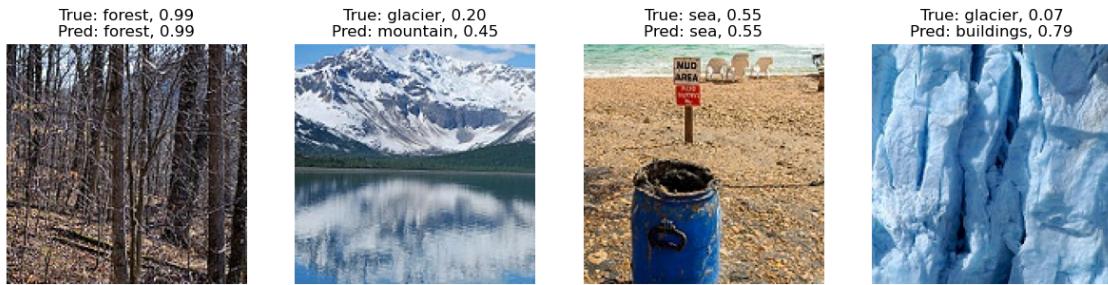
3.3.1 Make predictions on the test dataset

```
[31]: show_images_with_predictions(model_underfitting)
```

```
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 8ms/step
```

```
2023-03-16 15:36:29.549993: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
1/1 [=====] - 0s 8ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 8ms/step
```



3.4 Evaluate the performance of the model

```
[32]: print(f'Accuracy: {model_underfitting.evaluate(test_ds)[1]}')
```

79/79 [=====] - 0s 6ms/step - loss: 1.8681 - accuracy:

0.4158

Accuracy: 0.41576629877090454

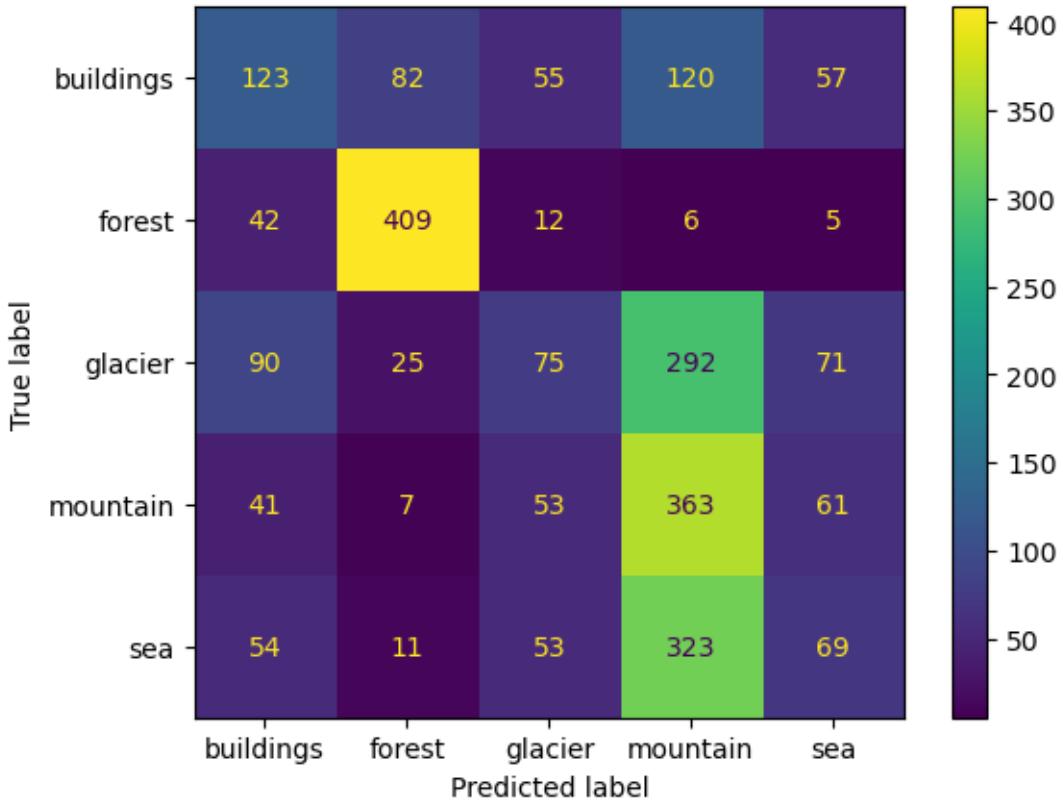
3.5 Confusion Matrix

The confusion matrix is another way to visualize the performance of the model and whether it is underfitting or not.

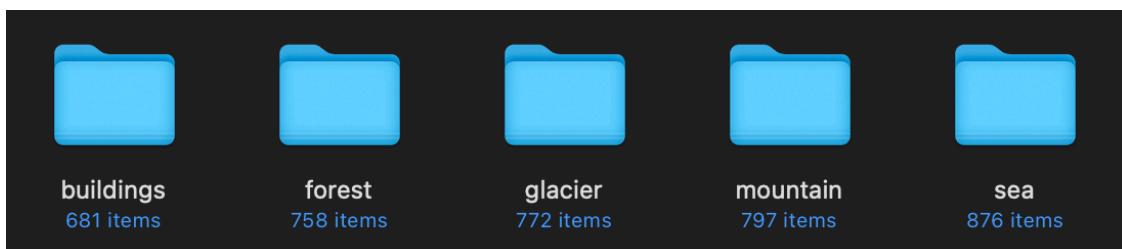
```
[33]: true_labels, pred_labels = show_confusion_matrix(model_underfitting)
```



```
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 8ms/step
1/1 [=====] - 0s 8ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 8ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 21ms/step
```



We noticed that the underfitting model tries to stick to one class for the most part (vertical line in the matrix). We suspect the reason for this being the amount of images assigned to that class in the training set. The classifier knows one class best and tries to match the most images to this class as well. Depending on who ran the notebook, it might be a different class the classifier tries to stick to, as we deleted images of the dataset manually, resulting in different class sizes per person.



3.6 Classification Report

```
[34]: print('Classification Report:\n', metrics.
    classification_report(y_true=true_labels, y_pred=pred_labels))
```

```
Classification Report:
precision    recall   f1-score   support
          0       0       0       0
          1       1       1       1
          2       2       2       2
          3       3       3       3
          4       4       4       4
```

0	0.35	0.28	0.31	437
1	0.77	0.86	0.81	474
2	0.30	0.14	0.19	553
3	0.33	0.69	0.45	525
4	0.26	0.14	0.18	510
accuracy			0.42	2499
macro avg		0.40	0.42	0.39
weighted avg		0.40	0.42	0.38
				2499

The classification report shows precision, recall, f1-score and support. Looking at the classification report above, for our underfitting model, we can see it is pretty bad. With a precision (proportion of true positive (TP) predictions among all positive predictions) rate low for all classes, same with the recall score (proportion of TP among all TP and FP), resulting in a low f1-score (mean of precision and recall), it means the model is not good enough for our task.

Support is the number of actual occurrences (images) of a class in the specified dataset. Unbalanced support in the training data may indicate structural weaknesses in the reported results of the classifier and could mean that stratified sampling or rebalancing is needed, which might be the case with our dataset in combination with a weak model like this (see confusion matrix above, last part).

4 Overfitting Model

This model will be too complex and will adapt perfectly on the training images but will perform poorly on the test data.

The learning curve of the overfitting model will show that the accuracy on the training set is high and improving over time, while the accuracy on the validation set is lower and plateauing or even decreasing over time.

```
[8]: show_images("Example Images", augment=False)
```

Example Images



4.1 Building the Model

The SparseCategoricalCrossentropy calculates the loss of a categorical model with labels provided as integer values.

```
[9]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

      def create_model_overfitting():
          model = Sequential()

          model.add(Conv2D(32, (3,3), input_shape= (img_size,img_size,3), activation='relu', padding = 'same')) #padding = same size output
          model.add(MaxPooling2D())

          model.add(Conv2D(64, (3,3), activation = 'relu', padding = 'same'))
          model.add(MaxPooling2D())

          model.add(Conv2D(128, (3,3), activation = 'relu', padding = 'same'))
          model.add(MaxPooling2D())
```

```

model.add(Conv2D(256, (3,3), activation = 'relu', padding = 'same'))
model.add(MaxPooling2D())

model.add(Conv2D(512, (3,3), activation = 'relu', padding = 'same'))
model.add(MaxPooling2D())

model.add(Flatten())

model.add(Dense(256, activation='relu'))
model.add(Dense(5, activation = 'softmax'))

model.compile(optimizer = 'adam', loss='sparse_categorical_crossentropy',
               metrics=['accuracy'])

model._name = "overfitting_model"

return model

```

[13]: overfitting_model = create_model_overfitting()
overfitting_model.summary()

Model: "overfitting_model"

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 150, 150, 32)	896
max_pooling2d_5 (MaxPooling 2D)	(None, 75, 75, 32)	0
conv2d_6 (Conv2D)	(None, 75, 75, 64)	18496
max_pooling2d_6 (MaxPooling 2D)	(None, 37, 37, 64)	0
conv2d_7 (Conv2D)	(None, 37, 37, 128)	73856
max_pooling2d_7 (MaxPooling 2D)	(None, 18, 18, 128)	0
conv2d_8 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_8 (MaxPooling 2D)	(None, 9, 9, 256)	0
conv2d_9 (Conv2D)	(None, 9, 9, 512)	1180160

```

max_pooling2d_9 (MaxPooling  (None, 4, 4, 512)      0
2D)

flatten_1 (Flatten)      (None, 8192)      0

dense_2 (Dense)         (None, 256)       2097408

dense_3 (Dense)         (None, 5)        1285

=====
Total params: 3,667,269
Trainable params: 3,667,269
Non-trainable params: 0
-----
```

Having this many trainable parameters means the model will be able to adapt i.e. fit to the training data much better than the underfitting model ever could.

4.2 Training the Model

Learnings: In each epoch all images are sent into the CNN and through backpropagation the weights get adapted. The output of the `model.fit()` process outputs e.g.

```
Epoch 1/20
35/163 [=====]
```

which indicates that the CNN is currently training the 35th batch of 162 batches of the first epoch. If we use a `batch_size` of 32 we can then calculate the number of images included in the training step:

$$163 * 32 = 5216$$

The number of epochs defines how many times the CNN sees the **entire** dataset but in a different order.

Another outcome that we've encountered is that the `model.fit()` process needs at least 30 epochs to converge sufficiently so that the overfitting phenomenon is visible in the loss and accuracy graph.

```
[14]: history = overfitting_model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=50,
)
```

```
Epoch 1/50
2023-03-17 12:12:02.524255: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

134/134 [=====] - ETA: 0s - loss: 4.2958 - accuracy: 0.4351
```

```
2023-03-17 12:12:06.237958: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

134/134 [=====] - 4s 30ms/step - loss: 4.2958 -
accuracy: 0.4351 - val_loss: 1.0116 - val_accuracy: 0.6004
Epoch 2/50
134/134 [=====] - 4s 27ms/step - loss: 1.0456 -
accuracy: 0.5602 - val_loss: 0.9129 - val_accuracy: 0.6210
Epoch 3/50
134/134 [=====] - 4s 26ms/step - loss: 0.9135 -
accuracy: 0.6235 - val_loss: 0.7844 - val_accuracy: 0.6679
Epoch 4/50
134/134 [=====] - 4s 27ms/step - loss: 0.8120 -
accuracy: 0.6840 - val_loss: 0.7338 - val_accuracy: 0.7430
Epoch 5/50
134/134 [=====] - 4s 27ms/step - loss: 0.7020 -
accuracy: 0.7309 - val_loss: 0.8673 - val_accuracy: 0.6660
Epoch 6/50
134/134 [=====] - 4s 27ms/step - loss: 0.6380 -
accuracy: 0.7750 - val_loss: 0.8998 - val_accuracy: 0.7261
Epoch 7/50
134/134 [=====] - 4s 27ms/step - loss: 0.5671 -
accuracy: 0.7815 - val_loss: 0.8938 - val_accuracy: 0.7298
Epoch 8/50
134/134 [=====] - 4s 26ms/step - loss: 0.5067 -
accuracy: 0.8078 - val_loss: 0.7391 - val_accuracy: 0.7786
Epoch 9/50
134/134 [=====] - 4s 27ms/step - loss: 0.4435 -
accuracy: 0.8350 - val_loss: 0.6521 - val_accuracy: 0.7899
Epoch 10/50
134/134 [=====] - 4s 27ms/step - loss: 0.3514 -
accuracy: 0.8762 - val_loss: 0.9332 - val_accuracy: 0.7373
Epoch 11/50
134/134 [=====] - 4s 26ms/step - loss: 0.3088 -
accuracy: 0.8851 - val_loss: 0.9467 - val_accuracy: 0.7561
Epoch 12/50
134/134 [=====] - 4s 26ms/step - loss: 0.2898 -
accuracy: 0.8959 - val_loss: 0.9348 - val_accuracy: 0.7561
Epoch 13/50
134/134 [=====] - 4s 26ms/step - loss: 0.1674 -
accuracy: 0.9353 - val_loss: 0.8706 - val_accuracy: 0.7899
Epoch 14/50
134/134 [=====] - 4s 26ms/step - loss: 0.3473 -
accuracy: 0.8795 - val_loss: 0.9807 - val_accuracy: 0.7542
Epoch 15/50
134/134 [=====] - 4s 27ms/step - loss: 0.2139 -
accuracy: 0.9316 - val_loss: 0.9757 - val_accuracy: 0.7861
```

```
Epoch 16/50
134/134 [=====] - 4s 27ms/step - loss: 0.1208 -
accuracy: 0.9573 - val_loss: 1.2155 - val_accuracy: 0.7598
Epoch 17/50
134/134 [=====] - 4s 27ms/step - loss: 0.1340 -
accuracy: 0.9545 - val_loss: 1.4216 - val_accuracy: 0.7392
Epoch 18/50
134/134 [=====] - 4s 27ms/step - loss: 0.2194 -
accuracy: 0.9301 - val_loss: 1.1514 - val_accuracy: 0.7617
Epoch 19/50
134/134 [=====] - 4s 27ms/step - loss: 0.2282 -
accuracy: 0.9250 - val_loss: 1.5378 - val_accuracy: 0.7523
Epoch 20/50
134/134 [=====] - 4s 28ms/step - loss: 0.1363 -
accuracy: 0.9531 - val_loss: 1.5480 - val_accuracy: 0.7129
Epoch 21/50
134/134 [=====] - 4s 27ms/step - loss: 0.1230 -
accuracy: 0.9625 - val_loss: 1.2634 - val_accuracy: 0.7580
Epoch 22/50
134/134 [=====] - 4s 27ms/step - loss: 0.0766 -
accuracy: 0.9775 - val_loss: 1.6506 - val_accuracy: 0.7430
Epoch 23/50
134/134 [=====] - 4s 27ms/step - loss: 0.2123 -
accuracy: 0.9325 - val_loss: 1.2913 - val_accuracy: 0.7448
Epoch 24/50
134/134 [=====] - 4s 27ms/step - loss: 0.1714 -
accuracy: 0.9512 - val_loss: 1.7036 - val_accuracy: 0.7317
Epoch 25/50
134/134 [=====] - 4s 27ms/step - loss: 0.0898 -
accuracy: 0.9733 - val_loss: 1.3596 - val_accuracy: 0.7711
Epoch 26/50
134/134 [=====] - 4s 27ms/step - loss: 0.1433 -
accuracy: 0.9583 - val_loss: 1.6258 - val_accuracy: 0.7223
Epoch 27/50
134/134 [=====] - 4s 27ms/step - loss: 0.1014 -
accuracy: 0.9709 - val_loss: 2.1031 - val_accuracy: 0.7280
Epoch 28/50
134/134 [=====] - 4s 27ms/step - loss: 0.1946 -
accuracy: 0.9423 - val_loss: 1.7409 - val_accuracy: 0.7054
Epoch 29/50
134/134 [=====] - 4s 27ms/step - loss: 0.1263 -
accuracy: 0.9583 - val_loss: 1.4039 - val_accuracy: 0.7711
Epoch 30/50
134/134 [=====] - 4s 27ms/step - loss: 0.0313 -
accuracy: 0.9897 - val_loss: 1.3771 - val_accuracy: 0.7692
Epoch 31/50
134/134 [=====] - 4s 27ms/step - loss: 0.0306 -
accuracy: 0.9892 - val_loss: 1.8029 - val_accuracy: 0.7561
```

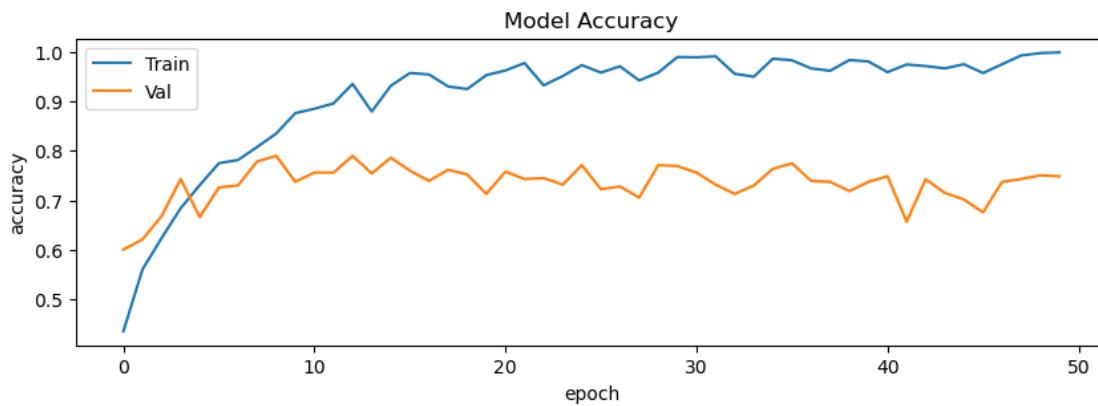
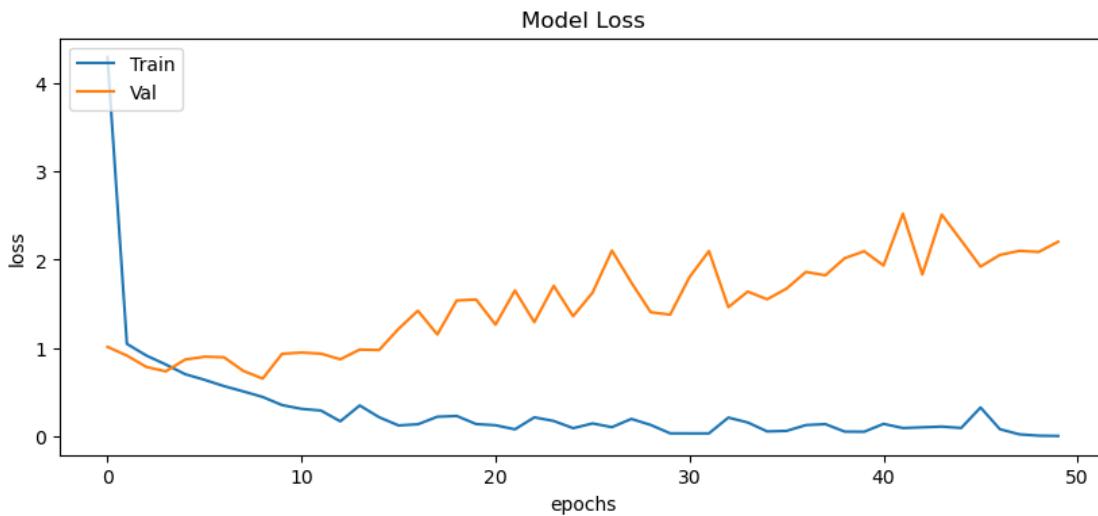
Epoch 32/50
134/134 [=====] - 4s 26ms/step - loss: 0.0305 -
accuracy: 0.9911 - val_loss: 2.0971 - val_accuracy: 0.7317
Epoch 33/50
134/134 [=====] - 4s 27ms/step - loss: 0.2099 -
accuracy: 0.9559 - val_loss: 1.4614 - val_accuracy: 0.7129
Epoch 34/50
134/134 [=====] - 4s 27ms/step - loss: 0.1543 -
accuracy: 0.9498 - val_loss: 1.6383 - val_accuracy: 0.7298
Epoch 35/50
134/134 [=====] - 4s 27ms/step - loss: 0.0530 -
accuracy: 0.9864 - val_loss: 1.5504 - val_accuracy: 0.7636
Epoch 36/50
134/134 [=====] - 4s 26ms/step - loss: 0.0589 -
accuracy: 0.9831 - val_loss: 1.6722 - val_accuracy: 0.7749
Epoch 37/50
134/134 [=====] - 4s 27ms/step - loss: 0.1248 -
accuracy: 0.9667 - val_loss: 1.8611 - val_accuracy: 0.7392
Epoch 38/50
134/134 [=====] - 4s 27ms/step - loss: 0.1355 -
accuracy: 0.9620 - val_loss: 1.8228 - val_accuracy: 0.7373
Epoch 39/50
134/134 [=====] - 4s 27ms/step - loss: 0.0507 -
accuracy: 0.9836 - val_loss: 2.0171 - val_accuracy: 0.7186
Epoch 40/50
134/134 [=====] - 4s 27ms/step - loss: 0.0492 -
accuracy: 0.9808 - val_loss: 2.0954 - val_accuracy: 0.7373
Epoch 41/50
134/134 [=====] - 4s 27ms/step - loss: 0.1384 -
accuracy: 0.9592 - val_loss: 1.9338 - val_accuracy: 0.7486
Epoch 42/50
134/134 [=====] - 4s 27ms/step - loss: 0.0911 -
accuracy: 0.9747 - val_loss: 2.5219 - val_accuracy: 0.6567
Epoch 43/50
134/134 [=====] - 4s 27ms/step - loss: 0.0995 -
accuracy: 0.9714 - val_loss: 1.8321 - val_accuracy: 0.7430
Epoch 44/50
134/134 [=====] - 4s 27ms/step - loss: 0.1076 -
accuracy: 0.9667 - val_loss: 2.5106 - val_accuracy: 0.7148
Epoch 45/50
134/134 [=====] - 4s 27ms/step - loss: 0.0920 -
accuracy: 0.9752 - val_loss: 2.2181 - val_accuracy: 0.7017
Epoch 46/50
134/134 [=====] - 4s 27ms/step - loss: 0.3241 -
accuracy: 0.9573 - val_loss: 1.9214 - val_accuracy: 0.6754
Epoch 47/50
134/134 [=====] - 4s 27ms/step - loss: 0.0777 -
accuracy: 0.9752 - val_loss: 2.0541 - val_accuracy: 0.7373

```

Epoch 48/50
134/134 [=====] - 4s 27ms/step - loss: 0.0199 -
accuracy: 0.9930 - val_loss: 2.0999 - val_accuracy: 0.7430
Epoch 49/50
134/134 [=====] - 4s 27ms/step - loss: 0.0048 -
accuracy: 0.9977 - val_loss: 2.0878 - val_accuracy: 0.7505
Epoch 50/50
134/134 [=====] - 4s 27ms/step - loss: 0.0014 -
accuracy: 0.9991 - val_loss: 2.2037 - val_accuracy: 0.7486

```

```
[15]: visualize_results(history)
```



The accuracy graph makes visible that the train accuracy will come really close or actually is 1 during the training phase. That means the classifier has adapted itself very well to the training dataset. Which also means it has adapted itself very well to the noise.

The validation dataset contains pictures that the classifier has not seen before, and as the model fits too much to the training data, it cannot predict the classes of the validation data well enough. The loss graph shows the same phenomenon. In that case, the validation loss increases, while the training loss decreases.

```
[16]: overfitting_model.evaluate(test_ds)
```

```
157/157 [=====] - 2s 12ms/step - loss: 2.1794 -
accuracy: 0.7467
```

```
[16]: [2.1793887615203857, 0.7466986775398254]
```

Again the model does not perform as well on new data (this time the testing dataset) as it performed on the training data.

```
[18]: show_images_with_predictions(overfitting_model)
```

```
1/1 [=====] - 0s 13ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
```



Unlike the underfitting model, the overfitting model is occasionally 100% sure an image will belong to a certain class. Which is sometimes true, sometimes not (visualized above for the test dataset).

4.3 Confusion Matrix

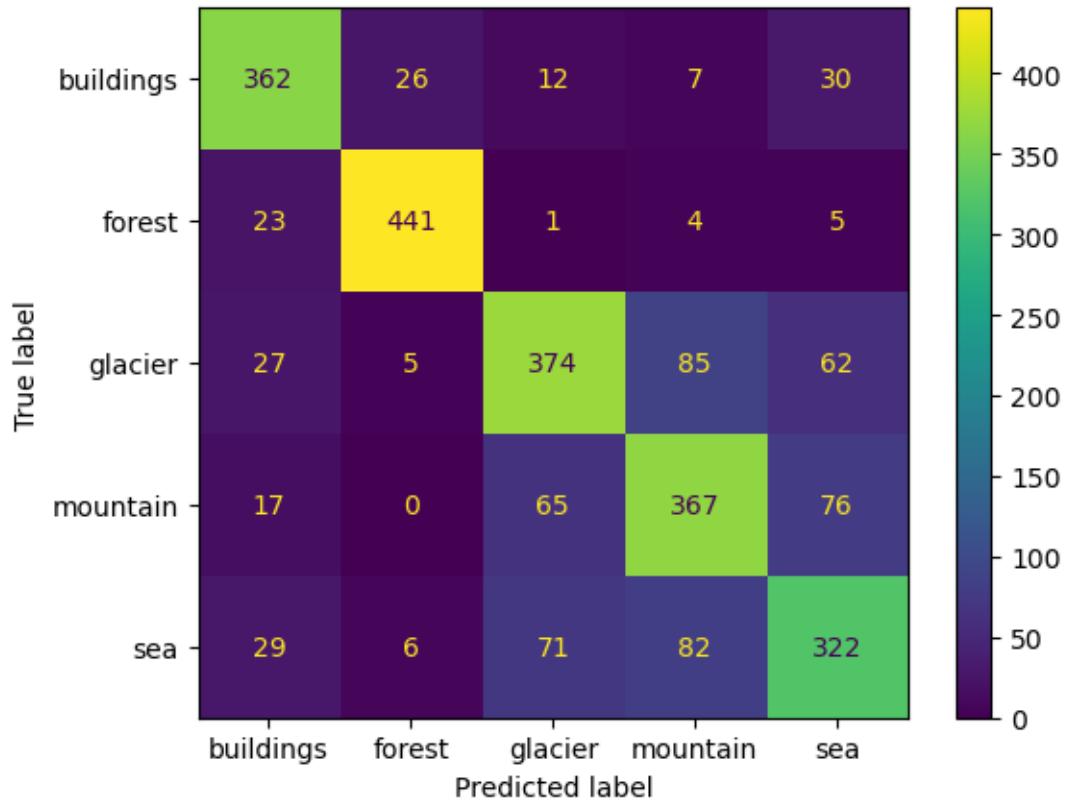
```
[20]: true_labels, pred_labels = show_confusion_matrix(overfitting_model)
```

```
1/1 [=====] - 0s 12ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 12ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
```

1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 12ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 12ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 14ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 12ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step

1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 14ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 9ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 13ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step

```
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 12ms/step
```



The confusion matrix of the overfitting model looks not too bad either. Some predictions are pretty accurate, others are less but in general the model is not performing too bad.

4.4 Classification Report

```
[21]: print('Classification Report:\n', metrics.
    classification_report(y_true=true_labels, y_pred=pred_labels))
```

Classification Report:				
	precision	recall	f1-score	support
0	0.79	0.83	0.81	437
1	0.92	0.93	0.93	474
2	0.72	0.68	0.70	553
3	0.67	0.70	0.69	525
4	0.65	0.63	0.64	510

accuracy		0.75	2499
macro avg	0.75	0.75	2499
weighted avg	0.75	0.75	2499

5 Optimized Model

This model does not show under- or overfitting and performs well on both, training and testing data. Afterwards, a brief description on how to tackle the challenges of an optimal model complexity.

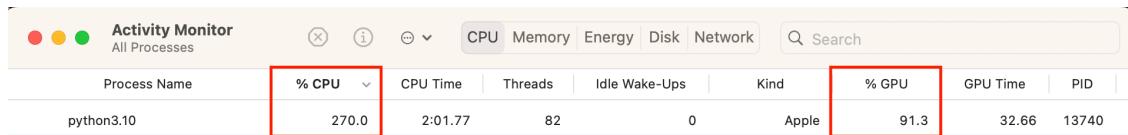
To address underfitting, one approach is to increase the complexity of the model by adding more layers or increasing the number of filters in each layer.

To address overfitting, we can try several approaches. One approach is to simplify the model by removing some layers or decreasing the number of filters in each layer. Another approach is to use less epochs for example (which not always the best approach to be fair).

Adding dropout or weight decay can help to address both of the above mentioned issues. We can also try adjusting the hyperparameters such as learning rate, batch size, or number of epochs.

5.1 GPU Acceleration

We found out that training a model on images can be a very time consuming task, especially if many images are involved. Randomly we stumbled upon an article that stated about “GPU Acceleration” which sounded pretty interesting. Afterwards, we read a bit about the topic and found out that it is possible on our Apple M1 Pro / M2 Pro chips. After a lot of trial and error we ended up using the tensorflow-macos and tensorflow-metal packages but with specific versions because the newest versions did not work. With the both packages installed correctly, both the CPU and GPU are being used for training. The CPU on the M2 Pro chip has 12 cores and its GPU has 19 cores. The total of 31 cores over 12 cores already indicates that usage of the GPU would lead to massive performance improvements. The following image shows the CPU and GPU usage during training of the optimized model:

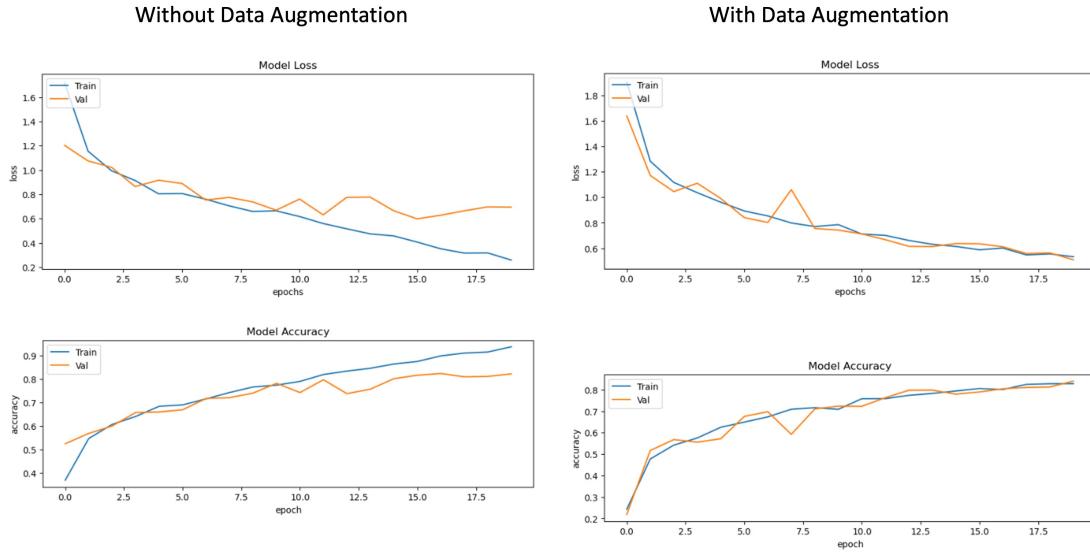


When we were utilizing the CPU only, training took ~100s per epoch for the overfitting model (there were ~11'000'000 trainable parameters, our overfitting model has changed since). After enabling GPU acceleration we came down to astonishing ~10s per epoch. This factor of 10 improvement was even more than we hoped for and enabled us to test more parameters, because training took about 2 minutes instead of the usual 20 minutes.

5.2 Data Augmentation

Data Augmentation has proven to be a useful tool for model generalization. We’ve compared the same model with and without data augmentation and the results show that the model without data augmentation tends to overfit while the model with data augmentation doesn’t:

Usage of Data Augmentation



The reason for this is that data augmentation reduces the generalization error. It introduces new samples to the dataset, by rotating, zooming, and otherwise altering existing images. This increases the size of the dataset, forces the model to learn features of the dataset in a different way, and thus reduces the model's tendency to overfit (more on that later).

During this project we've also encountered some downsides of data augmentation. One apparent downside is that it drastically slows down the training because the images are being preprocessed and the amount of data increases. Because of this reason, we've turned off data augmentation for the KFold Cross Validation.

```
[22]: show_images("Augmented Images (flipped, rotated, zoomed)", augment=True)
```

Augmented Images (flipped, rotated, zoomed)



5.3 Building the Model

As a baseline we used the same model as the overfitting one but we add some extras to reduce the overfitting behavior.

5.3.1 Regularization

Regularization is used to reduce the impact of the weights. The weights then end up having less impact on the loss function which determines the error between the actual label and predicted label. This reduces complexity of the model and therefore reduces overfitting. We are adding the same regularization parameter to all of the Conv2D layers. We are using L2 (Ridge) regularization since it is predetermined from the task.

Dropout Layers The benefit of using dropout is, as mentioned earlier, that no node in the network will be assigned with high parameter values, as a result the parameter values will be dependent on multiple nodes of the previous layer, but instead each node of the previous layer forwards robust and standalone information. E.g. `dropout(0.2)` turns off the nodes of a layer (sets their weights to 0) at a probability of 0.2. Dropout is applied per batch size, so the parameters get tuned for a full batch size with missing input parameters which results in a lower generalization error. Therefore, optimizing dropout parameters along with batch size would make sense for

parameter optimizing using KFold Cross Validation. The results of this experiment can be found below under the topic “Batch Size”.

5.3.2 Generalization

To improve generalization of the model, data augmentation is a useful tool. With data augmentation we can add artificial effects to the images such as flipping and rotating. Through these effects, the images always appear differently each time they are being used in a training step. Therefore the CNN doesn't adapt to the exact images, but rather learns about the relative features inside of an image.

5.3.3 Optimizer

For the optimized model we chose Adam over the competitors because it is the most common among Stochastic Gradient Descent (SGD). We tried out SGD itself, but it performed poorly compared to Adam. Adam optimization is a SGD method that is based on adaptive estimation of first-order (mean) and second-order (uncentered variance) moments. Its default implementation already provides a form of annealed learning, $\text{beta_1}=0.9$ for the first-order moment and $\text{beta_2}=0.999$ for the second-order moment.

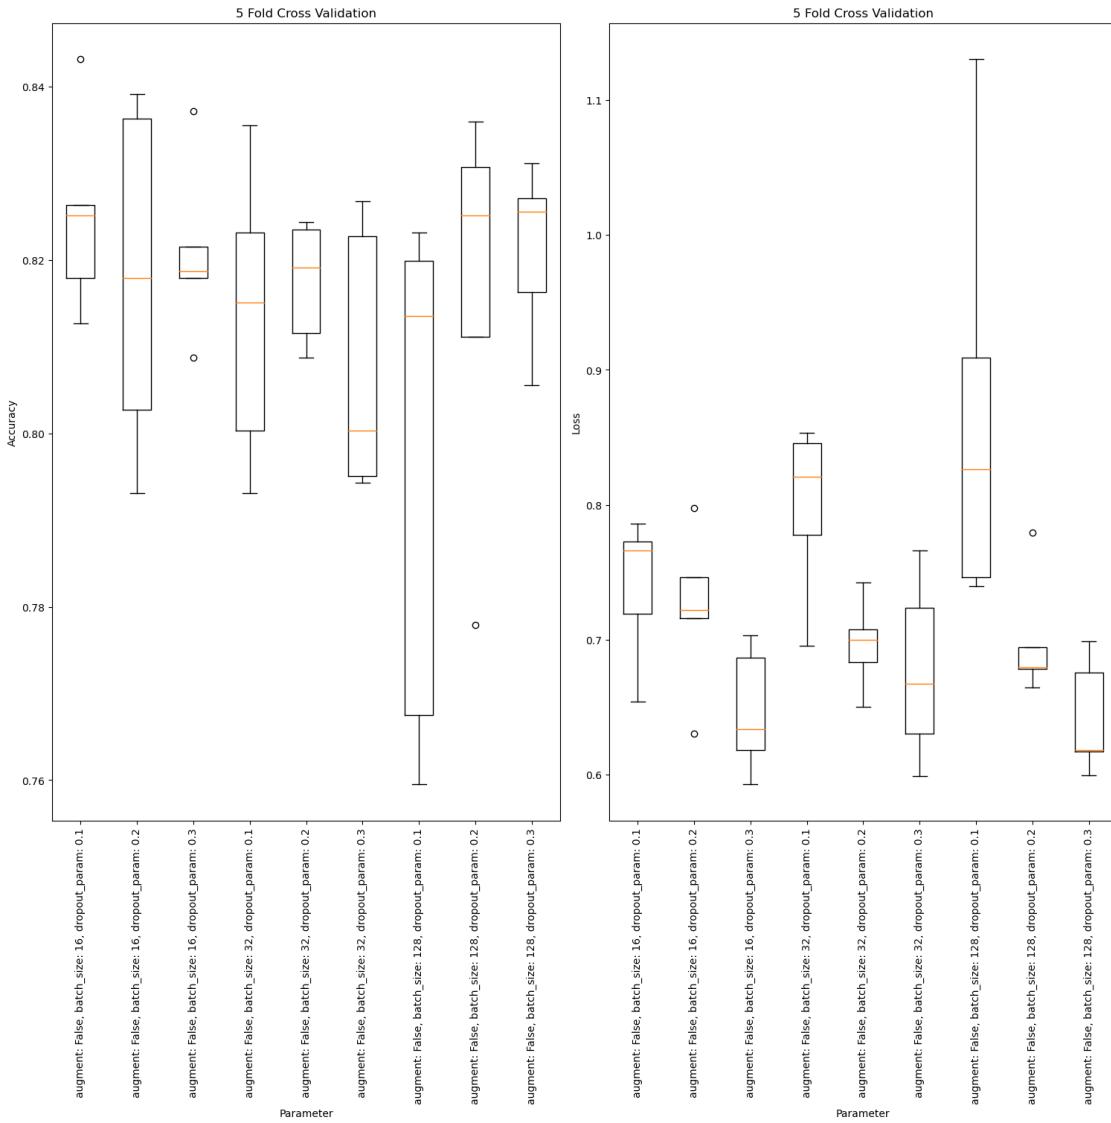
5.3.4 Activation Function

The [following article](#) states that ReLU is the overall the best suited activation function so based on this we decided to use ReLU for our optimized model. In general, non-saturated activation functions tend to be more efficient than saturated activation functions.

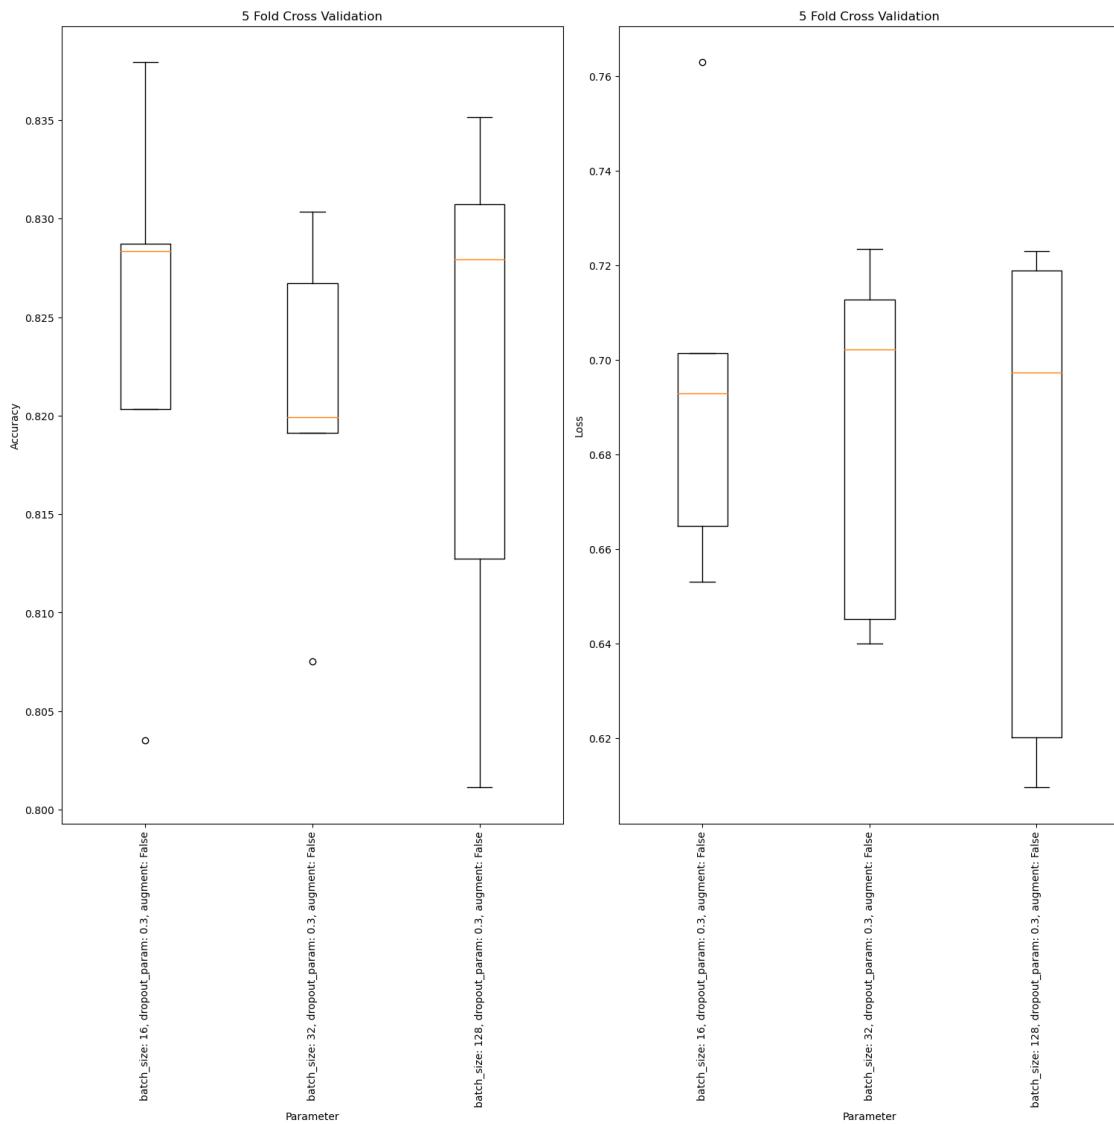
5.3.5 Batch Size

The batch size defines how many samples (in our case images) run through the neural network before the weights get adapted. It is recommended to use mini batches to update the neural network multiple times during an epoch. We've tried out different batch sizes with the same seed on the image generator. To find the optimal batch size we've used the KFold Cross Validation method to see which combination of dropout and batch size works the best. For the dropout parameter we tried out the values 0.1, 0.2 and 0.3 whereas for the batch size we tried out 16, 32 and 128.

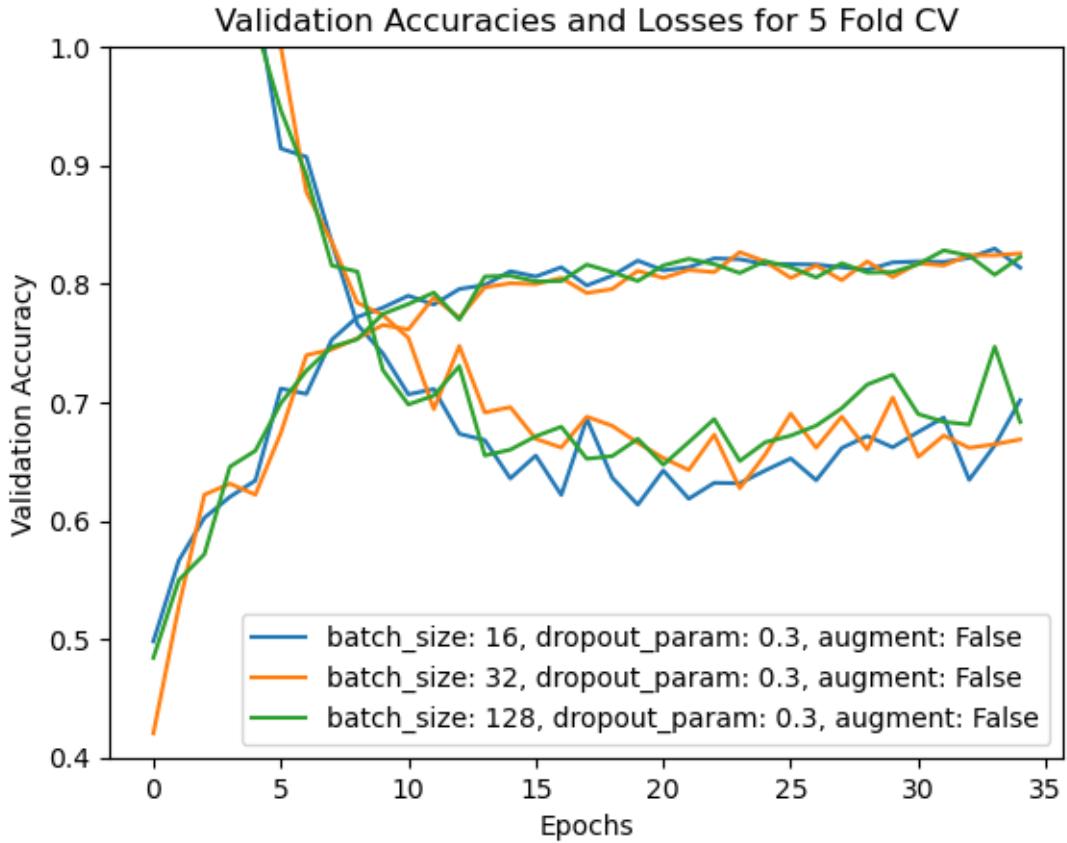
In a first run we obtained the following values:



On a first sight it wasn't clear which parameter is the best. But one thing that was clear is that the loss for dropout=0.3 is always the smallest so we used this parameter to test again for the different batch sizes and got the following diagram:



We also analyzed the training convergence of the three batch sizes:



With all those considerations we did not find an ideal parameter and would need to make more tests in order to find an ideal parameter. But for the sake of this project we took a batch size of 16 paired with a dropout of 30%.

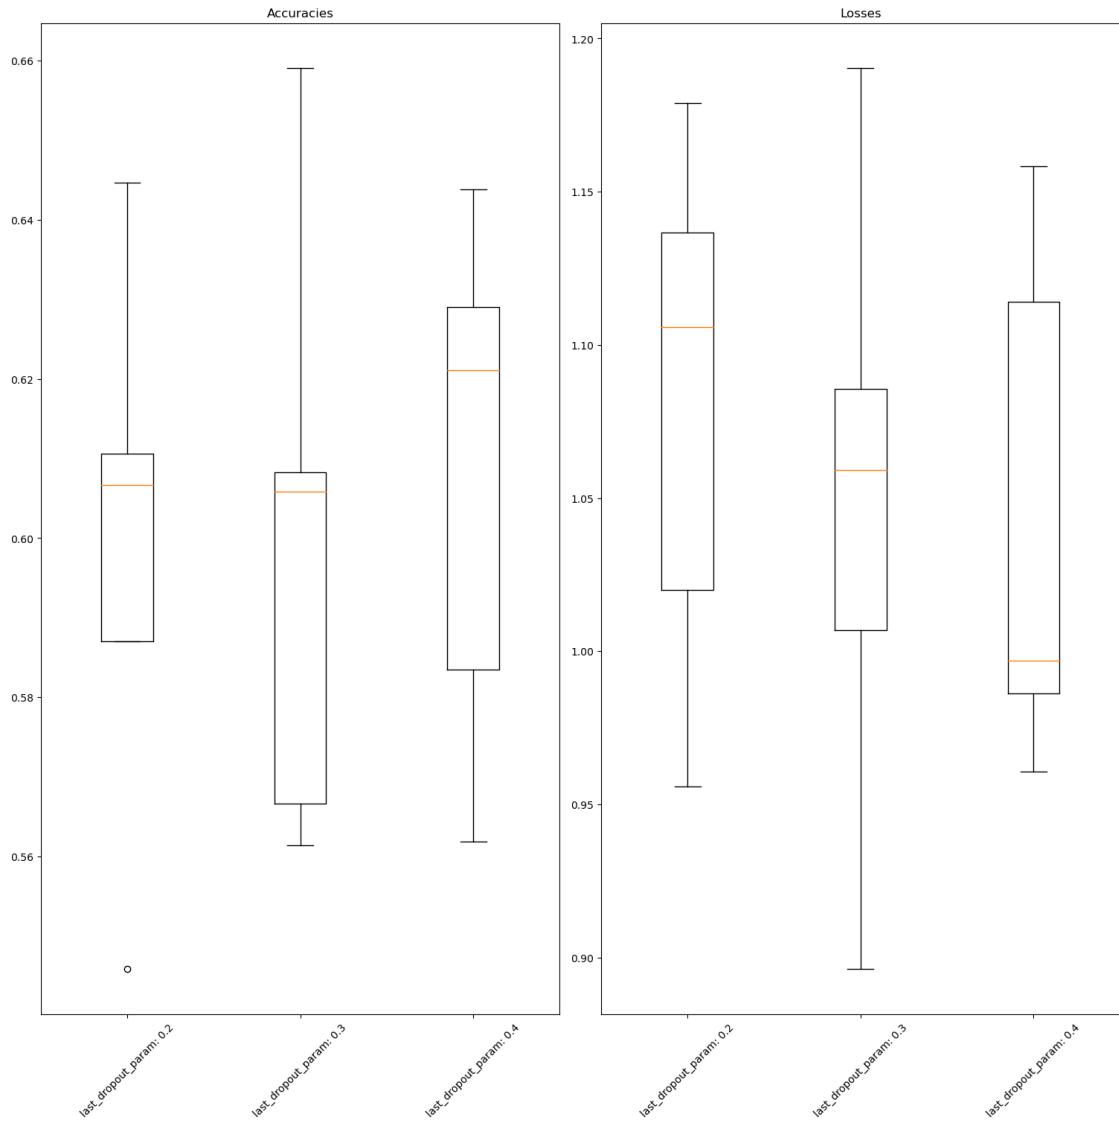
5.4 Layers

5.4.1 First Layer

The RGB channel values are in the $[0, 255]$ range. This is not ideal for a neural network because in general, the input values should be small. Therefore, we use the `BatchNormalization` layer as a normalization layer.

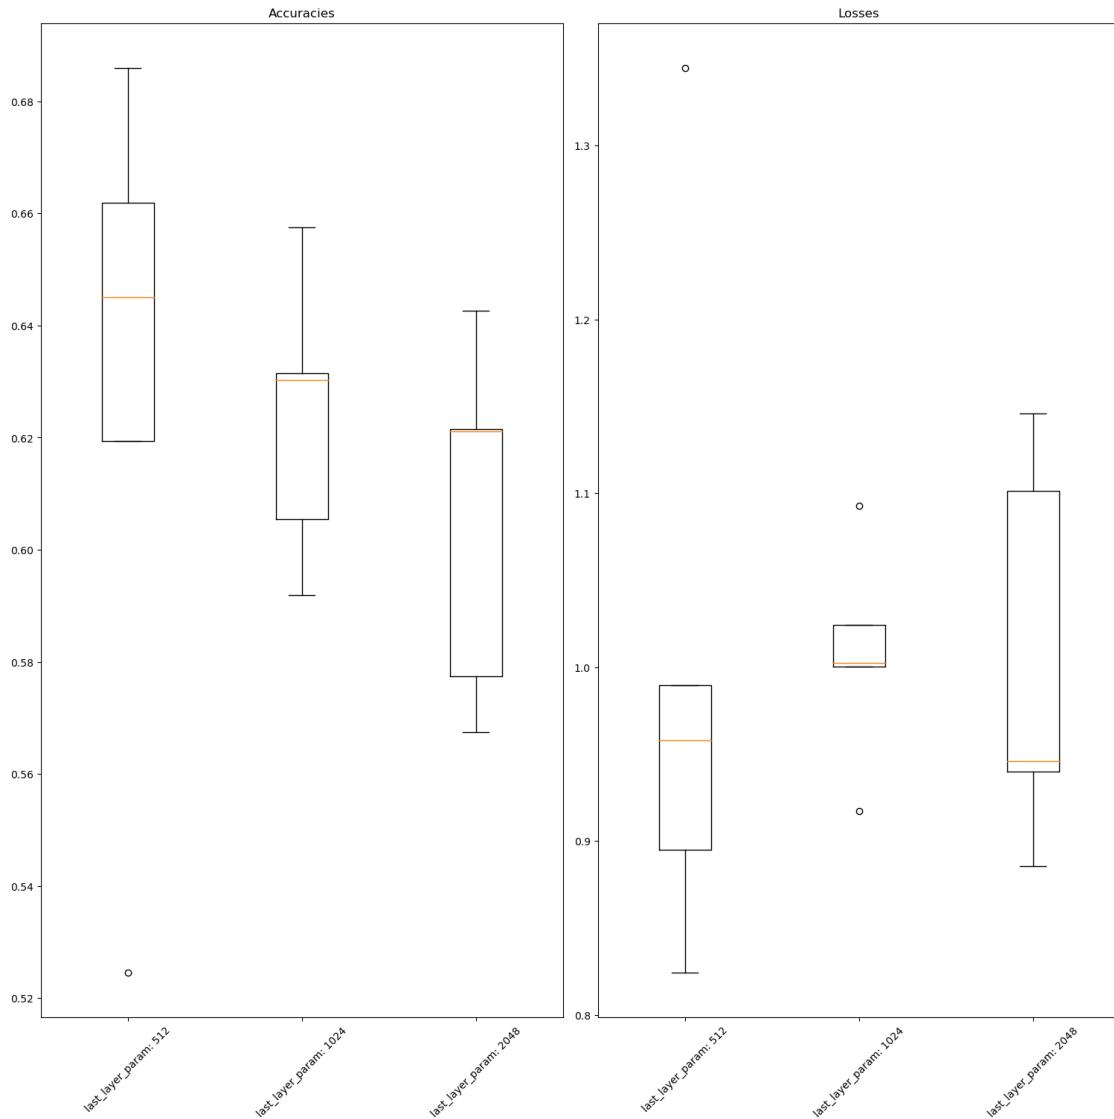
5.4.2 Dropout

With Cross-Validation we've tried out different dropout params (0.2, 0.3, 0.4) for the layer before the dense layers. The results are pretty similar but the 0.3 dropout seems to be more stable at the end of 30 epochs so we used this parameter for the model. The steep spikes of accuracy in the first few epochs are expected because due to dropout, some parameters are missing everytime and the model improves a lot over a short period of time because other parameters are being dropped out at every epoch.



5.4.3 Last Layer

For the last layer we tried out 512, 1024 and 2048 trainable parameters and the Cross-Validation result was that the accuracy and loss were not drastically affected, but the network with 2048 parameters in the last layer was much more stable than the others. Based on this fact, we stucked with using 2048 parameters in the last layer.



5.4.4 Loss Function

Because we use the SoftMax activation function as the output of the last Dense Layer, we get normalized probabilities, [0, 1]. The SparseCategoricalCrossentropy function's `from_logits=True` would expect Logits which are in the range of [-inf, +inf] and therefore, we use `from_logits=False`.

As the optimizer, we also tried out SGD with annealed learning but in the course of 20 epochs we reached poorer results than with adam.

To speed up the training with KFold Cross Validation, we've introduced a parameter to turn off data augmentation and another parameter to use more parameters since they seriously slow down training, which affects Cross Validation too much.

```
[24]: from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, BatchNormalization
       from tensorflow.keras.losses import SparseCategoricalCrossentropy
```

```

from tensorflow.keras.regularizers import l2

def create_model_optimized(l2_param=0.001, dropout_param=0.3,
                           last_layer_param=512, augment=True, weight_factor=1, batch_size=16):
    model = Sequential()

    import_models(batch_size)

    if augment:
        model.add(data_augmentation)
    model.add(BatchNormalization(input_shape=(img_size, img_size, 3)))

    model.add(Conv2D(32 * weight_factor, (3,3), activation = 'relu', padding = 'same'))
    model.add(MaxPooling2D())
    model.add(Dropout(dropout_param))

    model.add(Conv2D(64 * weight_factor, (3,3), activation = 'relu', padding = 'same',
                    kernel_regularizer=l2(l=l2_param)))
    model.add(MaxPooling2D())
    model.add(Dropout(dropout_param))

    model.add(Conv2D(128 * weight_factor, (3,3), activation = 'relu', padding = 'same',
                    kernel_regularizer=l2(l=l2_param)))
    model.add(MaxPooling2D())
    model.add(Dropout(dropout_param))

    model.add(Conv2D(256 * weight_factor, (3,3), activation = 'relu', padding = 'same',
                    kernel_regularizer=l2(l=l2_param)))
    model.add(MaxPooling2D())
    model.add(Dropout(dropout_param))

    model.add(Conv2D(last_layer_param * weight_factor, (3,3), activation = 'relu',
                    padding = 'same', kernel_regularizer=l2(l=l2_param)))
    model.add(MaxPooling2D())
    model.add(Dropout(0.4))

    model.add(Flatten())

    model.add(Dense(256, activation='relu'))
    model.add(Dense(5, activation = 'softmax'))

    model.compile(
        optimizer = 'adam',
        loss=SparseCategoricalCrossentropy(from_logits=False),
        metrics=['accuracy'])

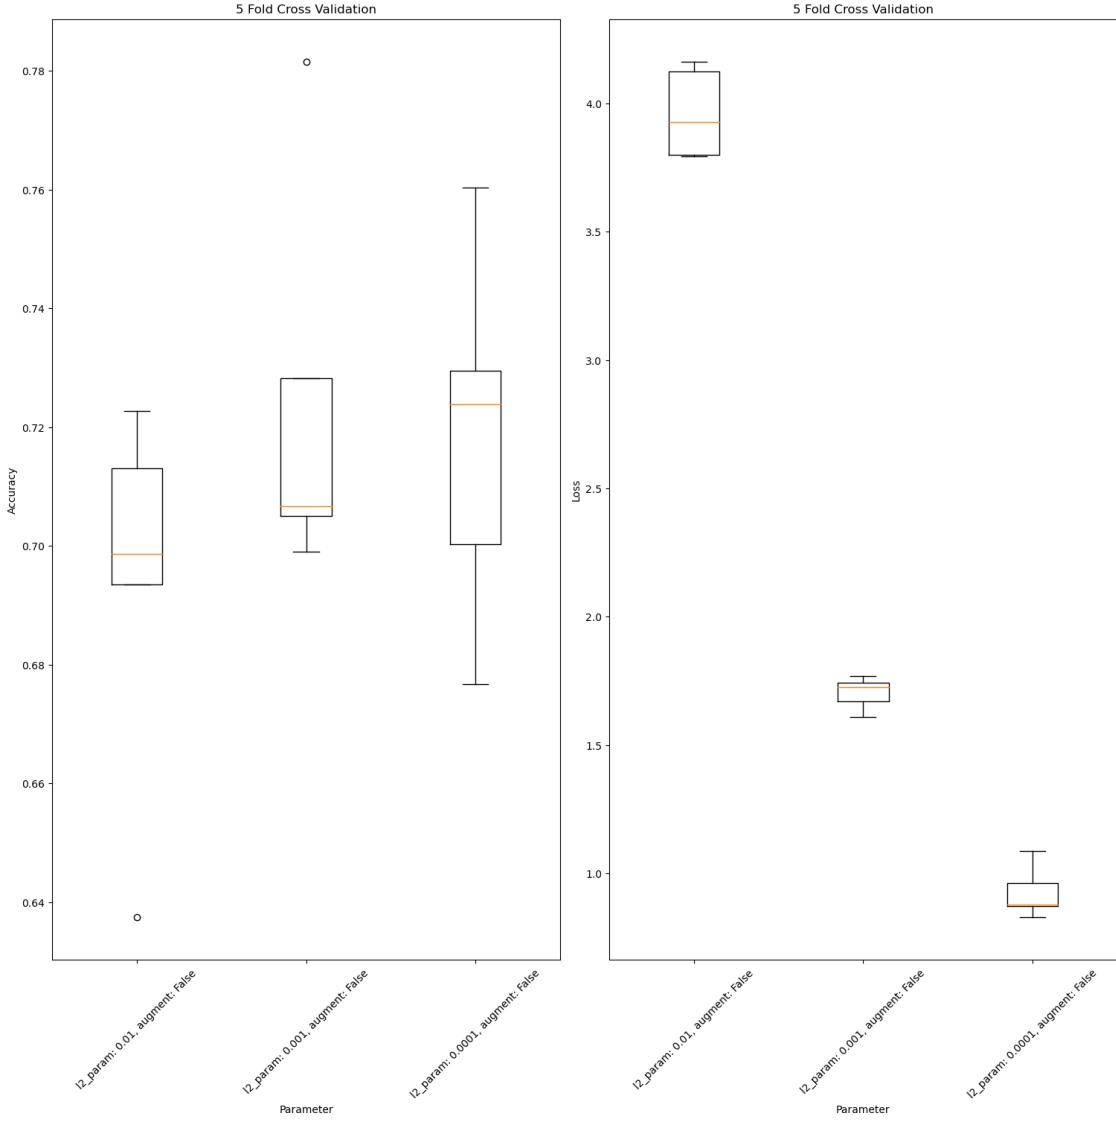
```

```
)  
  
    return model
```

5.5 Cross Validation for Parameter Optimization

For k-fold Cross Validation parameter optimization we are using a sklearn's KFold method paired with a custom functionality to fit the model with different parameter values. We tried to use GridSearchCV (as we did in the AI Foundation course) but we first thought that the model parameters (such as the regularization parameter of the 2nd Conv2D layer) are not accessible from the `param_grid` that is passed to GridSearchCV (which we later found out is not the case). We then came up with a custom KFold Cross Validation approach which is fully flexible and uses if-else statements in the model creation step. With this Cross Validation implementation we've tuned parameters such as the l2 regularization parameter, the number of trainable parameters in the last Conv2D layer or the percentage of dropout before the dense layers.

The following image visualizes the results of the l2 regularization parameter. We tested for the values 0.01, 0.001 and 0.0001 and ran a 5-fold Cross Validation with 10 epochs for each value. The regularizer is being used in three Conv2D layers of the model. As from the diagram below it is apparent that the accuracy improved the smaller the l2 regularization parameter was, and therefore, we decided to use $l2=0.0001$.



5.5.1 Development of the kFold CV Functionality

The part of splitting the full dataset into testing and validation data by the indices obtained by the KFold split was pretty tricky. We've encountered that it is not always convenient and easy to handle the tensorflow data. Especially without experience of the huge framework it was difficult to find information about splitting a dataset and we ended up trying out numerous ways. In the end we consulted ChatGPT which lead us in the right direction. One problem we've had is that the generated filtering functionality did not work directly and after reading the documentation of the filter method, it was not clear which parameters are being passed to this function. Also, it is surprising that tensorflow doesn't seem to have an out-of-the-box solution for such a dataset split. In the end we found out how the methods works and the implementation was successful.

The current solution is not optimal because the `.filter(... reduce_any(...))` is not performant and massively slows down the training which would be important to improve in a real project. Nevertheless, the implementation works and enabled us to make sophisticated comparisons of the model parameters.

During the usage of the intense training functionality of KFold Cross Validation we encountered a problem using tensorflow-metal with GPU acceleration. The problem is that the python process hangs up after approximately the 8-th fold of a 10 epoch training process. This was very inconvenient because we were not able to perform the parameter optimization in one step and needed to temporarily store the results for each parameter which we made by saving the results in a JSON file. Luckily, we were able to perform the 5-fold Cross Validation of at least one parameter before the process hung up. The issue of this behavior is known but does not seem to be fixed yet, as can be seen [here](#).

As already mentioned, we've turned off data augmentation for Cross Validation in favor of performance.

```
[25]: import itertools
from sklearn.model_selection import KFold
import tensorflow as tf
import json

param_grid=dict(
    # last_layer_param=[512, 1024, 2048]
    l2_param=[0.01, 0.001, 0.0001],
    augment=[False],
    # dropout_param=[0.2, 0.3, 0.4]
    # batch_size=[16, 32, 128],
)
epochs = 30
folds = 5

keys = list(param_grid.keys())
params = list(param_grid.get(x) for x in keys)
param_permutations = list(itertools.product(*params))

model_history = dict()

kfold = KFold(n_splits=folds, shuffle=True)

for perm in param_permutations:
    model_args = dict()
    for index in range(len(params)):
        key = keys[index]
        value = perm[index]
        model_args[key] = value

    # Set title according to the params for later visualization
    title = ", ".join(key + ": " + str(model_args[key]) for key in model_args.keys())

    print("CURRENTLY TRAINING THE MODEL WITH THE FOLLOWING PARAMS: " + title)
    model_history[title] = []
```

```

    for fold, (train_indices, val_indices) in enumerate(kfold.
        ↪split(train_images)):
        print(f'Fold {fold+1}/{folds}, (params={title})')

        X_train, y_train = train_images[train_indices], ↪
        ↪train_labels[train_indices]
        X_val, y_val = train_images[val_indices], train_labels[val_indices]

        model = create_model_optimized(**model_args)

        history = model.fit(
            X_train,
            y_train,
            validation_data=(X_val, y_val),
            epochs=epochs
        )

        # Evaluate the model on the validation dataset for this fold
        loss, acc = model.evaluate(test_ds)
        print(f'Test accuracy: {acc:.3f}\n')

        # Use title as key for easy usage
        model_history[title].append(dict(fold=fold, loss=loss, acc=acc, ↪
        ↪history=history))

    with open(f'model-history-{title}.json', 'w') as f:
        json.dump(model_history, f, default=lambda o: '<not serializable>')

```

CURRENTLY TRAINING THE MODEL WITH THE FOLLOWING PARAMS: l2_param: 0.01, augment: False
 Fold 1/5, (params=l2_param: 0.01, augment: False)
 Found 2666 files belonging to 5 classes.
 Using 2133 files for training.
 Found 2666 files belonging to 5 classes.
 Using 533 files for validation.
 Found 2499 files belonging to 5 classes.
 Epoch 1/30

2023-03-17 12:25:23.759974: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - ETA: 0s - loss: 4.9648 - accuracy: 0.3870

2023-03-17 12:25:28.883043: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

```
67/67 [=====] - 6s 71ms/step - loss: 4.9648 - accuracy: 0.3870 - val_loss: 3.6754 - val_accuracy: 0.3052
Epoch 2/30
67/67 [=====] - 4s 61ms/step - loss: 2.1367 - accuracy: 0.5375 - val_loss: 1.6243 - val_accuracy: 0.6049
Epoch 3/30
67/67 [=====] - 5s 71ms/step - loss: 1.4479 - accuracy: 0.5783 - val_loss: 1.3432 - val_accuracy: 0.5206
Epoch 4/30
67/67 [=====] - 4s 58ms/step - loss: 1.1888 - accuracy: 0.6116 - val_loss: 1.1556 - val_accuracy: 0.6030
Epoch 5/30
67/67 [=====] - 4s 56ms/step - loss: 1.0861 - accuracy: 0.6126 - val_loss: 1.0554 - val_accuracy: 0.6554
Epoch 6/30
67/67 [=====] - 4s 55ms/step - loss: 1.0312 - accuracy: 0.6243 - val_loss: 0.9720 - val_accuracy: 0.6498
Epoch 7/30
67/67 [=====] - 4s 56ms/step - loss: 0.9204 - accuracy: 0.6684 - val_loss: 1.0969 - val_accuracy: 0.5562
Epoch 8/30
67/67 [=====] - 4s 55ms/step - loss: 0.9208 - accuracy: 0.6642 - val_loss: 0.9733 - val_accuracy: 0.6667
Epoch 9/30
67/67 [=====] - 4s 56ms/step - loss: 0.8816 - accuracy: 0.6693 - val_loss: 0.8618 - val_accuracy: 0.6891
Epoch 10/30
67/67 [=====] - 4s 56ms/step - loss: 0.8300 - accuracy: 0.7148 - val_loss: 0.8697 - val_accuracy: 0.6929
Epoch 11/30
67/67 [=====] - 4s 56ms/step - loss: 0.8055 - accuracy: 0.7083 - val_loss: 0.9355 - val_accuracy: 0.6592
Epoch 12/30
67/67 [=====] - 4s 56ms/step - loss: 0.7884 - accuracy: 0.7242 - val_loss: 0.8417 - val_accuracy: 0.7060
Epoch 13/30
67/67 [=====] - 4s 56ms/step - loss: 0.7730 - accuracy: 0.7162 - val_loss: 0.7871 - val_accuracy: 0.7191
Epoch 14/30
67/67 [=====] - 4s 56ms/step - loss: 0.7218 - accuracy: 0.7561 - val_loss: 0.7186 - val_accuracy: 0.7715
Epoch 15/30
67/67 [=====] - 4s 56ms/step - loss: 0.7199 - accuracy: 0.7645 - val_loss: 0.7777 - val_accuracy: 0.7360
Epoch 16/30
67/67 [=====] - 4s 56ms/step - loss: 0.6617 - accuracy: 0.7908 - val_loss: 0.7372 - val_accuracy: 0.7697
Epoch 17/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 0.6707 - accuracy: 0.7781 - val_loss: 0.7059 - val_accuracy: 0.7903
Epoch 18/30
67/67 [=====] - 4s 56ms/step - loss: 0.6439 - accuracy: 0.8016 - val_loss: 0.7447 - val_accuracy: 0.7640
Epoch 19/30
67/67 [=====] - 4s 55ms/step - loss: 0.6226 - accuracy: 0.8077 - val_loss: 0.7769 - val_accuracy: 0.7566
Epoch 20/30
67/67 [=====] - 4s 56ms/step - loss: 0.5909 - accuracy: 0.8255 - val_loss: 0.7345 - val_accuracy: 0.7678
Epoch 21/30
67/67 [=====] - 4s 56ms/step - loss: 0.6121 - accuracy: 0.8049 - val_loss: 0.6859 - val_accuracy: 0.7865
Epoch 22/30
67/67 [=====] - 4s 56ms/step - loss: 0.6161 - accuracy: 0.8021 - val_loss: 0.6803 - val_accuracy: 0.7846
Epoch 23/30
67/67 [=====] - 4s 56ms/step - loss: 0.5710 - accuracy: 0.8265 - val_loss: 0.6597 - val_accuracy: 0.8127
Epoch 24/30
67/67 [=====] - 4s 56ms/step - loss: 0.6105 - accuracy: 0.8157 - val_loss: 0.7190 - val_accuracy: 0.7846
Epoch 25/30
67/67 [=====] - 4s 56ms/step - loss: 0.5600 - accuracy: 0.8307 - val_loss: 0.6725 - val_accuracy: 0.7996
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.5302 - accuracy: 0.8382 - val_loss: 0.6957 - val_accuracy: 0.7996
Epoch 27/30
67/67 [=====] - 4s 58ms/step - loss: 0.5483 - accuracy: 0.8419 - val_loss: 0.6593 - val_accuracy: 0.7978
Epoch 28/30
67/67 [=====] - 4s 56ms/step - loss: 0.5674 - accuracy: 0.8302 - val_loss: 0.6535 - val_accuracy: 0.8202
Epoch 29/30
67/67 [=====] - 4s 56ms/step - loss: 0.5272 - accuracy: 0.8410 - val_loss: 0.7283 - val_accuracy: 0.7622
Epoch 30/30
67/67 [=====] - 4s 56ms/step - loss: 0.5537 - accuracy: 0.8340 - val_loss: 0.7368 - val_accuracy: 0.7753
157/157 [=====] - 2s 12ms/step - loss: 0.7121 - accuracy: 0.7707
Test accuracy: 0.771
```

```
Fold 2/5, (params=l2_param: 0.01, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
```

```
Found 2666 files belonging to 5 classes.  
Using 533 files for validation.  
Found 2499 files belonging to 5 classes.  
Epoch 1/30  
  
2023-03-17 12:27:22.335396: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.  
  
67/67 [=====] - ETA: 0s - loss: 4.9318 - accuracy:  
0.3662  
  
2023-03-17 12:27:27.111074: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.  
  
67/67 [=====] - 6s 74ms/step - loss: 4.9318 - accuracy:  
0.3662 - val_loss: 2.7442 - val_accuracy: 0.5066  
Epoch 2/30  
67/67 [=====] - 4s 60ms/step - loss: 1.9792 - accuracy:  
0.5865 - val_loss: 1.6581 - val_accuracy: 0.5610  
Epoch 3/30  
67/67 [=====] - 4s 56ms/step - loss: 1.4150 - accuracy:  
0.5851 - val_loss: 1.6202 - val_accuracy: 0.3996  
Epoch 4/30  
67/67 [=====] - 4s 55ms/step - loss: 1.2005 - accuracy:  
0.6015 - val_loss: 1.1282 - val_accuracy: 0.5891  
Epoch 5/30  
67/67 [=====] - 4s 55ms/step - loss: 1.0768 - accuracy:  
0.6085 - val_loss: 1.1191 - val_accuracy: 0.6210  
Epoch 6/30  
67/67 [=====] - 4s 56ms/step - loss: 1.0188 - accuracy:  
0.6235 - val_loss: 1.0407 - val_accuracy: 0.5872  
Epoch 7/30  
67/67 [=====] - 4s 55ms/step - loss: 0.9570 - accuracy:  
0.6432 - val_loss: 0.9397 - val_accuracy: 0.6585  
Epoch 8/30  
67/67 [=====] - 4s 55ms/step - loss: 0.8984 - accuracy:  
0.6624 - val_loss: 0.9014 - val_accuracy: 0.6623  
Epoch 9/30  
67/67 [=====] - 4s 55ms/step - loss: 0.8996 - accuracy:  
0.6676 - val_loss: 0.9330 - val_accuracy: 0.6642  
Epoch 10/30  
67/67 [=====] - 4s 55ms/step - loss: 0.8234 - accuracy:  
0.7000 - val_loss: 0.7931 - val_accuracy: 0.6998  
Epoch 11/30  
67/67 [=====] - 4s 57ms/step - loss: 0.8326 - accuracy:  
0.7018 - val_loss: 0.9119 - val_accuracy: 0.6923  
Epoch 12/30  
67/67 [=====] - 4s 56ms/step - loss: 0.8099 - accuracy:
```

```
0.7131 - val_loss: 0.8018 - val_accuracy: 0.7205
Epoch 13/30
67/67 [=====] - 4s 55ms/step - loss: 0.7680 - accuracy:
0.7412 - val_loss: 0.7813 - val_accuracy: 0.7054
Epoch 14/30
67/67 [=====] - 4s 55ms/step - loss: 0.7170 - accuracy:
0.7647 - val_loss: 0.7342 - val_accuracy: 0.7355
Epoch 15/30
67/67 [=====] - 4s 56ms/step - loss: 0.6986 - accuracy:
0.7684 - val_loss: 0.8536 - val_accuracy: 0.6754
Epoch 16/30
67/67 [=====] - 4s 57ms/step - loss: 0.7021 - accuracy:
0.7773 - val_loss: 0.7080 - val_accuracy: 0.7580
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.6652 - accuracy:
0.7806 - val_loss: 0.7081 - val_accuracy: 0.7617
Epoch 18/30
67/67 [=====] - 4s 56ms/step - loss: 0.6664 - accuracy:
0.7843 - val_loss: 0.7575 - val_accuracy: 0.7392
Epoch 19/30
67/67 [=====] - 4s 56ms/step - loss: 0.6288 - accuracy:
0.7998 - val_loss: 0.7191 - val_accuracy: 0.7711
Epoch 20/30
67/67 [=====] - 4s 56ms/step - loss: 0.6475 - accuracy:
0.7881 - val_loss: 0.6539 - val_accuracy: 0.7711
Epoch 21/30
67/67 [=====] - 4s 55ms/step - loss: 0.6497 - accuracy:
0.7900 - val_loss: 0.8422 - val_accuracy: 0.6904
Epoch 22/30
67/67 [=====] - 4s 55ms/step - loss: 0.6194 - accuracy:
0.8129 - val_loss: 0.7331 - val_accuracy: 0.7617
Epoch 23/30
67/67 [=====] - 4s 56ms/step - loss: 0.6251 - accuracy:
0.7937 - val_loss: 0.7211 - val_accuracy: 0.7411
Epoch 24/30
67/67 [=====] - 4s 55ms/step - loss: 0.5972 - accuracy:
0.8111 - val_loss: 0.6866 - val_accuracy: 0.7730
Epoch 25/30
67/67 [=====] - 4s 55ms/step - loss: 0.6210 - accuracy:
0.8008 - val_loss: 0.6232 - val_accuracy: 0.7936
Epoch 26/30
67/67 [=====] - 4s 55ms/step - loss: 0.6098 - accuracy:
0.8059 - val_loss: 0.7470 - val_accuracy: 0.7373
Epoch 27/30
67/67 [=====] - 4s 56ms/step - loss: 0.6019 - accuracy:
0.8078 - val_loss: 0.6359 - val_accuracy: 0.8030
Epoch 28/30
67/67 [=====] - 4s 55ms/step - loss: 0.5989 - accuracy:
```

```
0.8111 - val_loss: 0.7676 - val_accuracy: 0.7317
Epoch 29/30
67/67 [=====] - 4s 55ms/step - loss: 0.5863 - accuracy:
0.8143 - val_loss: 0.6468 - val_accuracy: 0.7861
Epoch 30/30
67/67 [=====] - 4s 55ms/step - loss: 0.5701 - accuracy:
0.8204 - val_loss: 0.6556 - val_accuracy: 0.7861
157/157 [=====] - 2s 12ms/step - loss: 0.6746 -
accuracy: 0.7839
Test accuracy: 0.784

Fold 3/5, (params=l2_param: 0.01, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30

2023-03-17 12:29:19.332675: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - ETA: 0s - loss: 4.8179 - accuracy:
0.3755

2023-03-17 12:29:23.660407: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - 5s 67ms/step - loss: 4.8179 - accuracy:
0.3755 - val_loss: 2.7680 - val_accuracy: 0.3959
Epoch 2/30
67/67 [=====] - 4s 57ms/step - loss: 1.8944 - accuracy:
0.5570 - val_loss: 1.5708 - val_accuracy: 0.5216
Epoch 3/30
67/67 [=====] - 4s 58ms/step - loss: 1.2862 - accuracy:
0.5987 - val_loss: 1.1491 - val_accuracy: 0.5985
Epoch 4/30
67/67 [=====] - 4s 56ms/step - loss: 1.1109 - accuracy:
0.6118 - val_loss: 1.1426 - val_accuracy: 0.6173
Epoch 5/30
67/67 [=====] - 4s 56ms/step - loss: 1.0137 - accuracy:
0.6446 - val_loss: 1.0357 - val_accuracy: 0.6004
Epoch 6/30
67/67 [=====] - 4s 58ms/step - loss: 0.9434 - accuracy:
0.6479 - val_loss: 1.0367 - val_accuracy: 0.6135
Epoch 7/30
67/67 [=====] - 4s 55ms/step - loss: 0.9176 - accuracy:
0.6578 - val_loss: 0.9916 - val_accuracy: 0.6116
```

```
Epoch 8/30
67/67 [=====] - 4s 56ms/step - loss: 0.8843 - accuracy: 0.6723 - val_loss: 0.8964 - val_accuracy: 0.6548
Epoch 9/30
67/67 [=====] - 4s 56ms/step - loss: 0.8235 - accuracy: 0.7154 - val_loss: 0.9600 - val_accuracy: 0.6529
Epoch 10/30
67/67 [=====] - 4s 55ms/step - loss: 0.7873 - accuracy: 0.7262 - val_loss: 0.8523 - val_accuracy: 0.6717
Epoch 11/30
67/67 [=====] - 4s 56ms/step - loss: 0.7840 - accuracy: 0.7393 - val_loss: 0.7963 - val_accuracy: 0.7336
Epoch 12/30
67/67 [=====] - 4s 57ms/step - loss: 0.7460 - accuracy: 0.7557 - val_loss: 0.7566 - val_accuracy: 0.7336
Epoch 13/30
67/67 [=====] - 4s 57ms/step - loss: 0.7507 - accuracy: 0.7482 - val_loss: 0.8775 - val_accuracy: 0.6942
Epoch 14/30
67/67 [=====] - 4s 55ms/step - loss: 0.7063 - accuracy: 0.7703 - val_loss: 0.7406 - val_accuracy: 0.7448
Epoch 15/30
67/67 [=====] - 4s 55ms/step - loss: 0.7213 - accuracy: 0.7590 - val_loss: 0.9704 - val_accuracy: 0.6642
Epoch 16/30
67/67 [=====] - 4s 55ms/step - loss: 0.6776 - accuracy: 0.7759 - val_loss: 0.7693 - val_accuracy: 0.7598
Epoch 17/30
67/67 [=====] - 4s 55ms/step - loss: 0.6709 - accuracy: 0.7815 - val_loss: 0.6656 - val_accuracy: 0.7767
Epoch 18/30
67/67 [=====] - 4s 55ms/step - loss: 0.6670 - accuracy: 0.7768 - val_loss: 0.7662 - val_accuracy: 0.7711
Epoch 19/30
67/67 [=====] - 4s 56ms/step - loss: 0.6663 - accuracy: 0.7834 - val_loss: 0.8297 - val_accuracy: 0.7054
Epoch 20/30
67/67 [=====] - 4s 56ms/step - loss: 0.6335 - accuracy: 0.7942 - val_loss: 0.7374 - val_accuracy: 0.7767
Epoch 21/30
67/67 [=====] - 4s 55ms/step - loss: 0.6062 - accuracy: 0.8097 - val_loss: 0.7336 - val_accuracy: 0.7467
Epoch 22/30
67/67 [=====] - 4s 56ms/step - loss: 0.5892 - accuracy: 0.8134 - val_loss: 0.6988 - val_accuracy: 0.7749
Epoch 23/30
67/67 [=====] - 4s 55ms/step - loss: 0.6097 - accuracy: 0.8068 - val_loss: 0.6393 - val_accuracy: 0.8086
```

```
Epoch 24/30
67/67 [=====] - 4s 55ms/step - loss: 0.5880 - accuracy: 0.8134 - val_loss: 0.7555 - val_accuracy: 0.7411
Epoch 25/30
67/67 [=====] - 4s 56ms/step - loss: 0.6095 - accuracy: 0.8022 - val_loss: 0.7018 - val_accuracy: 0.7730
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.6187 - accuracy: 0.7918 - val_loss: 0.6890 - val_accuracy: 0.7842
Epoch 27/30
67/67 [=====] - 4s 55ms/step - loss: 0.5941 - accuracy: 0.8078 - val_loss: 0.7447 - val_accuracy: 0.7711
Epoch 28/30
67/67 [=====] - 4s 57ms/step - loss: 0.5895 - accuracy: 0.8097 - val_loss: 0.7302 - val_accuracy: 0.7767
Epoch 29/30
67/67 [=====] - 4s 55ms/step - loss: 0.5749 - accuracy: 0.8092 - val_loss: 0.6732 - val_accuracy: 0.7917
Epoch 30/30
67/67 [=====] - 4s 55ms/step - loss: 0.5853 - accuracy: 0.8242 - val_loss: 0.6994 - val_accuracy: 0.7936
157/157 [=====] - 2s 12ms/step - loss: 0.7339 - accuracy: 0.7695
Test accuracy: 0.770
```

```
Fold 4/5, (params=l2_param: 0.01, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
```

```
2023-03-17 12:31:15.873418: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
67/67 [=====] - ETA: 0s - loss: 5.2221 - accuracy: 0.3090
```

```
2023-03-17 12:31:20.795341: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
67/67 [=====] - 6s 75ms/step - loss: 5.2221 - accuracy: 0.3090 - val_loss: 3.0902 - val_accuracy: 0.4484
```

```
Epoch 2/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 2.2742 - accuracy: 0.5396 - val_loss: 1.8939 - val_accuracy: 0.5216
```

```
Epoch 3/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 1.5269 - accuracy: 0.5945 - val_loss: 1.3554 - val_accuracy: 0.5872
Epoch 4/30
67/67 [=====] - 4s 56ms/step - loss: 1.2336 - accuracy: 0.6128 - val_loss: 1.1270 - val_accuracy: 0.6492
Epoch 5/30
67/67 [=====] - 4s 56ms/step - loss: 1.1012 - accuracy: 0.6226 - val_loss: 1.0409 - val_accuracy: 0.6473
Epoch 6/30
67/67 [=====] - 4s 56ms/step - loss: 1.0086 - accuracy: 0.6404 - val_loss: 1.0072 - val_accuracy: 0.6417
Epoch 7/30
67/67 [=====] - 4s 56ms/step - loss: 0.9314 - accuracy: 0.6624 - val_loss: 0.9471 - val_accuracy: 0.6904
Epoch 8/30
67/67 [=====] - 4s 56ms/step - loss: 0.9173 - accuracy: 0.6676 - val_loss: 0.9026 - val_accuracy: 0.6642
Epoch 9/30
67/67 [=====] - 4s 55ms/step - loss: 0.8645 - accuracy: 0.6779 - val_loss: 0.8800 - val_accuracy: 0.6717
Epoch 10/30
67/67 [=====] - 4s 55ms/step - loss: 0.8376 - accuracy: 0.7028 - val_loss: 0.8512 - val_accuracy: 0.7129
Epoch 11/30
67/67 [=====] - 4s 55ms/step - loss: 0.8023 - accuracy: 0.7150 - val_loss: 1.0238 - val_accuracy: 0.6604
Epoch 12/30
67/67 [=====] - 4s 56ms/step - loss: 0.7381 - accuracy: 0.7548 - val_loss: 0.7271 - val_accuracy: 0.7992
Epoch 13/30
67/67 [=====] - 4s 55ms/step - loss: 0.7690 - accuracy: 0.7290 - val_loss: 0.7710 - val_accuracy: 0.7824
Epoch 14/30
67/67 [=====] - 4s 56ms/step - loss: 0.7117 - accuracy: 0.7595 - val_loss: 0.7729 - val_accuracy: 0.7561
Epoch 15/30
67/67 [=====] - 4s 56ms/step - loss: 0.6887 - accuracy: 0.7754 - val_loss: 0.8041 - val_accuracy: 0.7411
Epoch 16/30
67/67 [=====] - 4s 56ms/step - loss: 0.7010 - accuracy: 0.7689 - val_loss: 0.7180 - val_accuracy: 0.7711
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.6809 - accuracy: 0.7642 - val_loss: 0.7638 - val_accuracy: 0.7711
Epoch 18/30
67/67 [=====] - 4s 56ms/step - loss: 0.6601 - accuracy: 0.7867 - val_loss: 0.7080 - val_accuracy: 0.7786
Epoch 19/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 0.6172 - accuracy: 0.8003 - val_loss: 0.7063 - val_accuracy: 0.7917
Epoch 20/30
67/67 [=====] - 4s 56ms/step - loss: 0.6323 - accuracy: 0.7918 - val_loss: 0.9511 - val_accuracy: 0.6811
Epoch 21/30
67/67 [=====] - 4s 57ms/step - loss: 0.6087 - accuracy: 0.8026 - val_loss: 0.7474 - val_accuracy: 0.7411
Epoch 22/30
67/67 [=====] - 4s 56ms/step - loss: 0.6009 - accuracy: 0.8054 - val_loss: 0.7073 - val_accuracy: 0.7824
Epoch 23/30
67/67 [=====] - 4s 56ms/step - loss: 0.6404 - accuracy: 0.7932 - val_loss: 0.6189 - val_accuracy: 0.8105
Epoch 24/30
67/67 [=====] - 4s 56ms/step - loss: 0.6056 - accuracy: 0.8054 - val_loss: 0.6407 - val_accuracy: 0.8218
Epoch 25/30
67/67 [=====] - 4s 56ms/step - loss: 0.5839 - accuracy: 0.8012 - val_loss: 0.7471 - val_accuracy: 0.7542
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.6215 - accuracy: 0.8003 - val_loss: 0.6723 - val_accuracy: 0.7974
Epoch 27/30
67/67 [=====] - 4s 56ms/step - loss: 0.6047 - accuracy: 0.8078 - val_loss: 0.6842 - val_accuracy: 0.7917
Epoch 28/30
67/67 [=====] - 4s 56ms/step - loss: 0.5478 - accuracy: 0.8284 - val_loss: 0.6928 - val_accuracy: 0.7917
Epoch 29/30
67/67 [=====] - 4s 57ms/step - loss: 0.5511 - accuracy: 0.8186 - val_loss: 0.6841 - val_accuracy: 0.7786
Epoch 30/30
67/67 [=====] - 4s 56ms/step - loss: 0.5477 - accuracy: 0.8195 - val_loss: 0.8072 - val_accuracy: 0.7280
157/157 [=====] - 2s 12ms/step - loss: 0.7918 - accuracy: 0.7435
Test accuracy: 0.743
```

```
Fold 5/5, (params=l2_param: 0.01, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
```

2023-03-17 12:33:13.179414: I

```
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - ETA: 0s - loss: 4.7657 - accuracy: 0.3671

2023-03-17 12:33:18.257457: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - 6s 73ms/step - loss: 4.7657 - accuracy: 0.3671 - val_loss: 2.6357 - val_accuracy: 0.4334
Epoch 2/30
67/67 [=====] - 4s 59ms/step - loss: 1.8615 - accuracy: 0.5471 - val_loss: 1.4149 - val_accuracy: 0.5722
Epoch 3/30
67/67 [=====] - 4s 56ms/step - loss: 1.2836 - accuracy: 0.5992 - val_loss: 1.1767 - val_accuracy: 0.6135
Epoch 4/30
67/67 [=====] - 4s 56ms/step - loss: 1.0755 - accuracy: 0.6221 - val_loss: 1.0827 - val_accuracy: 0.6041
Epoch 5/30
67/67 [=====] - 4s 56ms/step - loss: 1.0099 - accuracy: 0.6137 - val_loss: 0.9634 - val_accuracy: 0.6435
Epoch 6/30
67/67 [=====] - 4s 56ms/step - loss: 0.9689 - accuracy: 0.6442 - val_loss: 1.0463 - val_accuracy: 0.6116
Epoch 7/30
67/67 [=====] - 4s 56ms/step - loss: 0.9468 - accuracy: 0.6446 - val_loss: 0.9927 - val_accuracy: 0.6417
Epoch 8/30
67/67 [=====] - 4s 56ms/step - loss: 0.9046 - accuracy: 0.6573 - val_loss: 0.9391 - val_accuracy: 0.6285
Epoch 9/30
67/67 [=====] - 4s 56ms/step - loss: 0.8546 - accuracy: 0.6803 - val_loss: 0.9049 - val_accuracy: 0.6717
Epoch 10/30
67/67 [=====] - 4s 55ms/step - loss: 0.8371 - accuracy: 0.6798 - val_loss: 0.9128 - val_accuracy: 0.6604
Epoch 11/30
67/67 [=====] - 4s 56ms/step - loss: 0.8167 - accuracy: 0.7023 - val_loss: 0.8361 - val_accuracy: 0.6867
Epoch 12/30
67/67 [=====] - 4s 56ms/step - loss: 0.7953 - accuracy: 0.7164 - val_loss: 0.8395 - val_accuracy: 0.7167
Epoch 13/30
67/67 [=====] - 4s 55ms/step - loss: 0.7705 - accuracy: 0.7112 - val_loss: 0.8042 - val_accuracy: 0.7392
Epoch 14/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 0.7561 - accuracy: 0.7248 - val_loss: 0.8301 - val_accuracy: 0.7298
Epoch 15/30
67/67 [=====] - 4s 55ms/step - loss: 0.7396 - accuracy: 0.7436 - val_loss: 0.7337 - val_accuracy: 0.7523
Epoch 16/30
67/67 [=====] - 4s 55ms/step - loss: 0.6839 - accuracy: 0.7675 - val_loss: 0.7122 - val_accuracy: 0.7411
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.7210 - accuracy: 0.7557 - val_loss: 0.7880 - val_accuracy: 0.7186
Epoch 18/30
67/67 [=====] - 4s 56ms/step - loss: 0.6823 - accuracy: 0.7675 - val_loss: 0.8727 - val_accuracy: 0.6604
Epoch 19/30
67/67 [=====] - 4s 56ms/step - loss: 0.6735 - accuracy: 0.7712 - val_loss: 0.7388 - val_accuracy: 0.7561
Epoch 20/30
67/67 [=====] - 4s 56ms/step - loss: 0.6901 - accuracy: 0.7750 - val_loss: 0.7632 - val_accuracy: 0.7373
Epoch 21/30
67/67 [=====] - 4s 56ms/step - loss: 0.6639 - accuracy: 0.7811 - val_loss: 0.7214 - val_accuracy: 0.7580
Epoch 22/30
67/67 [=====] - 4s 55ms/step - loss: 0.6495 - accuracy: 0.7904 - val_loss: 0.7538 - val_accuracy: 0.7317
Epoch 23/30
67/67 [=====] - 4s 57ms/step - loss: 0.6445 - accuracy: 0.7843 - val_loss: 0.7707 - val_accuracy: 0.7298
Epoch 24/30
67/67 [=====] - 4s 56ms/step - loss: 0.6310 - accuracy: 0.7895 - val_loss: 0.7602 - val_accuracy: 0.7129
Epoch 25/30
67/67 [=====] - 4s 56ms/step - loss: 0.6439 - accuracy: 0.7853 - val_loss: 0.7497 - val_accuracy: 0.7430
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.6310 - accuracy: 0.7904 - val_loss: 0.7664 - val_accuracy: 0.7467
Epoch 27/30
67/67 [=====] - 4s 55ms/step - loss: 0.6470 - accuracy: 0.7862 - val_loss: 0.8676 - val_accuracy: 0.7111
Epoch 28/30
67/67 [=====] - 4s 55ms/step - loss: 0.6021 - accuracy: 0.7942 - val_loss: 0.6969 - val_accuracy: 0.7692
Epoch 29/30
67/67 [=====] - 4s 56ms/step - loss: 0.5857 - accuracy: 0.8092 - val_loss: 0.7182 - val_accuracy: 0.7505
Epoch 30/30
```

```
67/67 [=====] - 4s 55ms/step - loss: 0.5832 - accuracy: 0.8045 - val_loss: 0.6600 - val_accuracy: 0.7880
157/157 [=====] - 2s 12ms/step - loss: 0.6712 - accuracy: 0.7815
Test accuracy: 0.782

CURRENTLY TRAINING THE MODEL WITH THE FOLLOWING PARAMS: l2_param: 0.001,
augment: False
Fold 1/5, (params=l2_param: 0.001, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30

2023-03-17 12:35:11.708035: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - ETA: 0s - loss: 1.9798 - accuracy: 0.4020

2023-03-17 12:35:17.471308: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - 7s 88ms/step - loss: 1.9798 - accuracy: 0.4020 - val_loss: 1.5573 - val_accuracy: 0.5543
Epoch 2/30
67/67 [=====] - 4s 58ms/step - loss: 1.4534 - accuracy: 0.5516 - val_loss: 1.3568 - val_accuracy: 0.5674
Epoch 3/30
67/67 [=====] - 4s 58ms/step - loss: 1.2097 - accuracy: 0.6187 - val_loss: 1.1476 - val_accuracy: 0.6479
Epoch 4/30
67/67 [=====] - 4s 56ms/step - loss: 1.0647 - accuracy: 0.6506 - val_loss: 1.1454 - val_accuracy: 0.6367
Epoch 5/30
67/67 [=====] - 4s 56ms/step - loss: 0.9403 - accuracy: 0.6975 - val_loss: 1.0088 - val_accuracy: 0.6910
Epoch 6/30
67/67 [=====] - 4s 55ms/step - loss: 0.8735 - accuracy: 0.7162 - val_loss: 0.8930 - val_accuracy: 0.6873
Epoch 7/30
67/67 [=====] - 4s 56ms/step - loss: 0.7860 - accuracy: 0.7575 - val_loss: 0.8039 - val_accuracy: 0.7491
Epoch 8/30
67/67 [=====] - 4s 56ms/step - loss: 0.7465 - accuracy: 0.7645 - val_loss: 0.8097 - val_accuracy: 0.7622
```

Epoch 9/30
67/67 [=====] - 4s 55ms/step - loss: 0.6999 - accuracy: 0.7894 - val_loss: 0.7522 - val_accuracy: 0.7903
Epoch 10/30
67/67 [=====] - 4s 55ms/step - loss: 0.6331 - accuracy: 0.8058 - val_loss: 0.7700 - val_accuracy: 0.7659
Epoch 11/30
67/67 [=====] - 4s 57ms/step - loss: 0.6061 - accuracy: 0.8096 - val_loss: 0.7201 - val_accuracy: 0.7903
Epoch 12/30
67/67 [=====] - 4s 56ms/step - loss: 0.5934 - accuracy: 0.8218 - val_loss: 0.7288 - val_accuracy: 0.7790
Epoch 13/30
67/67 [=====] - 4s 56ms/step - loss: 0.5737 - accuracy: 0.8143 - val_loss: 0.7249 - val_accuracy: 0.7846
Epoch 14/30
67/67 [=====] - 4s 57ms/step - loss: 0.5679 - accuracy: 0.8283 - val_loss: 0.7185 - val_accuracy: 0.8165
Epoch 15/30
67/67 [=====] - 4s 56ms/step - loss: 0.5391 - accuracy: 0.8377 - val_loss: 0.7621 - val_accuracy: 0.7753
Epoch 16/30
67/67 [=====] - 4s 56ms/step - loss: 0.4932 - accuracy: 0.8569 - val_loss: 0.6647 - val_accuracy: 0.8109
Epoch 17/30
67/67 [=====] - 4s 55ms/step - loss: 0.4939 - accuracy: 0.8532 - val_loss: 0.7199 - val_accuracy: 0.7865
Epoch 18/30
67/67 [=====] - 4s 56ms/step - loss: 0.4639 - accuracy: 0.8635 - val_loss: 0.6299 - val_accuracy: 0.8165
Epoch 19/30
67/67 [=====] - 4s 55ms/step - loss: 0.4964 - accuracy: 0.8565 - val_loss: 0.7134 - val_accuracy: 0.7790
Epoch 20/30
67/67 [=====] - 4s 55ms/step - loss: 0.4502 - accuracy: 0.8705 - val_loss: 0.6631 - val_accuracy: 0.8146
Epoch 21/30
67/67 [=====] - 4s 57ms/step - loss: 0.4663 - accuracy: 0.8574 - val_loss: 0.6665 - val_accuracy: 0.8071
Epoch 22/30
67/67 [=====] - 4s 56ms/step - loss: 0.4237 - accuracy: 0.8790 - val_loss: 0.6562 - val_accuracy: 0.8333
Epoch 23/30
67/67 [=====] - 4s 56ms/step - loss: 0.4132 - accuracy: 0.8827 - val_loss: 0.6605 - val_accuracy: 0.8277
Epoch 24/30
67/67 [=====] - 4s 56ms/step - loss: 0.4272 - accuracy: 0.8757 - val_loss: 0.7410 - val_accuracy: 0.7921

```
Epoch 25/30
67/67 [=====] - 4s 56ms/step - loss: 0.3938 - accuracy: 0.8921 - val_loss: 0.6622 - val_accuracy: 0.8333
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.4040 - accuracy: 0.8907 - val_loss: 0.7017 - val_accuracy: 0.7978
Epoch 27/30
67/67 [=====] - 4s 57ms/step - loss: 0.3640 - accuracy: 0.9099 - val_loss: 0.6751 - val_accuracy: 0.8202
Epoch 28/30
67/67 [=====] - 4s 56ms/step - loss: 0.3614 - accuracy: 0.9071 - val_loss: 0.6160 - val_accuracy: 0.8408
Epoch 29/30
67/67 [=====] - 4s 56ms/step - loss: 0.3739 - accuracy: 0.9006 - val_loss: 0.6300 - val_accuracy: 0.8371
Epoch 30/30
67/67 [=====] - 4s 56ms/step - loss: 0.3993 - accuracy: 0.8921 - val_loss: 0.7082 - val_accuracy: 0.8221
157/157 [=====] - 2s 13ms/step - loss: 0.6990 - accuracy: 0.8155
Test accuracy: 0.816
```

```
Fold 2/5, (params=l2_param: 0.001, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
```

```
2023-03-17 12:37:10.722467: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
67/67 [=====] - ETA: 0s - loss: 1.9336 - accuracy: 0.4158
```

```
2023-03-17 12:37:16.016441: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
67/67 [=====] - 6s 77ms/step - loss: 1.9336 - accuracy: 0.4158 - val_loss: 1.8823 - val_accuracy: 0.3902
```

```
Epoch 2/30
```

```
67/67 [=====] - 4s 61ms/step - loss: 1.4092 - accuracy: 0.5659 - val_loss: 1.5406 - val_accuracy: 0.4165
```

```
Epoch 3/30
```

```
67/67 [=====] - 4s 57ms/step - loss: 1.1790 - accuracy: 0.6240 - val_loss: 1.1195 - val_accuracy: 0.6360
```

```
Epoch 4/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 1.0628 - accuracy: 0.6414 - val_loss: 1.0485 - val_accuracy: 0.6473
Epoch 5/30
67/67 [=====] - 4s 56ms/step - loss: 0.9665 - accuracy: 0.6615 - val_loss: 1.0879 - val_accuracy: 0.5328
Epoch 6/30
67/67 [=====] - 4s 56ms/step - loss: 0.8684 - accuracy: 0.6976 - val_loss: 0.9175 - val_accuracy: 0.6998
Epoch 7/30
67/67 [=====] - 4s 56ms/step - loss: 0.7863 - accuracy: 0.7361 - val_loss: 0.8926 - val_accuracy: 0.6979
Epoch 8/30
67/67 [=====] - 4s 56ms/step - loss: 0.7545 - accuracy: 0.7595 - val_loss: 0.8169 - val_accuracy: 0.7598
Epoch 9/30
67/67 [=====] - 4s 56ms/step - loss: 0.6489 - accuracy: 0.7876 - val_loss: 0.7397 - val_accuracy: 0.7880
Epoch 10/30
67/67 [=====] - 4s 55ms/step - loss: 0.6252 - accuracy: 0.7965 - val_loss: 0.6610 - val_accuracy: 0.8068
Epoch 11/30
67/67 [=====] - 4s 55ms/step - loss: 0.5809 - accuracy: 0.8237 - val_loss: 0.8026 - val_accuracy: 0.7636
Epoch 12/30
67/67 [=====] - 4s 56ms/step - loss: 0.5509 - accuracy: 0.8322 - val_loss: 0.8995 - val_accuracy: 0.7261
Epoch 13/30
67/67 [=====] - 4s 56ms/step - loss: 0.5083 - accuracy: 0.8523 - val_loss: 0.6137 - val_accuracy: 0.8255
Epoch 14/30
67/67 [=====] - 4s 56ms/step - loss: 0.5142 - accuracy: 0.8434 - val_loss: 0.6159 - val_accuracy: 0.8274
Epoch 15/30
67/67 [=====] - 4s 56ms/step - loss: 0.4962 - accuracy: 0.8486 - val_loss: 0.7153 - val_accuracy: 0.7636
Epoch 16/30
67/67 [=====] - 4s 56ms/step - loss: 0.4881 - accuracy: 0.8575 - val_loss: 0.6653 - val_accuracy: 0.7936
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.4694 - accuracy: 0.8617 - val_loss: 0.6008 - val_accuracy: 0.8293
Epoch 18/30
67/67 [=====] - 4s 56ms/step - loss: 0.4625 - accuracy: 0.8579 - val_loss: 0.6569 - val_accuracy: 0.7955
Epoch 19/30
67/67 [=====] - 4s 56ms/step - loss: 0.4298 - accuracy: 0.8823 - val_loss: 0.8099 - val_accuracy: 0.7692
Epoch 20/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 0.4561 - accuracy: 0.8711 - val_loss: 0.8112 - val_accuracy: 0.7824
Epoch 21/30
67/67 [=====] - 4s 56ms/step - loss: 0.4190 - accuracy: 0.8790 - val_loss: 0.5870 - val_accuracy: 0.8330
Epoch 22/30
67/67 [=====] - 4s 56ms/step - loss: 0.3873 - accuracy: 0.8955 - val_loss: 0.6391 - val_accuracy: 0.8405
Epoch 23/30
67/67 [=====] - 4s 57ms/step - loss: 0.3714 - accuracy: 0.8945 - val_loss: 0.6277 - val_accuracy: 0.8330
Epoch 24/30
67/67 [=====] - 4s 55ms/step - loss: 0.3634 - accuracy: 0.9062 - val_loss: 0.9849 - val_accuracy: 0.7505
Epoch 25/30
67/67 [=====] - 4s 55ms/step - loss: 0.3755 - accuracy: 0.9067 - val_loss: 0.6610 - val_accuracy: 0.8293
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.3857 - accuracy: 0.9030 - val_loss: 0.6457 - val_accuracy: 0.8161
Epoch 27/30
67/67 [=====] - 4s 56ms/step - loss: 0.3471 - accuracy: 0.9147 - val_loss: 0.6004 - val_accuracy: 0.8555
Epoch 28/30
67/67 [=====] - 4s 56ms/step - loss: 0.3532 - accuracy: 0.9165 - val_loss: 0.6122 - val_accuracy: 0.8330
Epoch 29/30
67/67 [=====] - 4s 56ms/step - loss: 0.3212 - accuracy: 0.9278 - val_loss: 0.6420 - val_accuracy: 0.8330
Epoch 30/30
67/67 [=====] - 4s 56ms/step - loss: 0.3357 - accuracy: 0.9231 - val_loss: 0.6199 - val_accuracy: 0.8424
157/157 [=====] - 2s 14ms/step - loss: 0.6399 - accuracy: 0.8355
Test accuracy: 0.836
```

```
Fold 3/5, (params=l2_param: 0.001, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
```

```
2023-03-17 12:39:08.933497: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
67/67 [=====] - ETA: 0s - loss: 2.0535 - accuracy:
```

0.3619

2023-03-17 12:39:14.482312: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - 7s 83ms/step - loss: 2.0535 - accuracy:
0.3619 - val_loss: 1.6272 - val_accuracy: 0.4897

Epoch 2/30

67/67 [=====] - 4s 58ms/step - loss: 1.4798 - accuracy:
0.5509 - val_loss: 1.5980 - val_accuracy: 0.4634

Epoch 3/30

67/67 [=====] - 4s 57ms/step - loss: 1.2203 - accuracy:
0.6067 - val_loss: 1.1709 - val_accuracy: 0.5966

Epoch 4/30

67/67 [=====] - 4s 56ms/step - loss: 1.0557 - accuracy:
0.6578 - val_loss: 1.1456 - val_accuracy: 0.5797

Epoch 5/30

67/67 [=====] - 4s 57ms/step - loss: 0.9689 - accuracy:
0.6906 - val_loss: 1.0111 - val_accuracy: 0.6304

Epoch 6/30

67/67 [=====] - 4s 56ms/step - loss: 0.8976 - accuracy:
0.7135 - val_loss: 0.9161 - val_accuracy: 0.6886

Epoch 7/30

67/67 [=====] - 4s 56ms/step - loss: 0.7599 - accuracy:
0.7642 - val_loss: 0.8270 - val_accuracy: 0.7486

Epoch 8/30

67/67 [=====] - 4s 56ms/step - loss: 0.7085 - accuracy:
0.7876 - val_loss: 0.8046 - val_accuracy: 0.7580

Epoch 9/30

67/67 [=====] - 4s 56ms/step - loss: 0.6853 - accuracy:
0.7951 - val_loss: 0.7741 - val_accuracy: 0.7786

Epoch 10/30

67/67 [=====] - 4s 56ms/step - loss: 0.6357 - accuracy:
0.8214 - val_loss: 0.6988 - val_accuracy: 0.7917

Epoch 11/30

67/67 [=====] - 4s 55ms/step - loss: 0.5825 - accuracy:
0.8218 - val_loss: 0.7207 - val_accuracy: 0.7749

Epoch 12/30

67/67 [=====] - 4s 55ms/step - loss: 0.5766 - accuracy:
0.8293 - val_loss: 0.6883 - val_accuracy: 0.7786

Epoch 13/30

67/67 [=====] - 4s 55ms/step - loss: 0.5359 - accuracy:
0.8345 - val_loss: 0.7098 - val_accuracy: 0.7824

Epoch 14/30

67/67 [=====] - 4s 56ms/step - loss: 0.5310 - accuracy:
0.8401 - val_loss: 0.6900 - val_accuracy: 0.7655

Epoch 15/30

67/67 [=====] - 4s 55ms/step - loss: 0.4962 - accuracy:

```
0.8626 - val_loss: 0.6442 - val_accuracy: 0.8011
Epoch 16/30
67/67 [=====] - 4s 55ms/step - loss: 0.4615 - accuracy:
0.8664 - val_loss: 0.6788 - val_accuracy: 0.7824
Epoch 17/30
67/67 [=====] - 4s 55ms/step - loss: 0.4667 - accuracy:
0.8706 - val_loss: 0.7168 - val_accuracy: 0.7674
Epoch 18/30
67/67 [=====] - 4s 56ms/step - loss: 0.4465 - accuracy:
0.8753 - val_loss: 0.6636 - val_accuracy: 0.8143
Epoch 19/30
67/67 [=====] - 4s 55ms/step - loss: 0.4506 - accuracy:
0.8767 - val_loss: 0.6571 - val_accuracy: 0.7955
Epoch 20/30
67/67 [=====] - 4s 56ms/step - loss: 0.4042 - accuracy:
0.8837 - val_loss: 0.6786 - val_accuracy: 0.7936
Epoch 21/30
67/67 [=====] - 4s 55ms/step - loss: 0.4088 - accuracy:
0.8894 - val_loss: 0.7174 - val_accuracy: 0.7674
Epoch 22/30
67/67 [=====] - 4s 55ms/step - loss: 0.4155 - accuracy:
0.8926 - val_loss: 0.6395 - val_accuracy: 0.8068
Epoch 23/30
67/67 [=====] - 4s 55ms/step - loss: 0.3748 - accuracy:
0.9086 - val_loss: 0.6774 - val_accuracy: 0.8293
Epoch 24/30
67/67 [=====] - 4s 56ms/step - loss: 0.3833 - accuracy:
0.8997 - val_loss: 0.7693 - val_accuracy: 0.7899
Epoch 25/30
67/67 [=====] - 4s 55ms/step - loss: 0.3814 - accuracy:
0.9001 - val_loss: 0.6234 - val_accuracy: 0.8330
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.4118 - accuracy:
0.8926 - val_loss: 0.6500 - val_accuracy: 0.8161
Epoch 27/30
67/67 [=====] - 4s 56ms/step - loss: 0.3744 - accuracy:
0.9015 - val_loss: 0.6243 - val_accuracy: 0.8386
Epoch 28/30
67/67 [=====] - 4s 55ms/step - loss: 0.3530 - accuracy:
0.9142 - val_loss: 0.6894 - val_accuracy: 0.8255
Epoch 29/30
67/67 [=====] - 4s 55ms/step - loss: 0.3580 - accuracy:
0.9142 - val_loss: 0.6871 - val_accuracy: 0.8161
Epoch 30/30
67/67 [=====] - 4s 56ms/step - loss: 0.3421 - accuracy:
0.9208 - val_loss: 0.7036 - val_accuracy: 0.8143
157/157 [=====] - 2s 13ms/step - loss: 0.6729 -
accuracy: 0.8231
```

Test accuracy: 0.823

Fold 4/5, (params=l2_param: 0.001, augment: False)

Found 2666 files belonging to 5 classes.

Using 2133 files for training.

Found 2666 files belonging to 5 classes.

Using 533 files for validation.

Found 2499 files belonging to 5 classes.

Epoch 1/30

2023-03-17 12:41:06.907564: I

tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]

Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - ETA: 0s - loss: 2.0652 - accuracy: 0.3601

2023-03-17 12:41:12.313319: I

tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]

Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - 6s 81ms/step - loss: 2.0652 - accuracy: 0.3601 - val_loss: 1.7632 - val_accuracy: 0.5103

Epoch 2/30

67/67 [=====] - 4s 62ms/step - loss: 1.4500 - accuracy: 0.5762 - val_loss: 1.3145 - val_accuracy: 0.6098

Epoch 3/30

67/67 [=====] - 4s 59ms/step - loss: 1.1973 - accuracy: 0.6273 - val_loss: 1.2328 - val_accuracy: 0.5929

Epoch 4/30

67/67 [=====] - 4s 59ms/step - loss: 1.0619 - accuracy: 0.6690 - val_loss: 1.1512 - val_accuracy: 0.6266

Epoch 5/30

67/67 [=====] - 4s 58ms/step - loss: 0.9415 - accuracy: 0.7107 - val_loss: 0.9014 - val_accuracy: 0.7336

Epoch 6/30

67/67 [=====] - 4s 58ms/step - loss: 0.8415 - accuracy: 0.7553 - val_loss: 0.8930 - val_accuracy: 0.7298

Epoch 7/30

67/67 [=====] - 4s 60ms/step - loss: 0.7797 - accuracy: 0.7609 - val_loss: 0.8134 - val_accuracy: 0.7561

Epoch 8/30

67/67 [=====] - 4s 58ms/step - loss: 0.7236 - accuracy: 0.7834 - val_loss: 0.7647 - val_accuracy: 0.7598

Epoch 9/30

67/67 [=====] - 4s 58ms/step - loss: 0.6675 - accuracy: 0.8012 - val_loss: 0.7153 - val_accuracy: 0.7899

Epoch 10/30

67/67 [=====] - 4s 58ms/step - loss: 0.6445 - accuracy: 0.8022 - val_loss: 0.7506 - val_accuracy: 0.7561

```
Epoch 11/30
67/67 [=====] - 4s 58ms/step - loss: 0.6360 - accuracy: 0.8003 - val_loss: 0.7187 - val_accuracy: 0.7955
Epoch 12/30
67/67 [=====] - 4s 58ms/step - loss: 0.5841 - accuracy: 0.8200 - val_loss: 0.6918 - val_accuracy: 0.7861
Epoch 13/30
67/67 [=====] - 4s 58ms/step - loss: 0.5738 - accuracy: 0.8336 - val_loss: 0.6984 - val_accuracy: 0.8068
Epoch 14/30
67/67 [=====] - 4s 58ms/step - loss: 0.5418 - accuracy: 0.8279 - val_loss: 0.6726 - val_accuracy: 0.7936
Epoch 15/30
67/67 [=====] - 4s 58ms/step - loss: 0.5266 - accuracy: 0.8434 - val_loss: 0.6465 - val_accuracy: 0.8161
Epoch 16/30
67/67 [=====] - 4s 59ms/step - loss: 0.5270 - accuracy: 0.8490 - val_loss: 0.6518 - val_accuracy: 0.7992
Epoch 17/30
67/67 [=====] - 4s 59ms/step - loss: 0.5066 - accuracy: 0.8387 - val_loss: 0.5726 - val_accuracy: 0.8161
Epoch 18/30
67/67 [=====] - 4s 58ms/step - loss: 0.5289 - accuracy: 0.8429 - val_loss: 0.6421 - val_accuracy: 0.8349
Epoch 19/30
67/67 [=====] - 4s 59ms/step - loss: 0.5083 - accuracy: 0.8462 - val_loss: 0.6621 - val_accuracy: 0.8161
Epoch 20/30
67/67 [=====] - 4s 58ms/step - loss: 0.4361 - accuracy: 0.8776 - val_loss: 0.6277 - val_accuracy: 0.8086
Epoch 21/30
67/67 [=====] - 4s 58ms/step - loss: 0.4771 - accuracy: 0.8579 - val_loss: 0.6081 - val_accuracy: 0.8274
Epoch 22/30
67/67 [=====] - 4s 60ms/step - loss: 0.4375 - accuracy: 0.8697 - val_loss: 0.6879 - val_accuracy: 0.7767
Epoch 23/30
67/67 [=====] - 4s 58ms/step - loss: 0.4258 - accuracy: 0.8805 - val_loss: 0.6952 - val_accuracy: 0.8011
Epoch 24/30
67/67 [=====] - 4s 58ms/step - loss: 0.4084 - accuracy: 0.8926 - val_loss: 0.6679 - val_accuracy: 0.8105
Epoch 25/30
67/67 [=====] - 4s 58ms/step - loss: 0.3881 - accuracy: 0.8940 - val_loss: 0.6624 - val_accuracy: 0.8199
Epoch 26/30
67/67 [=====] - 4s 58ms/step - loss: 0.3938 - accuracy: 0.8837 - val_loss: 0.5513 - val_accuracy: 0.8349
```

```
Epoch 27/30
67/67 [=====] - 4s 58ms/step - loss: 0.3610 - accuracy: 0.9067 - val_loss: 0.8183 - val_accuracy: 0.7880
Epoch 28/30
67/67 [=====] - 4s 58ms/step - loss: 0.4122 - accuracy: 0.8875 - val_loss: 0.6000 - val_accuracy: 0.8405
Epoch 29/30
67/67 [=====] - 4s 58ms/step - loss: 0.3519 - accuracy: 0.9151 - val_loss: 0.6660 - val_accuracy: 0.8255
Epoch 30/30
67/67 [=====] - 4s 58ms/step - loss: 0.3463 - accuracy: 0.9048 - val_loss: 0.6641 - val_accuracy: 0.8105
157/157 [=====] - 2s 14ms/step - loss: 0.6981 - accuracy: 0.8043
Test accuracy: 0.804

Fold 5/5, (params=12_param: 0.001, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 12:43:10.737854: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
67/67 [=====] - ETA: 0s - loss: 2.0031 - accuracy: 0.3877

2023-03-17 12:43:16.031341: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
67/67 [=====] - 7s 79ms/step - loss: 2.0031 - accuracy: 0.3877 - val_loss: 1.7858 - val_accuracy: 0.4015
Epoch 2/30
67/67 [=====] - 4s 58ms/step - loss: 1.4182 - accuracy: 0.5565 - val_loss: 1.3718 - val_accuracy: 0.5178
Epoch 3/30
67/67 [=====] - 4s 59ms/step - loss: 1.2115 - accuracy: 0.6128 - val_loss: 1.1915 - val_accuracy: 0.5760
Epoch 4/30
67/67 [=====] - 4s 57ms/step - loss: 1.0337 - accuracy: 0.6620 - val_loss: 1.0080 - val_accuracy: 0.6435
Epoch 5/30
67/67 [=====] - 4s 56ms/step - loss: 0.9407 - accuracy: 0.6887 - val_loss: 0.9608 - val_accuracy: 0.6341
Epoch 6/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 0.8739 - accuracy: 0.6971 - val_loss: 0.9752 - val_accuracy: 0.6529
Epoch 7/30
67/67 [=====] - 4s 56ms/step - loss: 0.7741 - accuracy: 0.7600 - val_loss: 0.9096 - val_accuracy: 0.7111
Epoch 8/30
67/67 [=====] - 4s 55ms/step - loss: 0.6941 - accuracy: 0.7947 - val_loss: 0.7579 - val_accuracy: 0.7523
Epoch 9/30
67/67 [=====] - 4s 56ms/step - loss: 0.6434 - accuracy: 0.8050 - val_loss: 0.7443 - val_accuracy: 0.7448
Epoch 10/30
67/67 [=====] - 4s 56ms/step - loss: 0.6146 - accuracy: 0.8092 - val_loss: 0.6728 - val_accuracy: 0.7711
Epoch 11/30
67/67 [=====] - 4s 56ms/step - loss: 0.5881 - accuracy: 0.8256 - val_loss: 0.6190 - val_accuracy: 0.8161
Epoch 12/30
67/67 [=====] - 4s 55ms/step - loss: 0.5763 - accuracy: 0.8317 - val_loss: 0.7129 - val_accuracy: 0.7711
Epoch 13/30
67/67 [=====] - 4s 56ms/step - loss: 0.5410 - accuracy: 0.8481 - val_loss: 0.6051 - val_accuracy: 0.8011
Epoch 14/30
67/67 [=====] - 4s 56ms/step - loss: 0.5361 - accuracy: 0.8429 - val_loss: 0.6168 - val_accuracy: 0.8011
Epoch 15/30
67/67 [=====] - 4s 56ms/step - loss: 0.5209 - accuracy: 0.8556 - val_loss: 0.6380 - val_accuracy: 0.8086
Epoch 16/30
67/67 [=====] - 4s 56ms/step - loss: 0.4768 - accuracy: 0.8654 - val_loss: 0.6614 - val_accuracy: 0.7842
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.4521 - accuracy: 0.8711 - val_loss: 0.6038 - val_accuracy: 0.8143
Epoch 18/30
67/67 [=====] - 4s 55ms/step - loss: 0.4549 - accuracy: 0.8608 - val_loss: 0.6057 - val_accuracy: 0.8086
Epoch 19/30
67/67 [=====] - 4s 55ms/step - loss: 0.4422 - accuracy: 0.8814 - val_loss: 0.5777 - val_accuracy: 0.8180
Epoch 20/30
67/67 [=====] - 4s 55ms/step - loss: 0.4255 - accuracy: 0.8837 - val_loss: 0.5609 - val_accuracy: 0.8311
Epoch 21/30
67/67 [=====] - 4s 55ms/step - loss: 0.4094 - accuracy: 0.8870 - val_loss: 0.5721 - val_accuracy: 0.8274
Epoch 22/30
```

```
67/67 [=====] - 4s 55ms/step - loss: 0.4008 - accuracy: 0.8898 - val_loss: 0.6195 - val_accuracy: 0.8086
Epoch 23/30
67/67 [=====] - 4s 55ms/step - loss: 0.4119 - accuracy: 0.8884 - val_loss: 0.6134 - val_accuracy: 0.8068
Epoch 24/30
67/67 [=====] - 4s 55ms/step - loss: 0.3591 - accuracy: 0.9072 - val_loss: 0.6423 - val_accuracy: 0.8143
Epoch 25/30
67/67 [=====] - 4s 56ms/step - loss: 0.3906 - accuracy: 0.8908 - val_loss: 0.6098 - val_accuracy: 0.8218
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.3326 - accuracy: 0.9250 - val_loss: 0.6802 - val_accuracy: 0.8293
Epoch 27/30
67/67 [=====] - 4s 55ms/step - loss: 0.3702 - accuracy: 0.8983 - val_loss: 0.6003 - val_accuracy: 0.8293
Epoch 28/30
67/67 [=====] - 4s 56ms/step - loss: 0.3114 - accuracy: 0.9283 - val_loss: 0.6762 - val_accuracy: 0.8368
Epoch 29/30
67/67 [=====] - 4s 56ms/step - loss: 0.3520 - accuracy: 0.9161 - val_loss: 0.6129 - val_accuracy: 0.8368
Epoch 30/30
67/67 [=====] - 4s 56ms/step - loss: 0.3190 - accuracy: 0.9287 - val_loss: 0.6767 - val_accuracy: 0.8199
157/157 [=====] - 2s 13ms/step - loss: 0.6470 - accuracy: 0.8211
Test accuracy: 0.821
```

```
CURRENTLY TRAINING THE MODEL WITH THE FOLLOWING PARAMS: l2_param: 0.0001, augment: False
Fold 1/5, (params=l2_param: 0.0001, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
```

```
2023-03-17 12:45:08.723115: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
67/67 [=====] - ETA: 0s - loss: 1.5730 - accuracy: 0.3626
```

```
2023-03-17 12:45:14.966040: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
```

Plugin optimizer for device_type GPU is enabled.

```
67/67 [=====] - 7s 91ms/step - loss: 1.5730 - accuracy: 0.3626 - val_loss: 1.3814 - val_accuracy: 0.4307
Epoch 2/30
67/67 [=====] - 4s 61ms/step - loss: 1.0996 - accuracy: 0.5699 - val_loss: 1.2133 - val_accuracy: 0.4438
Epoch 3/30
67/67 [=====] - 4s 59ms/step - loss: 0.9944 - accuracy: 0.6191 - val_loss: 1.0394 - val_accuracy: 0.5899
Epoch 4/30
67/67 [=====] - 4s 57ms/step - loss: 0.9265 - accuracy: 0.6431 - val_loss: 0.9133 - val_accuracy: 0.6255
Epoch 5/30
67/67 [=====] - 4s 57ms/step - loss: 0.8249 - accuracy: 0.6890 - val_loss: 1.1216 - val_accuracy: 0.5150
Epoch 6/30
67/67 [=====] - 4s 57ms/step - loss: 0.7611 - accuracy: 0.7219 - val_loss: 0.7880 - val_accuracy: 0.7004
Epoch 7/30
67/67 [=====] - 4s 57ms/step - loss: 0.7167 - accuracy: 0.7430 - val_loss: 0.8542 - val_accuracy: 0.6835
Epoch 8/30
67/67 [=====] - 4s 56ms/step - loss: 0.6627 - accuracy: 0.7603 - val_loss: 0.6854 - val_accuracy: 0.7846
Epoch 9/30
67/67 [=====] - 4s 56ms/step - loss: 0.5749 - accuracy: 0.8096 - val_loss: 0.7473 - val_accuracy: 0.7360
Epoch 10/30
67/67 [=====] - 4s 56ms/step - loss: 0.5549 - accuracy: 0.8138 - val_loss: 0.6836 - val_accuracy: 0.7828
Epoch 11/30
67/67 [=====] - 4s 57ms/step - loss: 0.5113 - accuracy: 0.8368 - val_loss: 0.6459 - val_accuracy: 0.7959
Epoch 12/30
67/67 [=====] - 4s 57ms/step - loss: 0.4830 - accuracy: 0.8386 - val_loss: 0.5926 - val_accuracy: 0.8071
Epoch 13/30
67/67 [=====] - 4s 57ms/step - loss: 0.4359 - accuracy: 0.8551 - val_loss: 0.5742 - val_accuracy: 0.8258
Epoch 14/30
67/67 [=====] - 4s 58ms/step - loss: 0.4168 - accuracy: 0.8673 - val_loss: 0.6666 - val_accuracy: 0.8015
Epoch 15/30
67/67 [=====] - 4s 55ms/step - loss: 0.3911 - accuracy: 0.8748 - val_loss: 0.5495 - val_accuracy: 0.8333
Epoch 16/30
67/67 [=====] - 4s 56ms/step - loss: 0.3903 - accuracy:
```

```
0.8724 - val_loss: 0.6373 - val_accuracy: 0.8052
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.3571 - accuracy:
0.8931 - val_loss: 0.5579 - val_accuracy: 0.8352
Epoch 18/30
67/67 [=====] - 4s 56ms/step - loss: 0.3687 - accuracy:
0.8795 - val_loss: 0.6688 - val_accuracy: 0.7884
Epoch 19/30
67/67 [=====] - 4s 56ms/step - loss: 0.3112 - accuracy:
0.9099 - val_loss: 0.6069 - val_accuracy: 0.8071
Epoch 20/30
67/67 [=====] - 4s 57ms/step - loss: 0.2871 - accuracy:
0.9114 - val_loss: 0.5525 - val_accuracy: 0.8221
Epoch 21/30
67/67 [=====] - 4s 56ms/step - loss: 0.3096 - accuracy:
0.9048 - val_loss: 0.6176 - val_accuracy: 0.8202
Epoch 22/30
67/67 [=====] - 4s 55ms/step - loss: 0.3209 - accuracy:
0.8987 - val_loss: 0.6694 - val_accuracy: 0.8165
Epoch 23/30
67/67 [=====] - 4s 56ms/step - loss: 0.2667 - accuracy:
0.9254 - val_loss: 0.6493 - val_accuracy: 0.8071
Epoch 24/30
67/67 [=====] - 4s 55ms/step - loss: 0.2644 - accuracy:
0.9306 - val_loss: 0.6461 - val_accuracy: 0.8071
Epoch 25/30
67/67 [=====] - 4s 56ms/step - loss: 0.2334 - accuracy:
0.9414 - val_loss: 0.6570 - val_accuracy: 0.8258
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.2526 - accuracy:
0.9325 - val_loss: 0.6563 - val_accuracy: 0.8277
Epoch 27/30
67/67 [=====] - 4s 65ms/step - loss: 0.2434 - accuracy:
0.9339 - val_loss: 0.7488 - val_accuracy: 0.7978
Epoch 28/30
67/67 [=====] - 4s 59ms/step - loss: 0.2377 - accuracy:
0.9418 - val_loss: 0.6618 - val_accuracy: 0.8277
Epoch 29/30
67/67 [=====] - 4s 57ms/step - loss: 0.2300 - accuracy:
0.9414 - val_loss: 0.6724 - val_accuracy: 0.8296
Epoch 30/30
67/67 [=====] - 4s 57ms/step - loss: 0.2435 - accuracy:
0.9395 - val_loss: 0.6533 - val_accuracy: 0.8333
157/157 [=====] - 2s 14ms/step - loss: 0.7173 -
accuracy: 0.8059
Test accuracy: 0.806
```

Fold 2/5, (params=l2_param: 0.0001, augment: False)

```
Found 2666 files belonging to 5 classes.  
Using 2133 files for training.  
Found 2666 files belonging to 5 classes.  
Using 533 files for validation.  
Found 2499 files belonging to 5 classes.  
Epoch 1/30  
  
2023-03-17 12:47:09.951077: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.  
  
67/67 [=====] - ETA: 0s - loss: 1.5481 - accuracy:  
0.3526  
  
2023-03-17 12:47:16.183735: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.  
  
67/67 [=====] - 7s 94ms/step - loss: 1.5481 - accuracy:  
0.3526 - val_loss: 1.2381 - val_accuracy: 0.5197  
Epoch 2/30  
67/67 [=====] - 4s 58ms/step - loss: 1.0672 - accuracy:  
0.5692 - val_loss: 1.1017 - val_accuracy: 0.5891  
Epoch 3/30  
67/67 [=====] - 4s 57ms/step - loss: 0.9311 - accuracy:  
0.6507 - val_loss: 0.9716 - val_accuracy: 0.5966  
Epoch 4/30  
67/67 [=====] - 4s 58ms/step - loss: 0.8588 - accuracy:  
0.6742 - val_loss: 0.8897 - val_accuracy: 0.6323  
Epoch 5/30  
67/67 [=====] - 4s 56ms/step - loss: 0.7542 - accuracy:  
0.7323 - val_loss: 0.8801 - val_accuracy: 0.6548  
Epoch 6/30  
67/67 [=====] - 4s 57ms/step - loss: 0.7084 - accuracy:  
0.7473 - val_loss: 0.9267 - val_accuracy: 0.6379  
Epoch 7/30  
67/67 [=====] - 4s 56ms/step - loss: 0.6370 - accuracy:  
0.7881 - val_loss: 0.7856 - val_accuracy: 0.7355  
Epoch 8/30  
67/67 [=====] - 4s 57ms/step - loss: 0.5916 - accuracy:  
0.7979 - val_loss: 0.7103 - val_accuracy: 0.7617  
Epoch 9/30  
67/67 [=====] - 4s 56ms/step - loss: 0.5634 - accuracy:  
0.8143 - val_loss: 0.6877 - val_accuracy: 0.7936  
Epoch 10/30  
67/67 [=====] - 4s 56ms/step - loss: 0.5539 - accuracy:  
0.8181 - val_loss: 0.6721 - val_accuracy: 0.7861  
Epoch 11/30  
67/67 [=====] - 4s 56ms/step - loss: 0.4908 - accuracy:  
0.8429 - val_loss: 0.6468 - val_accuracy: 0.7936
```

Epoch 12/30
67/67 [=====] - 4s 56ms/step - loss: 0.4512 - accuracy: 0.8537 - val_loss: 0.5622 - val_accuracy: 0.8349
Epoch 13/30
67/67 [=====] - 4s 56ms/step - loss: 0.4186 - accuracy: 0.8570 - val_loss: 0.6321 - val_accuracy: 0.8049
Epoch 14/30
67/67 [=====] - 4s 56ms/step - loss: 0.4087 - accuracy: 0.8758 - val_loss: 0.5514 - val_accuracy: 0.8386
Epoch 15/30
67/67 [=====] - 4s 56ms/step - loss: 0.3868 - accuracy: 0.8800 - val_loss: 0.5706 - val_accuracy: 0.8386
Epoch 16/30
67/67 [=====] - 4s 58ms/step - loss: 0.4149 - accuracy: 0.8659 - val_loss: 0.5613 - val_accuracy: 0.8349
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.3629 - accuracy: 0.8814 - val_loss: 0.5982 - val_accuracy: 0.8161
Epoch 18/30
67/67 [=====] - 4s 55ms/step - loss: 0.3217 - accuracy: 0.8973 - val_loss: 0.6089 - val_accuracy: 0.8218
Epoch 19/30
67/67 [=====] - 4s 55ms/step - loss: 0.2983 - accuracy: 0.9114 - val_loss: 0.7011 - val_accuracy: 0.8180
Epoch 20/30
67/67 [=====] - 4s 56ms/step - loss: 0.3162 - accuracy: 0.9048 - val_loss: 0.6723 - val_accuracy: 0.8143
Epoch 21/30
67/67 [=====] - 4s 55ms/step - loss: 0.3173 - accuracy: 0.8987 - val_loss: 0.6740 - val_accuracy: 0.8293
Epoch 22/30
67/67 [=====] - 4s 55ms/step - loss: 0.2747 - accuracy: 0.9212 - val_loss: 0.7909 - val_accuracy: 0.7786
Epoch 23/30
67/67 [=====] - 4s 56ms/step - loss: 0.2646 - accuracy: 0.9212 - val_loss: 0.7566 - val_accuracy: 0.7880
Epoch 24/30
67/67 [=====] - 4s 55ms/step - loss: 0.2812 - accuracy: 0.9222 - val_loss: 0.6480 - val_accuracy: 0.8368
Epoch 25/30
67/67 [=====] - 4s 56ms/step - loss: 0.2801 - accuracy: 0.9161 - val_loss: 0.7031 - val_accuracy: 0.8368
Epoch 26/30
67/67 [=====] - 4s 55ms/step - loss: 0.2534 - accuracy: 0.9301 - val_loss: 0.6177 - val_accuracy: 0.8443
Epoch 27/30
67/67 [=====] - 4s 55ms/step - loss: 0.2260 - accuracy: 0.9428 - val_loss: 0.7341 - val_accuracy: 0.8161

```
Epoch 28/30
67/67 [=====] - 4s 55ms/step - loss: 0.2096 - accuracy: 0.9451 - val_loss: 0.7083 - val_accuracy: 0.8386
Epoch 29/30
67/67 [=====] - 4s 56ms/step - loss: 0.2064 - accuracy: 0.9526 - val_loss: 0.8358 - val_accuracy: 0.8236
Epoch 30/30
67/67 [=====] - 4s 56ms/step - loss: 0.2106 - accuracy: 0.9512 - val_loss: 0.8049 - val_accuracy: 0.8236
157/157 [=====] - 2s 14ms/step - loss: 0.7644 - accuracy: 0.8127
Test accuracy: 0.813

Fold 3/5, (params=l2_param: 0.0001, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 12:49:10.784757: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - ETA: 0s - loss: 1.5362 - accuracy: 0.3830

2023-03-17 12:49:16.921557: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - 7s 87ms/step - loss: 1.5362 - accuracy: 0.3830 - val_loss: 1.3791 - val_accuracy: 0.4146
Epoch 2/30
67/67 [=====] - 4s 60ms/step - loss: 1.1000 - accuracy: 0.5663 - val_loss: 1.1001 - val_accuracy: 0.5797
Epoch 3/30
67/67 [=====] - 4s 59ms/step - loss: 0.9550 - accuracy: 0.6338 - val_loss: 0.9597 - val_accuracy: 0.6398
Epoch 4/30
67/67 [=====] - 4s 57ms/step - loss: 0.8535 - accuracy: 0.6718 - val_loss: 0.9326 - val_accuracy: 0.6829
Epoch 5/30
67/67 [=====] - 4s 57ms/step - loss: 0.7380 - accuracy: 0.7398 - val_loss: 0.7604 - val_accuracy: 0.7148
Epoch 6/30
67/67 [=====] - 4s 56ms/step - loss: 0.7041 - accuracy: 0.7557 - val_loss: 0.7059 - val_accuracy: 0.7636
Epoch 7/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 0.6094 - accuracy: 0.7993 - val_loss: 0.7579 - val_accuracy: 0.7111
Epoch 8/30
67/67 [=====] - 4s 57ms/step - loss: 0.5900 - accuracy: 0.8054 - val_loss: 0.6915 - val_accuracy: 0.7730
Epoch 9/30
67/67 [=====] - 4s 57ms/step - loss: 0.5454 - accuracy: 0.8190 - val_loss: 0.8294 - val_accuracy: 0.7223
Epoch 10/30
67/67 [=====] - 4s 56ms/step - loss: 0.4981 - accuracy: 0.8340 - val_loss: 0.6702 - val_accuracy: 0.7936
Epoch 11/30
67/67 [=====] - 4s 56ms/step - loss: 0.4859 - accuracy: 0.8425 - val_loss: 0.7108 - val_accuracy: 0.7674
Epoch 12/30
67/67 [=====] - 4s 56ms/step - loss: 0.4448 - accuracy: 0.8584 - val_loss: 0.7226 - val_accuracy: 0.7467
Epoch 13/30
67/67 [=====] - 4s 56ms/step - loss: 0.4302 - accuracy: 0.8626 - val_loss: 0.7035 - val_accuracy: 0.7636
Epoch 14/30
67/67 [=====] - 4s 56ms/step - loss: 0.4359 - accuracy: 0.8598 - val_loss: 0.8340 - val_accuracy: 0.7561
Epoch 15/30
67/67 [=====] - 4s 56ms/step - loss: 0.3878 - accuracy: 0.8758 - val_loss: 0.5261 - val_accuracy: 0.8255
Epoch 16/30
67/67 [=====] - 4s 57ms/step - loss: 0.3500 - accuracy: 0.8880 - val_loss: 0.6209 - val_accuracy: 0.7842
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.3455 - accuracy: 0.8931 - val_loss: 0.7103 - val_accuracy: 0.8011
Epoch 18/30
67/67 [=====] - 4s 56ms/step - loss: 0.2925 - accuracy: 0.9189 - val_loss: 0.6762 - val_accuracy: 0.8086
Epoch 19/30
67/67 [=====] - 4s 57ms/step - loss: 0.3321 - accuracy: 0.8950 - val_loss: 0.7815 - val_accuracy: 0.8068
Epoch 20/30
67/67 [=====] - 4s 55ms/step - loss: 0.3145 - accuracy: 0.9034 - val_loss: 0.8987 - val_accuracy: 0.7767
Epoch 21/30
67/67 [=====] - 4s 59ms/step - loss: 0.3017 - accuracy: 0.9048 - val_loss: 0.6980 - val_accuracy: 0.8030
Epoch 22/30
67/67 [=====] - 4s 57ms/step - loss: 0.2634 - accuracy: 0.9208 - val_loss: 0.7572 - val_accuracy: 0.7786
Epoch 23/30
```

```
67/67 [=====] - 4s 56ms/step - loss: 0.2289 - accuracy: 0.9348 - val_loss: 0.8588 - val_accuracy: 0.8049
Epoch 24/30
67/67 [=====] - 4s 56ms/step - loss: 0.2897 - accuracy: 0.9184 - val_loss: 0.7614 - val_accuracy: 0.7955
Epoch 25/30
67/67 [=====] - 4s 55ms/step - loss: 0.2532 - accuracy: 0.9330 - val_loss: 0.5970 - val_accuracy: 0.8405
Epoch 26/30
67/67 [=====] - 4s 56ms/step - loss: 0.2287 - accuracy: 0.9419 - val_loss: 0.7429 - val_accuracy: 0.8011
Epoch 27/30
67/67 [=====] - 4s 60ms/step - loss: 0.2001 - accuracy: 0.9522 - val_loss: 0.7491 - val_accuracy: 0.8011
Epoch 28/30
67/67 [=====] - 5s 68ms/step - loss: 0.1995 - accuracy: 0.9512 - val_loss: 0.7295 - val_accuracy: 0.8105
Epoch 29/30
67/67 [=====] - 4s 63ms/step - loss: 0.2555 - accuracy: 0.9409 - val_loss: 0.7010 - val_accuracy: 0.8161
Epoch 30/30
67/67 [=====] - 4s 57ms/step - loss: 0.2147 - accuracy: 0.9480 - val_loss: 0.8804 - val_accuracy: 0.7767
157/157 [=====] - 2s 14ms/step - loss: 0.7050 - accuracy: 0.8259
Test accuracy: 0.826
```

```
Fold 4/5, (params=l2_param: 0.0001, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
```

```
2023-03-17 12:51:12.893499: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
67/67 [=====] - ETA: 0s - loss: 1.4455 - accuracy: 0.4126
```

```
2023-03-17 12:51:19.283659: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
67/67 [=====] - 8s 91ms/step - loss: 1.4455 - accuracy: 0.4126 - val_loss: 1.1102 - val_accuracy: 0.5610
Epoch 2/30
67/67 [=====] - 4s 61ms/step - loss: 1.0841 - accuracy:
```

```
0.5631 - val_loss: 1.0961 - val_accuracy: 0.5535
Epoch 3/30
67/67 [=====] - 4s 59ms/step - loss: 0.9789 - accuracy:
0.6179 - val_loss: 1.0494 - val_accuracy: 0.5272
Epoch 4/30
67/67 [=====] - 4s 60ms/step - loss: 0.8646 - accuracy:
0.6648 - val_loss: 0.9450 - val_accuracy: 0.6191
Epoch 5/30
67/67 [=====] - 4s 57ms/step - loss: 0.8083 - accuracy:
0.7079 - val_loss: 0.7855 - val_accuracy: 0.7261
Epoch 6/30
67/67 [=====] - 4s 55ms/step - loss: 0.7613 - accuracy:
0.7393 - val_loss: 0.8453 - val_accuracy: 0.6848
Epoch 7/30
67/67 [=====] - 4s 56ms/step - loss: 0.6698 - accuracy:
0.7787 - val_loss: 0.6657 - val_accuracy: 0.7730
Epoch 8/30
67/67 [=====] - 4s 55ms/step - loss: 0.6232 - accuracy:
0.7937 - val_loss: 0.6648 - val_accuracy: 0.7749
Epoch 9/30
67/67 [=====] - 4s 56ms/step - loss: 0.5625 - accuracy:
0.8045 - val_loss: 0.6914 - val_accuracy: 0.7542
Epoch 10/30
67/67 [=====] - 4s 56ms/step - loss: 0.5485 - accuracy:
0.8261 - val_loss: 0.6468 - val_accuracy: 0.8030
Epoch 11/30
67/67 [=====] - 4s 55ms/step - loss: 0.5291 - accuracy:
0.8289 - val_loss: 0.6591 - val_accuracy: 0.7711
Epoch 12/30
67/67 [=====] - 4s 56ms/step - loss: 0.4707 - accuracy:
0.8495 - val_loss: 0.6082 - val_accuracy: 0.7936
Epoch 13/30
67/67 [=====] - 4s 55ms/step - loss: 0.4378 - accuracy:
0.8579 - val_loss: 0.6585 - val_accuracy: 0.7824
Epoch 14/30
67/67 [=====] - 4s 55ms/step - loss: 0.4173 - accuracy:
0.8706 - val_loss: 0.5410 - val_accuracy: 0.8218
Epoch 15/30
67/67 [=====] - 4s 57ms/step - loss: 0.4004 - accuracy:
0.8790 - val_loss: 0.4982 - val_accuracy: 0.8311
Epoch 16/30
67/67 [=====] - 4s 55ms/step - loss: 0.3751 - accuracy:
0.8767 - val_loss: 0.6141 - val_accuracy: 0.7899
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.3875 - accuracy:
0.8842 - val_loss: 0.6508 - val_accuracy: 0.7842
Epoch 18/30
67/67 [=====] - 4s 55ms/step - loss: 0.3432 - accuracy:
```

```
0.8940 - val_loss: 0.5668 - val_accuracy: 0.8199
Epoch 19/30
67/67 [=====] - 4s 55ms/step - loss: 0.2990 - accuracy:
0.9100 - val_loss: 0.5216 - val_accuracy: 0.8386
Epoch 20/30
67/67 [=====] - 4s 55ms/step - loss: 0.2995 - accuracy:
0.9142 - val_loss: 0.5342 - val_accuracy: 0.8499
Epoch 21/30
67/67 [=====] - 4s 59ms/step - loss: 0.3050 - accuracy:
0.9067 - val_loss: 0.6152 - val_accuracy: 0.8161
Epoch 22/30
67/67 [=====] - 4s 56ms/step - loss: 0.2785 - accuracy:
0.9245 - val_loss: 0.6983 - val_accuracy: 0.7899
Epoch 23/30
67/67 [=====] - 4s 55ms/step - loss: 0.2498 - accuracy:
0.9348 - val_loss: 0.6482 - val_accuracy: 0.8124
Epoch 24/30
67/67 [=====] - 4s 55ms/step - loss: 0.2353 - accuracy:
0.9367 - val_loss: 0.5583 - val_accuracy: 0.8480
Epoch 25/30
67/67 [=====] - 4s 56ms/step - loss: 0.2573 - accuracy:
0.9278 - val_loss: 0.6219 - val_accuracy: 0.8218
Epoch 26/30
67/67 [=====] - 4s 55ms/step - loss: 0.2326 - accuracy:
0.9376 - val_loss: 0.6001 - val_accuracy: 0.8274
Epoch 27/30
67/67 [=====] - 4s 57ms/step - loss: 0.2290 - accuracy:
0.9447 - val_loss: 0.6712 - val_accuracy: 0.7992
Epoch 28/30
67/67 [=====] - 4s 57ms/step - loss: 0.1887 - accuracy:
0.9597 - val_loss: 0.6148 - val_accuracy: 0.8424
Epoch 29/30
67/67 [=====] - 4s 57ms/step - loss: 0.2001 - accuracy:
0.9508 - val_loss: 0.6754 - val_accuracy: 0.8462
Epoch 30/30
67/67 [=====] - 4s 55ms/step - loss: 0.2381 - accuracy:
0.9423 - val_loss: 0.7007 - val_accuracy: 0.7917
157/157 [=====] - 2s 14ms/step - loss: 0.7698 -
accuracy: 0.7967
Test accuracy: 0.797
```

```
Fold 5/5, (params=l2_param: 0.0001, augment: False)
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
```

```
2023-03-17 12:53:13.032039: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - ETA: 0s - loss: 1.6421 - accuracy:
0.3301

2023-03-17 12:53:18.986531: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

67/67 [=====] - 7s 87ms/step - loss: 1.6421 - accuracy:
0.3301 - val_loss: 1.3768 - val_accuracy: 0.4653
Epoch 2/30
67/67 [=====] - 4s 61ms/step - loss: 1.0746 - accuracy:
0.5804 - val_loss: 1.0755 - val_accuracy: 0.5891
Epoch 3/30
67/67 [=====] - 4s 57ms/step - loss: 0.9444 - accuracy:
0.6343 - val_loss: 1.0247 - val_accuracy: 0.6116
Epoch 4/30
67/67 [=====] - 4s 57ms/step - loss: 0.8539 - accuracy:
0.6845 - val_loss: 0.8503 - val_accuracy: 0.7298
Epoch 5/30
67/67 [=====] - 4s 56ms/step - loss: 0.7522 - accuracy:
0.7267 - val_loss: 0.7998 - val_accuracy: 0.7317
Epoch 6/30
67/67 [=====] - 4s 56ms/step - loss: 0.7284 - accuracy:
0.7492 - val_loss: 1.0769 - val_accuracy: 0.6079
Epoch 7/30
67/67 [=====] - 4s 56ms/step - loss: 0.6781 - accuracy:
0.7609 - val_loss: 0.7382 - val_accuracy: 0.7730
Epoch 8/30
67/67 [=====] - 4s 56ms/step - loss: 0.6061 - accuracy:
0.7979 - val_loss: 0.7153 - val_accuracy: 0.7711
Epoch 9/30
67/67 [=====] - 4s 55ms/step - loss: 0.5613 - accuracy:
0.8186 - val_loss: 0.7045 - val_accuracy: 0.7767
Epoch 10/30
67/67 [=====] - 4s 55ms/step - loss: 0.5177 - accuracy:
0.8308 - val_loss: 0.6873 - val_accuracy: 0.7842
Epoch 11/30
67/67 [=====] - 4s 55ms/step - loss: 0.4980 - accuracy:
0.8359 - val_loss: 0.6142 - val_accuracy: 0.8161
Epoch 12/30
67/67 [=====] - 4s 56ms/step - loss: 0.4635 - accuracy:
0.8509 - val_loss: 0.8132 - val_accuracy: 0.7448
Epoch 13/30
67/67 [=====] - 4s 55ms/step - loss: 0.4432 - accuracy:
0.8561 - val_loss: 0.6860 - val_accuracy: 0.7655
```

```
Epoch 14/30
67/67 [=====] - 4s 55ms/step - loss: 0.4026 - accuracy: 0.8729 - val_loss: 0.5816 - val_accuracy: 0.8311
Epoch 15/30
67/67 [=====] - 4s 55ms/step - loss: 0.4268 - accuracy: 0.8640 - val_loss: 0.6108 - val_accuracy: 0.8068
Epoch 16/30
67/67 [=====] - 4s 55ms/step - loss: 0.3587 - accuracy: 0.8833 - val_loss: 0.7276 - val_accuracy: 0.7992
Epoch 17/30
67/67 [=====] - 4s 56ms/step - loss: 0.3939 - accuracy: 0.8687 - val_loss: 0.5608 - val_accuracy: 0.8274
Epoch 18/30
67/67 [=====] - 4s 55ms/step - loss: 0.3158 - accuracy: 0.8978 - val_loss: 0.6326 - val_accuracy: 0.8293
Epoch 19/30
67/67 [=====] - 4s 56ms/step - loss: 0.2954 - accuracy: 0.9039 - val_loss: 0.6170 - val_accuracy: 0.8255
Epoch 20/30
67/67 [=====] - 4s 55ms/step - loss: 0.2903 - accuracy: 0.9105 - val_loss: 0.6603 - val_accuracy: 0.7992
Epoch 21/30
67/67 [=====] - 4s 55ms/step - loss: 0.2792 - accuracy: 0.9147 - val_loss: 0.6412 - val_accuracy: 0.8236
Epoch 22/30
67/67 [=====] - 4s 55ms/step - loss: 0.2597 - accuracy: 0.9264 - val_loss: 0.6620 - val_accuracy: 0.8143
Epoch 23/30
67/67 [=====] - 4s 55ms/step - loss: 0.2494 - accuracy: 0.9348 - val_loss: 0.7287 - val_accuracy: 0.8068
Epoch 24/30
67/67 [=====] - 4s 55ms/step - loss: 0.2648 - accuracy: 0.9231 - val_loss: 0.6689 - val_accuracy: 0.8218
Epoch 25/30
67/67 [=====] - 4s 55ms/step - loss: 0.2377 - accuracy: 0.9339 - val_loss: 0.6906 - val_accuracy: 0.8368
Epoch 26/30
67/67 [=====] - 4s 55ms/step - loss: 0.2153 - accuracy: 0.9451 - val_loss: 0.7946 - val_accuracy: 0.8011
Epoch 27/30
67/67 [=====] - 4s 55ms/step - loss: 0.2585 - accuracy: 0.9306 - val_loss: 0.8594 - val_accuracy: 0.7749
Epoch 28/30
67/67 [=====] - 4s 55ms/step - loss: 0.2273 - accuracy: 0.9391 - val_loss: 0.6566 - val_accuracy: 0.8499
Epoch 29/30
67/67 [=====] - 4s 56ms/step - loss: 0.1753 - accuracy: 0.9597 - val_loss: 0.6903 - val_accuracy: 0.8405
```

```

Epoch 30/30
67/67 [=====] - 4s 55ms/step - loss: 0.2018 - accuracy:
0.9437 - val_loss: 0.7379 - val_accuracy: 0.8293
157/157 [=====] - 2s 14ms/step - loss: 0.7187 -
accuracy: 0.8275
Test accuracy: 0.828

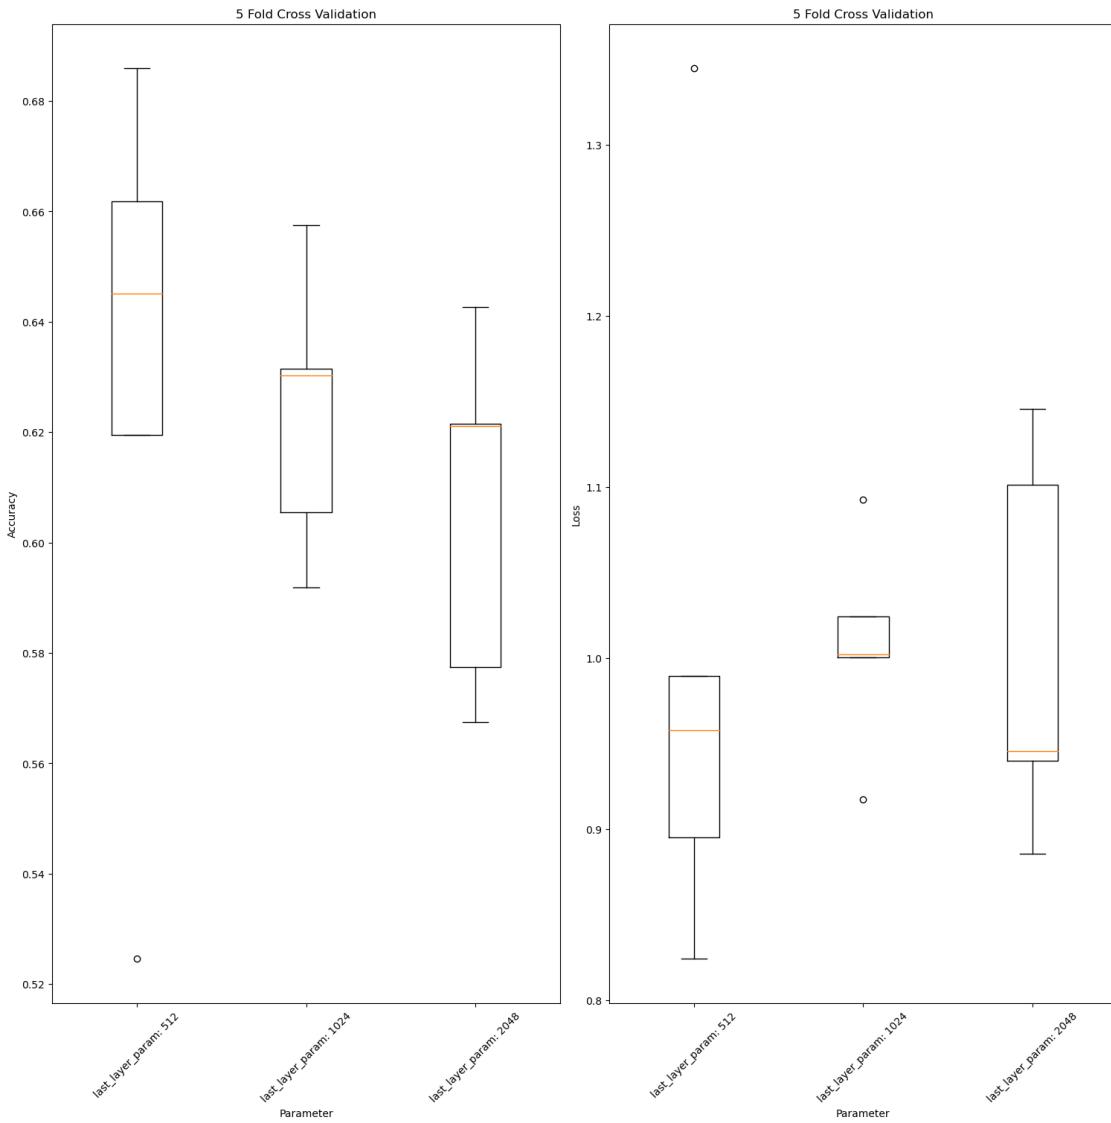
```

5.5.2 Visualizaton of the Parameter Performance

```
[4]: # The following data has been generated with the above custom KFold Cross Validation method
accuracies = dict()
losses = dict()
for key in model_history.keys():
    accuracies[key] = []
    losses[key] = []
    for entry in model_history[key]:
        losses[key].append(entry["loss"])
        accuracies[key].append(entry["acc"])

fig, ax = plt.subplots(1, 2, constrained_layout=True)
ax[0].set_title("5 Fold Cross Validation")
ax[0].set_ylabel("Accuracy")
ax[0].set_xlabel("Parameter")
ax[0].boxplot(accuracies.values())
ax[0].set_xticklabels(accuracies.keys(), rotation=45)
ax[1].set_title("5 Fold Cross Validation")
ax[1].set_ylabel("Loss")
ax[1].set_xlabel("Parameter")
ax[1].boxplot(losses.values())
ax[1].set_xticklabels(losses.keys(), rotation=45)

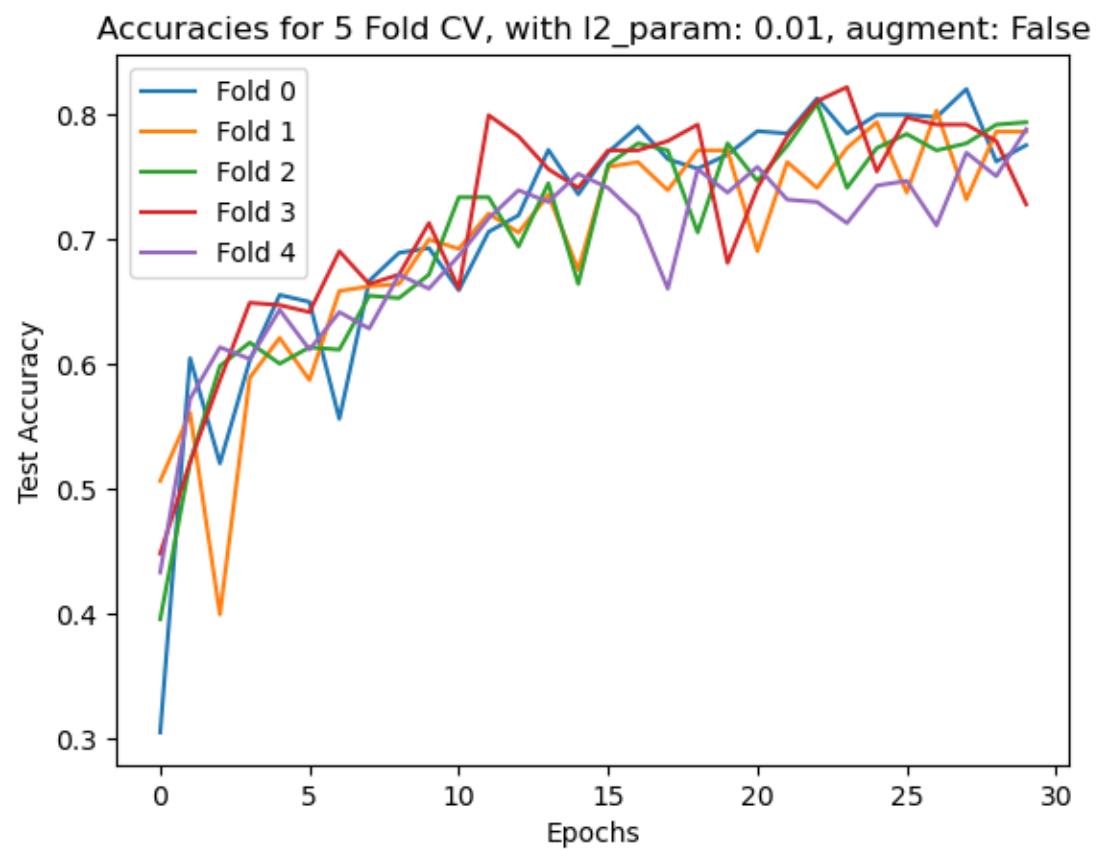
fig.set_figheight(15)
fig.set_figwidth(15)
```



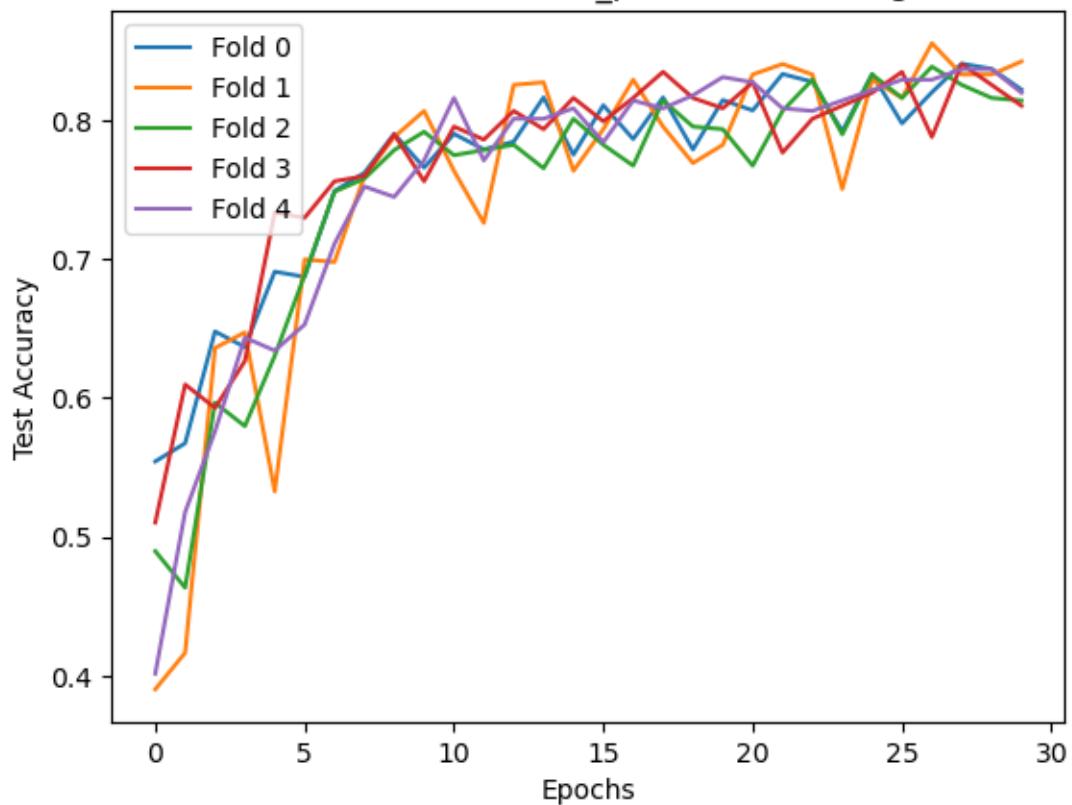
Visualization of the Test Accuracy for each Fold

```
[27]: for m_key in model_history:
    for item in model_history.get(m_key):
        for key in item:
            if key == 'history':
                history = item.get('history')
                fold = item.get("fold")
                plt.plot(history.history['val_accuracy'], label=f"Fold {fold}")
plt.title(f"Accuracies for {folds} Fold CV, with {m_key}")
plt.ylabel("Test Accuracy")
plt.xlabel("Epochs")
plt.legend(loc= 'upper left')
```

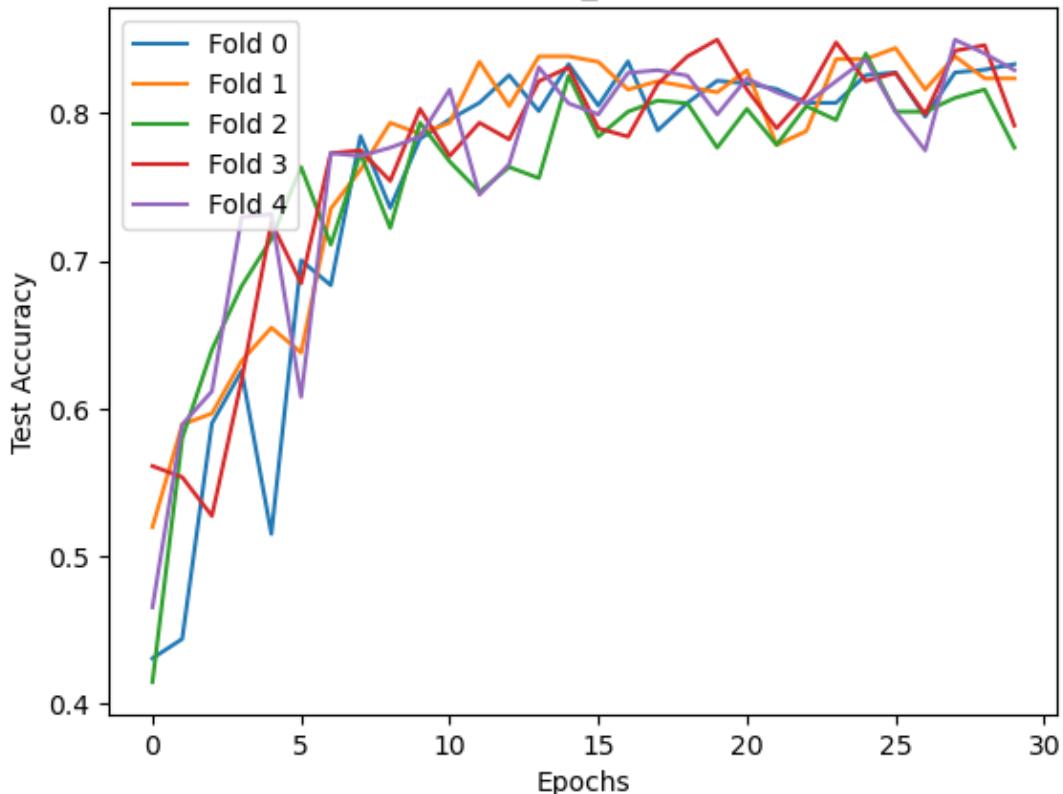
```
plt.show()
```



Accuracies for 5 Fold CV, with l2_param: 0.001, augment: False



Accuracies for 5 Fold CV, with l2_param: 0.0001, augment: False



Visualization of the Training Convergence

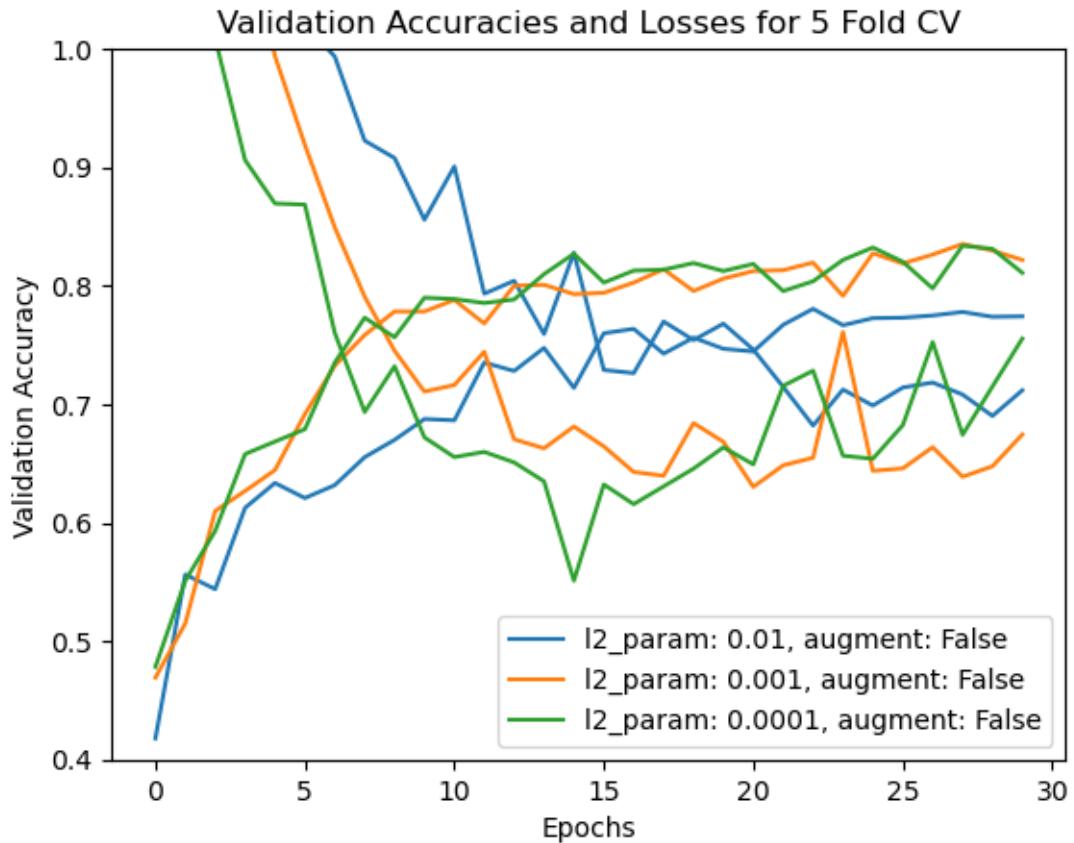
```
[28]: for m_key in model_history:
    accuracies = []
    losses = []
    for item in model_history.get(m_key):
        for key in item:
            if key == 'history':
                history = item.get('history')
                fold = item.get("fold")
                accuracies.append(history.history['val_accuracy'])
                losses.append(history.history['val_loss'])
    avg_acc = np.average(np.array(accuracies), axis=0)
    avg_loss = np.average(np.array(losses), axis=0)
    line = plt.plot(avg_acc, label=m_key)
    plt.plot(avg_loss, c=line[0].get_color())

plt.ylim(0.4, 1)
plt.title(f"Validation Accuracies and Losses for {folds} Fold CV")
plt.ylabel("Validation Accuracy")
```

```

plt.xlabel("Epochs")
plt.legend(loc= 'lower right')
plt.show()

```



For simplification, the above lines for accuracy and loss use the same color and can be differentiated over the course of multiple epochs such as the loss comes from the top and gets lower while the accuracy comes from the bottom and rises higher.

5.5.3 Update with GridSearchCV

Apparently, the first belief that the model parameters can't be optimized using GridSearchCV, was a misperception, because `model.get_params()` would not output the expected parameters... Beause the GridSearchCV implementation does somehow not work with our model, and because of the huge effort, we left the above implementation, but still provide a version using GridSearchCV, which is way more elegant. We believe that it is a bug of GridSearchCV that our model doesn't work with the function, because the accuracy does not improve at all, and if we use less layers, it seems to work. For the details, please look at the [Issue on GitHub](#) which we've created for this behavior.

To still use GridSearchCV, we've tried another implementation that uses `tf.data.Dataset` directly: [TensorCross](#). Our model works with this implementation, so we provide an example using this

method as well. At first, TensorCross was not compatible with tensorflow-macos, which we've also created a [GitHub Issue](#) for. That has been fixed very quickly so we were able to use the library.

```
[29]: from tensorcross.model_selection import GridSearchCV

param_grid=dict(
    # last_layer_param=[256, 512, 1024],
    l2_param=[0.01, 0.001, 0.0001],
    augment=[False],
    # last_dropout_param=[0.2, 0.3, 0.4]
)
grid_search_cv = GridSearchCV(
    model_fn=create_model_optimized,
    param_grid=param_grid,
    n_folds=5,
)

grid_search_cv.fit(
    dataset=full_train_dataset,
    epochs=30,
)
```

```
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30

2023-03-17 12:55:12.132427: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - ETA: 0s - loss: 5.0508 - accuracy: 0.3500

2023-03-17 12:55:18.835056: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - 8s 95ms/step - loss: 5.0508 - accuracy: 0.3500 - val_loss: 3.0621 - val_accuracy: 0.4180
Epoch 2/30
68/68 [=====] - 4s 58ms/step - loss: 2.1852 - accuracy: 0.5418 - val_loss: 1.7785 - val_accuracy: 0.6016
Epoch 3/30
68/68 [=====] - 4s 57ms/step - loss: 1.5145 - accuracy: 0.5952 - val_loss: 1.4351 - val_accuracy: 0.5547
Epoch 4/30
68/68 [=====] - 4s 56ms/step - loss: 1.2410 - accuracy:
```

```
0.6137 - val_loss: 1.2639 - val_accuracy: 0.5371
Epoch 5/30
68/68 [=====] - 4s 58ms/step - loss: 1.1287 - accuracy:
0.6188 - val_loss: 1.1877 - val_accuracy: 0.5430
Epoch 6/30
68/68 [=====] - 4s 57ms/step - loss: 1.0392 - accuracy:
0.6258 - val_loss: 1.0841 - val_accuracy: 0.6426
Epoch 7/30
68/68 [=====] - 4s 56ms/step - loss: 0.9872 - accuracy:
0.6235 - val_loss: 1.0027 - val_accuracy: 0.6250
Epoch 8/30
68/68 [=====] - 4s 57ms/step - loss: 0.9173 - accuracy:
0.6537 - val_loss: 0.9311 - val_accuracy: 0.6699
Epoch 9/30
68/68 [=====] - 4s 56ms/step - loss: 0.8909 - accuracy:
0.6727 - val_loss: 0.9047 - val_accuracy: 0.6758
Epoch 10/30
68/68 [=====] - 4s 57ms/step - loss: 0.8392 - accuracy:
0.6908 - val_loss: 0.8275 - val_accuracy: 0.6953
Epoch 11/30
68/68 [=====] - 4s 56ms/step - loss: 0.8490 - accuracy:
0.6713 - val_loss: 0.8802 - val_accuracy: 0.6934
Epoch 12/30
68/68 [=====] - 4s 56ms/step - loss: 0.8039 - accuracy:
0.6996 - val_loss: 0.8335 - val_accuracy: 0.7090
Epoch 13/30
68/68 [=====] - 4s 56ms/step - loss: 0.7820 - accuracy:
0.7173 - val_loss: 0.8101 - val_accuracy: 0.7500
Epoch 14/30
68/68 [=====] - 4s 56ms/step - loss: 0.7816 - accuracy:
0.7261 - val_loss: 0.8549 - val_accuracy: 0.7070
Epoch 15/30
68/68 [=====] - 4s 56ms/step - loss: 0.7282 - accuracy:
0.7553 - val_loss: 0.7177 - val_accuracy: 0.7637
Epoch 16/30
68/68 [=====] - 4s 57ms/step - loss: 0.6809 - accuracy:
0.7758 - val_loss: 0.7033 - val_accuracy: 0.7578
Epoch 17/30
68/68 [=====] - 4s 56ms/step - loss: 0.7385 - accuracy:
0.7549 - val_loss: 0.6834 - val_accuracy: 0.7949
Epoch 18/30
68/68 [=====] - 4s 57ms/step - loss: 0.6774 - accuracy:
0.7772 - val_loss: 0.7698 - val_accuracy: 0.7383
Epoch 19/30
68/68 [=====] - 4s 56ms/step - loss: 0.6427 - accuracy:
0.7939 - val_loss: 0.6977 - val_accuracy: 0.7793
Epoch 20/30
68/68 [=====] - 4s 56ms/step - loss: 0.6772 - accuracy:
```

```
0.7679 - val_loss: 0.7114 - val_accuracy: 0.7754
Epoch 21/30
68/68 [=====] - 4s 56ms/step - loss: 0.6223 - accuracy: 0.8013 - val_loss: 0.6183 - val_accuracy: 0.8105
Epoch 22/30
68/68 [=====] - 4s 56ms/step - loss: 0.6399 - accuracy: 0.7902 - val_loss: 0.6553 - val_accuracy: 0.8047
Epoch 23/30
68/68 [=====] - 4s 56ms/step - loss: 0.6217 - accuracy: 0.8050 - val_loss: 0.6574 - val_accuracy: 0.7930
Epoch 24/30
68/68 [=====] - 4s 55ms/step - loss: 0.6185 - accuracy: 0.8018 - val_loss: 0.7642 - val_accuracy: 0.7656
Epoch 25/30
68/68 [=====] - 4s 56ms/step - loss: 0.5871 - accuracy: 0.8134 - val_loss: 0.6598 - val_accuracy: 0.7949
Epoch 26/30
68/68 [=====] - 4s 56ms/step - loss: 0.6208 - accuracy: 0.8041 - val_loss: 0.8551 - val_accuracy: 0.7188
Epoch 27/30
68/68 [=====] - 4s 56ms/step - loss: 0.6192 - accuracy: 0.7990 - val_loss: 0.7401 - val_accuracy: 0.7832
Epoch 28/30
68/68 [=====] - 4s 56ms/step - loss: 0.5816 - accuracy: 0.8236 - val_loss: 0.6758 - val_accuracy: 0.8066
Epoch 29/30
68/68 [=====] - 4s 56ms/step - loss: 0.5898 - accuracy: 0.8213 - val_loss: 0.6576 - val_accuracy: 0.7988
Epoch 30/30
68/68 [=====] - 4s 56ms/step - loss: 0.5581 - accuracy: 0.8222 - val_loss: 0.7715 - val_accuracy: 0.7715
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 12:57:12.984752: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - ETA: 0s - loss: 4.5995 - accuracy: 0.3463
2023-03-17 12:57:18.620695: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - 7s 82ms/step - loss: 4.5995 - accuracy:
```

```
0.3463 - val_loss: 2.2423 - val_accuracy: 0.5332
Epoch 2/30
68/68 [=====] - 4s 64ms/step - loss: 1.7280 - accuracy:
0.5334 - val_loss: 1.3332 - val_accuracy: 0.6035
Epoch 3/30
68/68 [=====] - 4s 61ms/step - loss: 1.1741 - accuracy:
0.5956 - val_loss: 1.2122 - val_accuracy: 0.6055
Epoch 4/30
68/68 [=====] - 4s 58ms/step - loss: 1.0559 - accuracy:
0.6114 - val_loss: 1.0860 - val_accuracy: 0.6230
Epoch 5/30
68/68 [=====] - 4s 58ms/step - loss: 1.0245 - accuracy:
0.5942 - val_loss: 1.0052 - val_accuracy: 0.6348
Epoch 6/30
68/68 [=====] - 4s 57ms/step - loss: 0.9463 - accuracy:
0.6481 - val_loss: 0.9520 - val_accuracy: 0.6621
Epoch 7/30
68/68 [=====] - 4s 60ms/step - loss: 0.9458 - accuracy:
0.6430 - val_loss: 0.9868 - val_accuracy: 0.6230
Epoch 8/30
68/68 [=====] - 4s 57ms/step - loss: 0.9351 - accuracy:
0.6541 - val_loss: 0.8961 - val_accuracy: 0.6816
Epoch 9/30
68/68 [=====] - 4s 58ms/step - loss: 0.8790 - accuracy:
0.6750 - val_loss: 0.8152 - val_accuracy: 0.7090
Epoch 10/30
68/68 [=====] - 4s 57ms/step - loss: 0.8380 - accuracy:
0.6792 - val_loss: 0.7793 - val_accuracy: 0.7305
Epoch 11/30
68/68 [=====] - 4s 57ms/step - loss: 0.8186 - accuracy:
0.6908 - val_loss: 0.8721 - val_accuracy: 0.7031
Epoch 12/30
68/68 [=====] - 4s 57ms/step - loss: 0.7689 - accuracy:
0.7247 - val_loss: 0.8584 - val_accuracy: 0.7090
Epoch 13/30
68/68 [=====] - 4s 58ms/step - loss: 0.7408 - accuracy:
0.7349 - val_loss: 0.8255 - val_accuracy: 0.7109
Epoch 14/30
68/68 [=====] - 4s 57ms/step - loss: 0.7106 - accuracy:
0.7577 - val_loss: 0.6971 - val_accuracy: 0.8027
Epoch 15/30
68/68 [=====] - 4s 58ms/step - loss: 0.6797 - accuracy:
0.7804 - val_loss: 0.8157 - val_accuracy: 0.7109
Epoch 16/30
68/68 [=====] - 4s 57ms/step - loss: 0.7002 - accuracy:
0.7693 - val_loss: 0.7623 - val_accuracy: 0.7344
Epoch 17/30
68/68 [=====] - 4s 57ms/step - loss: 0.6812 - accuracy:
```

```
0.7786 - val_loss: 0.6335 - val_accuracy: 0.7871
Epoch 18/30
68/68 [=====] - 4s 57ms/step - loss: 0.6585 - accuracy:
0.7860 - val_loss: 0.6783 - val_accuracy: 0.7773
Epoch 19/30
68/68 [=====] - 4s 57ms/step - loss: 0.6493 - accuracy:
0.7967 - val_loss: 0.8155 - val_accuracy: 0.7383
Epoch 20/30
68/68 [=====] - 4s 57ms/step - loss: 0.6582 - accuracy:
0.7855 - val_loss: 0.7138 - val_accuracy: 0.7539
Epoch 21/30
68/68 [=====] - 4s 57ms/step - loss: 0.6450 - accuracy:
0.7939 - val_loss: 0.6099 - val_accuracy: 0.8125
Epoch 22/30
68/68 [=====] - 4s 57ms/step - loss: 0.6445 - accuracy:
0.7874 - val_loss: 0.7321 - val_accuracy: 0.7656
Epoch 23/30
68/68 [=====] - 4s 57ms/step - loss: 0.6215 - accuracy:
0.8069 - val_loss: 0.6168 - val_accuracy: 0.8066
Epoch 24/30
68/68 [=====] - 4s 57ms/step - loss: 0.6175 - accuracy:
0.7999 - val_loss: 0.5775 - val_accuracy: 0.8125
Epoch 25/30
68/68 [=====] - 4s 57ms/step - loss: 0.6046 - accuracy:
0.8022 - val_loss: 0.6884 - val_accuracy: 0.7793
Epoch 26/30
68/68 [=====] - 4s 57ms/step - loss: 0.5947 - accuracy:
0.8055 - val_loss: 0.7192 - val_accuracy: 0.7676
Epoch 27/30
68/68 [=====] - 4s 57ms/step - loss: 0.5996 - accuracy:
0.8045 - val_loss: 0.5394 - val_accuracy: 0.8496
Epoch 28/30
68/68 [=====] - 4s 57ms/step - loss: 0.5813 - accuracy:
0.8152 - val_loss: 0.5756 - val_accuracy: 0.8262
Epoch 29/30
68/68 [=====] - 4s 57ms/step - loss: 0.5725 - accuracy:
0.8199 - val_loss: 0.6967 - val_accuracy: 0.7832
Epoch 30/30
68/68 [=====] - 4s 57ms/step - loss: 0.6114 - accuracy:
0.8032 - val_loss: 0.6976 - val_accuracy: 0.7754
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
```

2023-03-17 12:59:14.414288: I

```
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - ETA: 0s - loss: 4.4490 - accuracy: 0.3821

2023-03-17 12:59:20.530114: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - 7s 93ms/step - loss: 4.4490 - accuracy: 0.3821 - val_loss: 2.6353 - val_accuracy: 0.3262
Epoch 2/30
68/68 [=====] - 4s 61ms/step - loss: 1.7182 - accuracy: 0.5608 - val_loss: 1.5304 - val_accuracy: 0.5176
Epoch 3/30
68/68 [=====] - 4s 59ms/step - loss: 1.2087 - accuracy: 0.6142 - val_loss: 1.2343 - val_accuracy: 0.5430
Epoch 4/30
68/68 [=====] - 4s 59ms/step - loss: 1.0884 - accuracy: 0.6184 - val_loss: 1.1352 - val_accuracy: 0.5859
Epoch 5/30
68/68 [=====] - 4s 60ms/step - loss: 1.0202 - accuracy: 0.6258 - val_loss: 1.0920 - val_accuracy: 0.5645
Epoch 6/30
68/68 [=====] - 4s 59ms/step - loss: 0.9748 - accuracy: 0.6416 - val_loss: 0.9982 - val_accuracy: 0.6367
Epoch 7/30
68/68 [=====] - 4s 61ms/step - loss: 0.9391 - accuracy: 0.6430 - val_loss: 0.9792 - val_accuracy: 0.6289
Epoch 8/30
68/68 [=====] - 4s 58ms/step - loss: 0.9365 - accuracy: 0.6527 - val_loss: 1.0192 - val_accuracy: 0.5801
Epoch 9/30
68/68 [=====] - 4s 58ms/step - loss: 0.8903 - accuracy: 0.6606 - val_loss: 0.9601 - val_accuracy: 0.6191
Epoch 10/30
68/68 [=====] - 4s 58ms/step - loss: 0.8581 - accuracy: 0.6722 - val_loss: 0.9501 - val_accuracy: 0.6113
Epoch 11/30
68/68 [=====] - 4s 58ms/step - loss: 0.8234 - accuracy: 0.6917 - val_loss: 0.9428 - val_accuracy: 0.6289
Epoch 12/30
68/68 [=====] - 4s 58ms/step - loss: 0.8084 - accuracy: 0.7052 - val_loss: 0.8520 - val_accuracy: 0.7031
Epoch 13/30
68/68 [=====] - 4s 58ms/step - loss: 0.7728 - accuracy: 0.7219 - val_loss: 0.9408 - val_accuracy: 0.6602
Epoch 14/30
```

```
68/68 [=====] - 4s 58ms/step - loss: 0.7804 - accuracy: 0.7196 - val_loss: 0.7882 - val_accuracy: 0.7598
Epoch 15/30
68/68 [=====] - 4s 58ms/step - loss: 0.7009 - accuracy: 0.7632 - val_loss: 0.7058 - val_accuracy: 0.7598
Epoch 16/30
68/68 [=====] - 4s 58ms/step - loss: 0.7313 - accuracy: 0.7409 - val_loss: 0.7914 - val_accuracy: 0.7324
Epoch 17/30
68/68 [=====] - 4s 58ms/step - loss: 0.6943 - accuracy: 0.7651 - val_loss: 0.8957 - val_accuracy: 0.6934
Epoch 18/30
68/68 [=====] - 4s 58ms/step - loss: 0.6563 - accuracy: 0.7864 - val_loss: 0.6622 - val_accuracy: 0.7715
Epoch 19/30
68/68 [=====] - 4s 58ms/step - loss: 0.6637 - accuracy: 0.7920 - val_loss: 0.7106 - val_accuracy: 0.7461
Epoch 20/30
68/68 [=====] - 4s 58ms/step - loss: 0.6666 - accuracy: 0.7874 - val_loss: 0.7508 - val_accuracy: 0.7461
Epoch 21/30
68/68 [=====] - 4s 59ms/step - loss: 0.6468 - accuracy: 0.7869 - val_loss: 0.8044 - val_accuracy: 0.7246
Epoch 22/30
68/68 [=====] - 4s 58ms/step - loss: 0.6089 - accuracy: 0.8045 - val_loss: 0.9317 - val_accuracy: 0.6582
Epoch 23/30
68/68 [=====] - 4s 58ms/step - loss: 0.6355 - accuracy: 0.7939 - val_loss: 0.7328 - val_accuracy: 0.7383
Epoch 24/30
68/68 [=====] - 4s 58ms/step - loss: 0.5907 - accuracy: 0.8106 - val_loss: 0.8067 - val_accuracy: 0.7129
Epoch 25/30
68/68 [=====] - 4s 58ms/step - loss: 0.6155 - accuracy: 0.8008 - val_loss: 0.8439 - val_accuracy: 0.6973
Epoch 26/30
68/68 [=====] - 4s 58ms/step - loss: 0.6228 - accuracy: 0.7994 - val_loss: 0.7180 - val_accuracy: 0.7793
Epoch 27/30
68/68 [=====] - 4s 58ms/step - loss: 0.5959 - accuracy: 0.8022 - val_loss: 0.6695 - val_accuracy: 0.7871
Epoch 28/30
68/68 [=====] - 4s 58ms/step - loss: 0.6056 - accuracy: 0.7981 - val_loss: 0.7603 - val_accuracy: 0.7734
Epoch 29/30
68/68 [=====] - 4s 58ms/step - loss: 0.5898 - accuracy: 0.8110 - val_loss: 0.6560 - val_accuracy: 0.7812
Epoch 30/30
```

```
68/68 [=====] - 4s 58ms/step - loss: 0.5900 - accuracy: 0.8148 - val_loss: 0.8220 - val_accuracy: 0.7188
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30

2023-03-17 13:01:17.636828: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - ETA: 0s - loss: 5.2465 - accuracy: 0.3078

2023-03-17 13:01:24.370882: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - 8s 92ms/step - loss: 5.2465 - accuracy: 0.3078 - val_loss: 2.9188 - val_accuracy: 0.5215
Epoch 2/30

68/68 [=====] - 4s 64ms/step - loss: 2.3490 - accuracy: 0.5330 - val_loss: 1.8311 - val_accuracy: 0.5449
Epoch 3/30

68/68 [=====] - 4s 64ms/step - loss: 1.5826 - accuracy: 0.5794 - val_loss: 1.4519 - val_accuracy: 0.5547
Epoch 4/30

68/68 [=====] - 4s 59ms/step - loss: 1.2900 - accuracy: 0.6026 - val_loss: 1.1869 - val_accuracy: 0.5938
Epoch 5/30

68/68 [=====] - 4s 61ms/step - loss: 1.1343 - accuracy: 0.6031 - val_loss: 1.0528 - val_accuracy: 0.6328
Epoch 6/30

68/68 [=====] - 4s 59ms/step - loss: 1.0393 - accuracy: 0.6291 - val_loss: 0.9621 - val_accuracy: 0.6621
Epoch 7/30

68/68 [=====] - 4s 60ms/step - loss: 0.9817 - accuracy: 0.6481 - val_loss: 0.8398 - val_accuracy: 0.7227
Epoch 8/30

68/68 [=====] - 4s 60ms/step - loss: 0.9401 - accuracy: 0.6551 - val_loss: 0.9715 - val_accuracy: 0.6445
Epoch 9/30

68/68 [=====] - 4s 58ms/step - loss: 0.8772 - accuracy: 0.6760 - val_loss: 0.8278 - val_accuracy: 0.7031
Epoch 10/30

68/68 [=====] - 4s 58ms/step - loss: 0.8305 - accuracy: 0.6978 - val_loss: 0.8260 - val_accuracy: 0.6934
Epoch 11/30
```

```
68/68 [=====] - 4s 58ms/step - loss: 0.8100 - accuracy: 0.7075 - val_loss: 0.8819 - val_accuracy: 0.6875
Epoch 12/30
68/68 [=====] - 4s 60ms/step - loss: 0.8089 - accuracy: 0.6955 - val_loss: 0.8585 - val_accuracy: 0.6621
Epoch 13/30
68/68 [=====] - 4s 59ms/step - loss: 0.7638 - accuracy: 0.7242 - val_loss: 0.7502 - val_accuracy: 0.7383
Epoch 14/30
68/68 [=====] - 4s 58ms/step - loss: 0.7278 - accuracy: 0.7470 - val_loss: 0.7962 - val_accuracy: 0.7188
Epoch 15/30
68/68 [=====] - 4s 59ms/step - loss: 0.7265 - accuracy: 0.7521 - val_loss: 0.7017 - val_accuracy: 0.7598
Epoch 16/30
68/68 [=====] - 4s 59ms/step - loss: 0.7116 - accuracy: 0.7581 - val_loss: 0.7620 - val_accuracy: 0.7598
Epoch 17/30
68/68 [=====] - 4s 58ms/step - loss: 0.6825 - accuracy: 0.7753 - val_loss: 0.7372 - val_accuracy: 0.7461
Epoch 18/30
68/68 [=====] - 4s 59ms/step - loss: 0.6897 - accuracy: 0.7702 - val_loss: 0.6860 - val_accuracy: 0.7852
Epoch 19/30
68/68 [=====] - 4s 58ms/step - loss: 0.6786 - accuracy: 0.7795 - val_loss: 0.7638 - val_accuracy: 0.7266
Epoch 20/30
68/68 [=====] - 4s 58ms/step - loss: 0.6716 - accuracy: 0.7716 - val_loss: 0.8188 - val_accuracy: 0.7207
Epoch 21/30
68/68 [=====] - 4s 58ms/step - loss: 0.6274 - accuracy: 0.7925 - val_loss: 0.7959 - val_accuracy: 0.7461
Epoch 22/30
68/68 [=====] - 4s 59ms/step - loss: 0.6210 - accuracy: 0.7864 - val_loss: 0.6548 - val_accuracy: 0.7969
Epoch 23/30
68/68 [=====] - 4s 58ms/step - loss: 0.5884 - accuracy: 0.8069 - val_loss: 0.7900 - val_accuracy: 0.7578
Epoch 24/30
68/68 [=====] - 4s 58ms/step - loss: 0.6215 - accuracy: 0.8008 - val_loss: 0.7428 - val_accuracy: 0.7793
Epoch 25/30
68/68 [=====] - 4s 59ms/step - loss: 0.6190 - accuracy: 0.7957 - val_loss: 0.5888 - val_accuracy: 0.8242
Epoch 26/30
68/68 [=====] - 4s 58ms/step - loss: 0.6149 - accuracy: 0.7962 - val_loss: 0.6889 - val_accuracy: 0.8008
Epoch 27/30
```

```
68/68 [=====] - 4s 58ms/step - loss: 0.6544 - accuracy: 0.7818 - val_loss: 0.6567 - val_accuracy: 0.8027
Epoch 28/30
68/68 [=====] - 4s 58ms/step - loss: 0.5960 - accuracy: 0.8078 - val_loss: 0.6039 - val_accuracy: 0.8066
Epoch 29/30
68/68 [=====] - 4s 58ms/step - loss: 0.6071 - accuracy: 0.8050 - val_loss: 0.6293 - val_accuracy: 0.7930
Epoch 30/30
68/68 [=====] - 4s 58ms/step - loss: 0.5710 - accuracy: 0.8138 - val_loss: 0.6938 - val_accuracy: 0.7949
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 13:03:23.175250: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - ETA: 0s - loss: 5.1000 - accuracy: 0.3542
2023-03-17 13:03:30.681677: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - 9s 103ms/step - loss: 5.1000 - accuracy: 0.3542 - val_loss: 3.0365 - val_accuracy: 0.4121
Epoch 2/30
68/68 [=====] - 4s 62ms/step - loss: 2.2208 - accuracy: 0.5590 - val_loss: 1.7622 - val_accuracy: 0.5859
Epoch 3/30
68/68 [=====] - 4s 60ms/step - loss: 1.5278 - accuracy: 0.5915 - val_loss: 1.3517 - val_accuracy: 0.5898
Epoch 4/30
68/68 [=====] - 4s 60ms/step - loss: 1.2150 - accuracy: 0.6295 - val_loss: 1.1282 - val_accuracy: 0.6211
Epoch 5/30
68/68 [=====] - 4s 63ms/step - loss: 1.1020 - accuracy: 0.6309 - val_loss: 1.0588 - val_accuracy: 0.6504
Epoch 6/30
68/68 [=====] - 4s 61ms/step - loss: 0.9977 - accuracy: 0.6537 - val_loss: 0.9652 - val_accuracy: 0.6523
Epoch 7/30
68/68 [=====] - 4s 61ms/step - loss: 0.9496 - accuracy: 0.6592 - val_loss: 0.9033 - val_accuracy: 0.6699
Epoch 8/30
```

```
68/68 [=====] - 4s 60ms/step - loss: 0.9207 - accuracy: 0.6588 - val_loss: 0.8986 - val_accuracy: 0.7031
Epoch 9/30
68/68 [=====] - 4s 59ms/step - loss: 0.8947 - accuracy: 0.6657 - val_loss: 0.8207 - val_accuracy: 0.6895
Epoch 10/30
68/68 [=====] - 4s 59ms/step - loss: 0.8520 - accuracy: 0.6834 - val_loss: 0.9140 - val_accuracy: 0.6816
Epoch 11/30
68/68 [=====] - 4s 61ms/step - loss: 0.7839 - accuracy: 0.7140 - val_loss: 0.7954 - val_accuracy: 0.7168
Epoch 12/30
68/68 [=====] - 4s 60ms/step - loss: 0.7893 - accuracy: 0.7210 - val_loss: 0.7835 - val_accuracy: 0.7461
Epoch 13/30
68/68 [=====] - 4s 59ms/step - loss: 0.7410 - accuracy: 0.7488 - val_loss: 0.7768 - val_accuracy: 0.7461
Epoch 14/30
68/68 [=====] - 4s 59ms/step - loss: 0.7687 - accuracy: 0.7400 - val_loss: 0.7323 - val_accuracy: 0.7500
Epoch 15/30
68/68 [=====] - 4s 59ms/step - loss: 0.6872 - accuracy: 0.7665 - val_loss: 0.7309 - val_accuracy: 0.7090
Epoch 16/30
68/68 [=====] - 4s 59ms/step - loss: 0.7187 - accuracy: 0.7623 - val_loss: 0.8046 - val_accuracy: 0.7051
Epoch 17/30
68/68 [=====] - 4s 59ms/step - loss: 0.6766 - accuracy: 0.7674 - val_loss: 0.6190 - val_accuracy: 0.8027
Epoch 18/30
68/68 [=====] - 4s 59ms/step - loss: 0.6854 - accuracy: 0.7753 - val_loss: 0.7531 - val_accuracy: 0.7285
Epoch 19/30
68/68 [=====] - 4s 59ms/step - loss: 0.6646 - accuracy: 0.7790 - val_loss: 0.6238 - val_accuracy: 0.8027
Epoch 20/30
68/68 [=====] - 4s 59ms/step - loss: 0.6534 - accuracy: 0.7832 - val_loss: 0.6626 - val_accuracy: 0.7734
Epoch 21/30
68/68 [=====] - 4s 59ms/step - loss: 0.6403 - accuracy: 0.7846 - val_loss: 0.6838 - val_accuracy: 0.7695
Epoch 22/30
68/68 [=====] - 4s 59ms/step - loss: 0.6344 - accuracy: 0.7869 - val_loss: 0.5818 - val_accuracy: 0.8281
Epoch 23/30
68/68 [=====] - 4s 60ms/step - loss: 0.6492 - accuracy: 0.7827 - val_loss: 0.6817 - val_accuracy: 0.7676
Epoch 24/30
```

```
68/68 [=====] - 4s 59ms/step - loss: 0.6010 - accuracy: 0.8041 - val_loss: 0.5844 - val_accuracy: 0.8184
Epoch 25/30
68/68 [=====] - 4s 59ms/step - loss: 0.6111 - accuracy: 0.7976 - val_loss: 0.6733 - val_accuracy: 0.7676
Epoch 26/30
68/68 [=====] - 4s 59ms/step - loss: 0.6123 - accuracy: 0.8036 - val_loss: 0.7212 - val_accuracy: 0.7520
Epoch 27/30
68/68 [=====] - 4s 59ms/step - loss: 0.5617 - accuracy: 0.8217 - val_loss: 0.6117 - val_accuracy: 0.7852
Epoch 28/30
68/68 [=====] - 4s 59ms/step - loss: 0.5943 - accuracy: 0.8087 - val_loss: 0.6032 - val_accuracy: 0.8105
Epoch 29/30
68/68 [=====] - 4s 59ms/step - loss: 0.5817 - accuracy: 0.8073 - val_loss: 0.7601 - val_accuracy: 0.7520
Epoch 30/30
68/68 [=====] - 4s 60ms/step - loss: 0.5778 - accuracy: 0.8148 - val_loss: 0.6220 - val_accuracy: 0.8184
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 13:05:30.707187: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - ETA: 0s - loss: 2.1097 - accuracy: 0.3422
2023-03-17 13:05:38.085341: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - 9s 102ms/step - loss: 2.1097 - accuracy: 0.3422 - val_loss: 1.6422 - val_accuracy: 0.5117
Epoch 2/30
68/68 [=====] - 4s 60ms/step - loss: 1.4444 - accuracy: 0.5729 - val_loss: 1.3870 - val_accuracy: 0.6035
Epoch 3/30
68/68 [=====] - 4s 59ms/step - loss: 1.2829 - accuracy: 0.5910 - val_loss: 1.1802 - val_accuracy: 0.6523
Epoch 4/30
68/68 [=====] - 4s 58ms/step - loss: 1.0910 - accuracy: 0.6444 - val_loss: 1.1047 - val_accuracy: 0.6699
Epoch 5/30
```

```
68/68 [=====] - 4s 59ms/step - loss: 1.0107 - accuracy: 0.6616 - val_loss: 0.9761 - val_accuracy: 0.6777
Epoch 6/30
68/68 [=====] - 4s 57ms/step - loss: 0.9409 - accuracy: 0.6903 - val_loss: 0.9482 - val_accuracy: 0.7246
Epoch 7/30
68/68 [=====] - 4s 59ms/step - loss: 0.8628 - accuracy: 0.7275 - val_loss: 0.8320 - val_accuracy: 0.7773
Epoch 8/30
68/68 [=====] - 4s 56ms/step - loss: 0.7778 - accuracy: 0.7725 - val_loss: 0.7099 - val_accuracy: 0.8027
Epoch 9/30
68/68 [=====] - 4s 56ms/step - loss: 0.6978 - accuracy: 0.7934 - val_loss: 0.7069 - val_accuracy: 0.8008
Epoch 10/30
68/68 [=====] - 4s 56ms/step - loss: 0.6675 - accuracy: 0.8073 - val_loss: 0.7151 - val_accuracy: 0.7969
Epoch 11/30
68/68 [=====] - 4s 56ms/step - loss: 0.6116 - accuracy: 0.8217 - val_loss: 0.6536 - val_accuracy: 0.8086
Epoch 12/30
68/68 [=====] - 4s 57ms/step - loss: 0.6058 - accuracy: 0.8180 - val_loss: 0.7413 - val_accuracy: 0.7852
Epoch 13/30
68/68 [=====] - 4s 56ms/step - loss: 0.5999 - accuracy: 0.8180 - val_loss: 0.5884 - val_accuracy: 0.8438
Epoch 14/30
68/68 [=====] - 4s 57ms/step - loss: 0.5596 - accuracy: 0.8347 - val_loss: 0.5661 - val_accuracy: 0.8262
Epoch 15/30
68/68 [=====] - 4s 56ms/step - loss: 0.5389 - accuracy: 0.8357 - val_loss: 0.6856 - val_accuracy: 0.8184
Epoch 16/30
68/68 [=====] - 4s 56ms/step - loss: 0.4691 - accuracy: 0.8589 - val_loss: 0.5737 - val_accuracy: 0.8457
Epoch 17/30
68/68 [=====] - 4s 56ms/step - loss: 0.5014 - accuracy: 0.8528 - val_loss: 0.5664 - val_accuracy: 0.8633
Epoch 18/30
68/68 [=====] - 4s 56ms/step - loss: 0.4672 - accuracy: 0.8575 - val_loss: 0.5133 - val_accuracy: 0.8672
Epoch 19/30
68/68 [=====] - 4s 56ms/step - loss: 0.4640 - accuracy: 0.8607 - val_loss: 0.5481 - val_accuracy: 0.8555
Epoch 20/30
68/68 [=====] - 4s 56ms/step - loss: 0.4384 - accuracy: 0.8709 - val_loss: 0.5567 - val_accuracy: 0.8594
Epoch 21/30
```

```
68/68 [=====] - 4s 56ms/step - loss: 0.4460 - accuracy: 0.8612 - val_loss: 0.5627 - val_accuracy: 0.8242
Epoch 22/30
68/68 [=====] - 4s 56ms/step - loss: 0.4117 - accuracy: 0.8793 - val_loss: 0.5725 - val_accuracy: 0.8457
Epoch 23/30
68/68 [=====] - 4s 56ms/step - loss: 0.4143 - accuracy: 0.8770 - val_loss: 0.4719 - val_accuracy: 0.8555
Epoch 24/30
68/68 [=====] - 4s 56ms/step - loss: 0.4080 - accuracy: 0.8756 - val_loss: 0.5059 - val_accuracy: 0.8730
Epoch 25/30
68/68 [=====] - 4s 56ms/step - loss: 0.4010 - accuracy: 0.8853 - val_loss: 0.6574 - val_accuracy: 0.8125
Epoch 26/30
68/68 [=====] - 4s 56ms/step - loss: 0.3840 - accuracy: 0.8928 - val_loss: 0.4799 - val_accuracy: 0.8809
Epoch 27/30
68/68 [=====] - 4s 56ms/step - loss: 0.3659 - accuracy: 0.8993 - val_loss: 0.4664 - val_accuracy: 0.8887
Epoch 28/30
68/68 [=====] - 4s 56ms/step - loss: 0.3282 - accuracy: 0.9136 - val_loss: 0.5536 - val_accuracy: 0.8672
Epoch 29/30
68/68 [=====] - 4s 56ms/step - loss: 0.3644 - accuracy: 0.9123 - val_loss: 0.4959 - val_accuracy: 0.8750
Epoch 30/30
68/68 [=====] - 4s 56ms/step - loss: 0.4080 - accuracy: 0.8863 - val_loss: 0.4388 - val_accuracy: 0.8789
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30

2023-03-17 13:07:32.270949: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - ETA: 0s - loss: 2.0019 - accuracy: 0.3807

2023-03-17 13:07:40.003485: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - 9s 106ms/step - loss: 2.0019 - accuracy: 0.3807 - val_loss: 1.7057 - val_accuracy: 0.5156
Epoch 2/30
```

```
68/68 [=====] - 4s 60ms/step - loss: 1.4218 - accuracy: 0.5617 - val_loss: 1.3673 - val_accuracy: 0.5996
Epoch 3/30
68/68 [=====] - 4s 60ms/step - loss: 1.2002 - accuracy: 0.6267 - val_loss: 1.1731 - val_accuracy: 0.6191
Epoch 4/30
68/68 [=====] - 4s 58ms/step - loss: 1.0450 - accuracy: 0.6769 - val_loss: 1.0168 - val_accuracy: 0.6953
Epoch 5/30
68/68 [=====] - 4s 60ms/step - loss: 0.9513 - accuracy: 0.7010 - val_loss: 0.9509 - val_accuracy: 0.7070
Epoch 6/30
68/68 [=====] - 4s 59ms/step - loss: 0.8919 - accuracy: 0.7205 - val_loss: 0.9181 - val_accuracy: 0.7246
Epoch 7/30
68/68 [=====] - 4s 57ms/step - loss: 0.7821 - accuracy: 0.7623 - val_loss: 0.8198 - val_accuracy: 0.7676
Epoch 8/30
68/68 [=====] - 4s 58ms/step - loss: 0.7294 - accuracy: 0.7869 - val_loss: 0.8049 - val_accuracy: 0.7617
Epoch 9/30
68/68 [=====] - 4s 57ms/step - loss: 0.6738 - accuracy: 0.8097 - val_loss: 0.7063 - val_accuracy: 0.7930
Epoch 10/30
68/68 [=====] - 4s 58ms/step - loss: 0.6347 - accuracy: 0.8148 - val_loss: 0.7231 - val_accuracy: 0.7871
Epoch 11/30
68/68 [=====] - 4s 58ms/step - loss: 0.6154 - accuracy: 0.8097 - val_loss: 0.6344 - val_accuracy: 0.8066
Epoch 12/30
68/68 [=====] - 4s 57ms/step - loss: 0.5982 - accuracy: 0.8236 - val_loss: 0.5952 - val_accuracy: 0.8359
Epoch 13/30
68/68 [=====] - 4s 57ms/step - loss: 0.5884 - accuracy: 0.8315 - val_loss: 0.7033 - val_accuracy: 0.7910
Epoch 14/30
68/68 [=====] - 4s 57ms/step - loss: 0.5439 - accuracy: 0.8398 - val_loss: 0.5538 - val_accuracy: 0.8535
Epoch 15/30
68/68 [=====] - 4s 57ms/step - loss: 0.4983 - accuracy: 0.8575 - val_loss: 0.5331 - val_accuracy: 0.8613
Epoch 16/30
68/68 [=====] - 4s 57ms/step - loss: 0.4840 - accuracy: 0.8607 - val_loss: 0.6412 - val_accuracy: 0.8008
Epoch 17/30
68/68 [=====] - 4s 59ms/step - loss: 0.5580 - accuracy: 0.8292 - val_loss: 0.5246 - val_accuracy: 0.8516
Epoch 18/30
```

```
68/68 [=====] - 4s 57ms/step - loss: 0.4847 - accuracy: 0.8584 - val_loss: 0.6772 - val_accuracy: 0.7949
Epoch 19/30
68/68 [=====] - 4s 57ms/step - loss: 0.5250 - accuracy: 0.8510 - val_loss: 0.5454 - val_accuracy: 0.8340
Epoch 20/30
68/68 [=====] - 4s 57ms/step - loss: 0.4597 - accuracy: 0.8621 - val_loss: 0.5029 - val_accuracy: 0.8613
Epoch 21/30
68/68 [=====] - 4s 57ms/step - loss: 0.4612 - accuracy: 0.8668 - val_loss: 0.5243 - val_accuracy: 0.8496
Epoch 22/30
68/68 [=====] - 4s 58ms/step - loss: 0.4426 - accuracy: 0.8816 - val_loss: 0.4797 - val_accuracy: 0.8672
Epoch 23/30
68/68 [=====] - 4s 57ms/step - loss: 0.4091 - accuracy: 0.8951 - val_loss: 0.4762 - val_accuracy: 0.8594
Epoch 24/30
68/68 [=====] - 4s 57ms/step - loss: 0.4256 - accuracy: 0.8835 - val_loss: 0.4567 - val_accuracy: 0.8652
Epoch 25/30
68/68 [=====] - 4s 57ms/step - loss: 0.3943 - accuracy: 0.8988 - val_loss: 0.4234 - val_accuracy: 0.8887
Epoch 26/30
68/68 [=====] - 4s 57ms/step - loss: 0.4089 - accuracy: 0.8960 - val_loss: 0.4533 - val_accuracy: 0.8672
Epoch 27/30
68/68 [=====] - 4s 57ms/step - loss: 0.3702 - accuracy: 0.9062 - val_loss: 0.4539 - val_accuracy: 0.8652
Epoch 28/30
68/68 [=====] - 4s 57ms/step - loss: 0.3740 - accuracy: 0.9048 - val_loss: 0.4812 - val_accuracy: 0.8594
Epoch 29/30
68/68 [=====] - 4s 57ms/step - loss: 0.3719 - accuracy: 0.9071 - val_loss: 0.4644 - val_accuracy: 0.8652
Epoch 30/30
68/68 [=====] - 4s 57ms/step - loss: 0.3715 - accuracy: 0.9099 - val_loss: 0.4068 - val_accuracy: 0.8848
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 13:09:35.947757: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

```
68/68 [=====] - ETA: 0s - loss: 2.1855 - accuracy: 0.3069  
2023-03-17 13:09:43.668200: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.  
68/68 [=====] - 9s 108ms/step - loss: 2.1855 - accuracy: 0.3069 - val_loss: 1.9315 - val_accuracy: 0.3926  
Epoch 2/30  
68/68 [=====] - 4s 63ms/step - loss: 1.5086 - accuracy: 0.5441 - val_loss: 1.3512 - val_accuracy: 0.5723  
Epoch 3/30  
68/68 [=====] - 4s 62ms/step - loss: 1.3058 - accuracy: 0.5994 - val_loss: 1.2796 - val_accuracy: 0.6016  
Epoch 4/30  
68/68 [=====] - 4s 58ms/step - loss: 1.1334 - accuracy: 0.6323 - val_loss: 1.1433 - val_accuracy: 0.6270  
Epoch 5/30  
68/68 [=====] - 4s 63ms/step - loss: 1.0209 - accuracy: 0.6597 - val_loss: 1.0820 - val_accuracy: 0.6367  
Epoch 6/30  
68/68 [=====] - 4s 61ms/step - loss: 0.9296 - accuracy: 0.6978 - val_loss: 0.9447 - val_accuracy: 0.6738  
Epoch 7/30  
68/68 [=====] - 4s 59ms/step - loss: 0.8492 - accuracy: 0.7303 - val_loss: 0.9377 - val_accuracy: 0.7148  
Epoch 8/30  
68/68 [=====] - 4s 61ms/step - loss: 0.8481 - accuracy: 0.7261 - val_loss: 0.8622 - val_accuracy: 0.7324  
Epoch 9/30  
68/68 [=====] - 4s 60ms/step - loss: 0.7442 - accuracy: 0.7809 - val_loss: 0.9573 - val_accuracy: 0.7168  
Epoch 10/30  
68/68 [=====] - 4s 59ms/step - loss: 0.7205 - accuracy: 0.7739 - val_loss: 0.6962 - val_accuracy: 0.7832  
Epoch 11/30  
68/68 [=====] - 4s 58ms/step - loss: 0.6731 - accuracy: 0.7925 - val_loss: 0.6979 - val_accuracy: 0.8008  
Epoch 12/30  
68/68 [=====] - 4s 59ms/step - loss: 0.6260 - accuracy: 0.8166 - val_loss: 0.6724 - val_accuracy: 0.7715  
Epoch 13/30  
68/68 [=====] - 4s 58ms/step - loss: 0.6212 - accuracy: 0.8143 - val_loss: 0.7093 - val_accuracy: 0.7812  
Epoch 14/30  
68/68 [=====] - 4s 58ms/step - loss: 0.5810 - accuracy: 0.8213 - val_loss: 0.6044 - val_accuracy: 0.8145  
Epoch 15/30
```

```
68/68 [=====] - 4s 60ms/step - loss: 0.5412 - accuracy: 0.8408 - val_loss: 0.6165 - val_accuracy: 0.8203
Epoch 16/30
68/68 [=====] - 4s 58ms/step - loss: 0.5647 - accuracy: 0.8357 - val_loss: 0.5844 - val_accuracy: 0.8105
Epoch 17/30
68/68 [=====] - 4s 58ms/step - loss: 0.5140 - accuracy: 0.8468 - val_loss: 0.6114 - val_accuracy: 0.8145
Epoch 18/30
68/68 [=====] - 4s 58ms/step - loss: 0.5395 - accuracy: 0.8389 - val_loss: 0.5903 - val_accuracy: 0.8027
Epoch 19/30
68/68 [=====] - 4s 59ms/step - loss: 0.5016 - accuracy: 0.8491 - val_loss: 0.5261 - val_accuracy: 0.8203
Epoch 20/30
68/68 [=====] - 4s 58ms/step - loss: 0.4559 - accuracy: 0.8644 - val_loss: 0.5209 - val_accuracy: 0.8301
Epoch 21/30
68/68 [=====] - 4s 57ms/step - loss: 0.4662 - accuracy: 0.8640 - val_loss: 0.4599 - val_accuracy: 0.8730
Epoch 22/30
68/68 [=====] - 4s 58ms/step - loss: 0.4330 - accuracy: 0.8751 - val_loss: 0.5978 - val_accuracy: 0.8320
Epoch 23/30
68/68 [=====] - 4s 58ms/step - loss: 0.4554 - accuracy: 0.8677 - val_loss: 0.4655 - val_accuracy: 0.8594
Epoch 24/30
68/68 [=====] - 4s 59ms/step - loss: 0.4116 - accuracy: 0.8830 - val_loss: 0.5172 - val_accuracy: 0.8379
Epoch 25/30
68/68 [=====] - 4s 58ms/step - loss: 0.4261 - accuracy: 0.8779 - val_loss: 0.4739 - val_accuracy: 0.8652
Epoch 26/30
68/68 [=====] - 4s 58ms/step - loss: 0.4224 - accuracy: 0.8756 - val_loss: 0.5054 - val_accuracy: 0.8477
Epoch 27/30
68/68 [=====] - 4s 58ms/step - loss: 0.3775 - accuracy: 0.8983 - val_loss: 0.5174 - val_accuracy: 0.8379
Epoch 28/30
68/68 [=====] - 4s 58ms/step - loss: 0.4347 - accuracy: 0.8812 - val_loss: 0.5594 - val_accuracy: 0.8320
Epoch 29/30
68/68 [=====] - 4s 58ms/step - loss: 0.3993 - accuracy: 0.8909 - val_loss: 0.4613 - val_accuracy: 0.8789
Epoch 30/30
68/68 [=====] - 4s 58ms/step - loss: 0.3458 - accuracy: 0.9081 - val_loss: 0.5165 - val_accuracy: 0.8398
Found 2666 files belonging to 5 classes.
```

```
Using 2133 files for training.  
Found 2666 files belonging to 5 classes.  
Using 533 files for validation.  
Found 2499 files belonging to 5 classes.  
Epoch 1/30  
  
2023-03-17 13:11:41.931011: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.  
  
68/68 [=====] - ETA: 0s - loss: 2.1260 - accuracy:  
0.3449  
  
2023-03-17 13:11:49.733414: I  
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]  
Plugin optimizer for device_type GPU is enabled.  
  
68/68 [=====] - 9s 106ms/step - loss: 2.1260 -  
accuracy: 0.3449 - val_loss: 1.7946 - val_accuracy: 0.4570  
Epoch 2/30  
68/68 [=====] - 4s 61ms/step - loss: 1.5034 - accuracy:  
0.5474 - val_loss: 1.5905 - val_accuracy: 0.4551  
Epoch 3/30  
68/68 [=====] - 4s 63ms/step - loss: 1.2749 - accuracy:  
0.6054 - val_loss: 1.2254 - val_accuracy: 0.5664  
Epoch 4/30  
68/68 [=====] - 4s 62ms/step - loss: 1.1485 - accuracy:  
0.6439 - val_loss: 1.0756 - val_accuracy: 0.6484  
Epoch 5/30  
68/68 [=====] - 4s 59ms/step - loss: 1.0140 - accuracy:  
0.6871 - val_loss: 1.0978 - val_accuracy: 0.5996  
Epoch 6/30  
68/68 [=====] - 4s 61ms/step - loss: 0.9119 - accuracy:  
0.7163 - val_loss: 1.0031 - val_accuracy: 0.6777  
Epoch 7/30  
68/68 [=====] - 4s 59ms/step - loss: 0.8408 - accuracy:  
0.7396 - val_loss: 0.8733 - val_accuracy: 0.7207  
Epoch 8/30  
68/68 [=====] - 4s 61ms/step - loss: 0.7924 - accuracy:  
0.7623 - val_loss: 0.7495 - val_accuracy: 0.7969  
Epoch 9/30  
68/68 [=====] - 4s 59ms/step - loss: 0.7406 - accuracy:  
0.7851 - val_loss: 0.7683 - val_accuracy: 0.7715  
Epoch 10/30  
68/68 [=====] - 4s 63ms/step - loss: 0.6842 - accuracy:  
0.8032 - val_loss: 0.7138 - val_accuracy: 0.7891  
Epoch 11/30  
68/68 [=====] - 4s 60ms/step - loss: 0.6332 - accuracy:  
0.8171 - val_loss: 0.6630 - val_accuracy: 0.8047  
Epoch 12/30
```

```
68/68 [=====] - 4s 60ms/step - loss: 0.6081 - accuracy: 0.8217 - val_loss: 0.6530 - val_accuracy: 0.8008
Epoch 13/30
68/68 [=====] - 4s 59ms/step - loss: 0.5959 - accuracy: 0.8296 - val_loss: 0.5644 - val_accuracy: 0.8398
Epoch 14/30
68/68 [=====] - 4s 58ms/step - loss: 0.6427 - accuracy: 0.8134 - val_loss: 0.6566 - val_accuracy: 0.8027
Epoch 15/30
68/68 [=====] - 4s 58ms/step - loss: 0.5406 - accuracy: 0.8449 - val_loss: 0.5043 - val_accuracy: 0.8633
Epoch 16/30
68/68 [=====] - 4s 59ms/step - loss: 0.5177 - accuracy: 0.8510 - val_loss: 0.5394 - val_accuracy: 0.8535
Epoch 17/30
68/68 [=====] - 4s 60ms/step - loss: 0.4990 - accuracy: 0.8579 - val_loss: 0.6108 - val_accuracy: 0.8242
Epoch 18/30
68/68 [=====] - 4s 58ms/step - loss: 0.4876 - accuracy: 0.8584 - val_loss: 0.5446 - val_accuracy: 0.8457
Epoch 19/30
68/68 [=====] - 4s 59ms/step - loss: 0.5106 - accuracy: 0.8500 - val_loss: 0.6013 - val_accuracy: 0.8203
Epoch 20/30
68/68 [=====] - 4s 58ms/step - loss: 0.4718 - accuracy: 0.8682 - val_loss: 0.5675 - val_accuracy: 0.8359
Epoch 21/30
68/68 [=====] - 4s 59ms/step - loss: 0.4643 - accuracy: 0.8747 - val_loss: 0.5372 - val_accuracy: 0.8516
Epoch 22/30
68/68 [=====] - 4s 58ms/step - loss: 0.4983 - accuracy: 0.8565 - val_loss: 0.4729 - val_accuracy: 0.8633
Epoch 23/30
68/68 [=====] - 4s 58ms/step - loss: 0.4268 - accuracy: 0.8858 - val_loss: 0.5199 - val_accuracy: 0.8555
Epoch 24/30
68/68 [=====] - 4s 59ms/step - loss: 0.4112 - accuracy: 0.8932 - val_loss: 0.4710 - val_accuracy: 0.8711
Epoch 25/30
68/68 [=====] - 4s 58ms/step - loss: 0.4073 - accuracy: 0.8867 - val_loss: 0.4829 - val_accuracy: 0.8574
Epoch 26/30
68/68 [=====] - 4s 58ms/step - loss: 0.4150 - accuracy: 0.8895 - val_loss: 0.4931 - val_accuracy: 0.8652
Epoch 27/30
68/68 [=====] - 4s 58ms/step - loss: 0.4025 - accuracy: 0.8955 - val_loss: 0.5566 - val_accuracy: 0.8398
Epoch 28/30
```

```
68/68 [=====] - 4s 58ms/step - loss: 0.3528 - accuracy: 0.9081 - val_loss: 0.4472 - val_accuracy: 0.8926
Epoch 29/30
68/68 [=====] - 4s 58ms/step - loss: 0.3790 - accuracy: 0.8960 - val_loss: 0.6180 - val_accuracy: 0.8203
Epoch 30/30
68/68 [=====] - 4s 58ms/step - loss: 0.3773 - accuracy: 0.9104 - val_loss: 0.4491 - val_accuracy: 0.8750
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 13:13:49.132308: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - ETA: 0s - loss: 2.0505 - accuracy: 0.3682
2023-03-17 13:13:57.152302: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - 9s 110ms/step - loss: 2.0505 - accuracy: 0.3682 - val_loss: 1.5902 - val_accuracy: 0.4766
Epoch 2/30
68/68 [=====] - 4s 65ms/step - loss: 1.4336 - accuracy: 0.5799 - val_loss: 1.2778 - val_accuracy: 0.5762
Epoch 3/30
68/68 [=====] - 4s 60ms/step - loss: 1.1790 - accuracy: 0.6360 - val_loss: 1.3159 - val_accuracy: 0.5781
Epoch 4/30
68/68 [=====] - 4s 62ms/step - loss: 1.0852 - accuracy: 0.6569 - val_loss: 1.0155 - val_accuracy: 0.6094
Epoch 5/30
68/68 [=====] - 4s 59ms/step - loss: 0.9671 - accuracy: 0.6862 - val_loss: 0.9952 - val_accuracy: 0.6191
Epoch 6/30
68/68 [=====] - 4s 63ms/step - loss: 0.9143 - accuracy: 0.6857 - val_loss: 1.0159 - val_accuracy: 0.6582
Epoch 7/30
68/68 [=====] - 4s 62ms/step - loss: 0.8779 - accuracy: 0.7024 - val_loss: 0.8949 - val_accuracy: 0.6816
Epoch 8/30
68/68 [=====] - 4s 62ms/step - loss: 0.8216 - accuracy: 0.7307 - val_loss: 1.0412 - val_accuracy: 0.6719
Epoch 9/30
```

```
68/68 [=====] - 4s 63ms/step - loss: 0.7702 - accuracy: 0.7461 - val_loss: 0.7489 - val_accuracy: 0.7227
Epoch 10/30
68/68 [=====] - 4s 62ms/step - loss: 0.7089 - accuracy: 0.7767 - val_loss: 0.7383 - val_accuracy: 0.7637
Epoch 11/30
68/68 [=====] - 4s 60ms/step - loss: 0.6913 - accuracy: 0.7734 - val_loss: 0.7925 - val_accuracy: 0.7207
Epoch 12/30
68/68 [=====] - 4s 61ms/step - loss: 0.6234 - accuracy: 0.8069 - val_loss: 0.7236 - val_accuracy: 0.7656
Epoch 13/30
68/68 [=====] - 4s 60ms/step - loss: 0.5958 - accuracy: 0.8069 - val_loss: 0.5786 - val_accuracy: 0.8398
Epoch 14/30
68/68 [=====] - 4s 60ms/step - loss: 0.6041 - accuracy: 0.8032 - val_loss: 0.6462 - val_accuracy: 0.7773
Epoch 15/30
68/68 [=====] - 4s 60ms/step - loss: 0.5472 - accuracy: 0.8333 - val_loss: 0.7948 - val_accuracy: 0.7578
Epoch 16/30
68/68 [=====] - 4s 60ms/step - loss: 0.5458 - accuracy: 0.8296 - val_loss: 0.7312 - val_accuracy: 0.7559
Epoch 17/30
68/68 [=====] - 4s 60ms/step - loss: 0.5092 - accuracy: 0.8408 - val_loss: 0.5569 - val_accuracy: 0.8359
Epoch 18/30
68/68 [=====] - 4s 60ms/step - loss: 0.4793 - accuracy: 0.8542 - val_loss: 0.5304 - val_accuracy: 0.8359
Epoch 19/30
68/68 [=====] - 4s 60ms/step - loss: 0.5229 - accuracy: 0.8380 - val_loss: 0.5456 - val_accuracy: 0.8320
Epoch 20/30
68/68 [=====] - 4s 60ms/step - loss: 0.4921 - accuracy: 0.8500 - val_loss: 0.5342 - val_accuracy: 0.8418
Epoch 21/30
68/68 [=====] - 4s 60ms/step - loss: 0.4618 - accuracy: 0.8630 - val_loss: 0.5094 - val_accuracy: 0.8457
Epoch 22/30
68/68 [=====] - 4s 60ms/step - loss: 0.4718 - accuracy: 0.8612 - val_loss: 0.5047 - val_accuracy: 0.8574
Epoch 23/30
68/68 [=====] - 4s 62ms/step - loss: 0.4404 - accuracy: 0.8663 - val_loss: 0.5164 - val_accuracy: 0.8496
Epoch 24/30
68/68 [=====] - 4s 60ms/step - loss: 0.4444 - accuracy: 0.8700 - val_loss: 0.4425 - val_accuracy: 0.8574
Epoch 25/30
```

```
68/68 [=====] - 4s 60ms/step - loss: 0.4109 - accuracy: 0.8807 - val_loss: 0.4441 - val_accuracy: 0.8730
Epoch 26/30
68/68 [=====] - 4s 61ms/step - loss: 0.4109 - accuracy: 0.8886 - val_loss: 0.4225 - val_accuracy: 0.8789
Epoch 27/30
68/68 [=====] - 4s 60ms/step - loss: 0.4325 - accuracy: 0.8733 - val_loss: 0.4653 - val_accuracy: 0.8770
Epoch 28/30
68/68 [=====] - 4s 61ms/step - loss: 0.3972 - accuracy: 0.8909 - val_loss: 0.4320 - val_accuracy: 0.8926
Epoch 29/30
68/68 [=====] - 4s 60ms/step - loss: 0.3899 - accuracy: 0.8946 - val_loss: 0.5088 - val_accuracy: 0.8516
Epoch 30/30
68/68 [=====] - 4s 61ms/step - loss: 0.3522 - accuracy: 0.9044 - val_loss: 0.6171 - val_accuracy: 0.8145
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 13:15:59.458353: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - ETA: 0s - loss: 1.6662 - accuracy: 0.2953
2023-03-17 13:16:07.750005: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - 10s 111ms/step - loss: 1.6662 - accuracy: 0.2953 - val_loss: 1.5346 - val_accuracy: 0.4219
Epoch 2/30
68/68 [=====] - 4s 64ms/step - loss: 1.1551 - accuracy: 0.5265 - val_loss: 1.1192 - val_accuracy: 0.5547
Epoch 3/30
68/68 [=====] - 4s 61ms/step - loss: 0.9796 - accuracy: 0.6184 - val_loss: 1.0054 - val_accuracy: 0.6289
Epoch 4/30
68/68 [=====] - 4s 59ms/step - loss: 0.8623 - accuracy: 0.6741 - val_loss: 0.9795 - val_accuracy: 0.6348
Epoch 5/30
68/68 [=====] - 4s 60ms/step - loss: 0.8145 - accuracy: 0.7033 - val_loss: 0.8323 - val_accuracy: 0.7305
Epoch 6/30
```

```
68/68 [=====] - 4s 57ms/step - loss: 0.7619 - accuracy: 0.7312 - val_loss: 0.8221 - val_accuracy: 0.7266
Epoch 7/30
68/68 [=====] - 4s 59ms/step - loss: 0.6908 - accuracy: 0.7577 - val_loss: 0.7503 - val_accuracy: 0.7500
Epoch 8/30
68/68 [=====] - 4s 56ms/step - loss: 0.6460 - accuracy: 0.7739 - val_loss: 0.7014 - val_accuracy: 0.7715
Epoch 9/30
68/68 [=====] - 4s 57ms/step - loss: 0.5936 - accuracy: 0.8092 - val_loss: 0.8519 - val_accuracy: 0.7383
Epoch 10/30
68/68 [=====] - 4s 57ms/step - loss: 0.5544 - accuracy: 0.8175 - val_loss: 0.6294 - val_accuracy: 0.7969
Epoch 11/30
68/68 [=====] - 4s 57ms/step - loss: 0.5479 - accuracy: 0.8217 - val_loss: 0.6049 - val_accuracy: 0.8164
Epoch 12/30
68/68 [=====] - 4s 58ms/step - loss: 0.4885 - accuracy: 0.8431 - val_loss: 0.7184 - val_accuracy: 0.7871
Epoch 13/30
68/68 [=====] - 4s 57ms/step - loss: 0.4872 - accuracy: 0.8496 - val_loss: 0.6154 - val_accuracy: 0.8164
Epoch 14/30
68/68 [=====] - 4s 56ms/step - loss: 0.4497 - accuracy: 0.8570 - val_loss: 0.5443 - val_accuracy: 0.8262
Epoch 15/30
68/68 [=====] - 4s 57ms/step - loss: 0.4063 - accuracy: 0.8672 - val_loss: 0.5552 - val_accuracy: 0.8320
Epoch 16/30
68/68 [=====] - 4s 57ms/step - loss: 0.5064 - accuracy: 0.8213 - val_loss: 0.5344 - val_accuracy: 0.8301
Epoch 17/30
68/68 [=====] - 4s 58ms/step - loss: 0.4040 - accuracy: 0.8719 - val_loss: 0.6168 - val_accuracy: 0.8164
Epoch 18/30
68/68 [=====] - 4s 60ms/step - loss: 0.4099 - accuracy: 0.8714 - val_loss: 0.5302 - val_accuracy: 0.8477
Epoch 19/30
68/68 [=====] - 4s 57ms/step - loss: 0.4050 - accuracy: 0.8807 - val_loss: 0.5165 - val_accuracy: 0.8340
Epoch 20/30
68/68 [=====] - 4s 61ms/step - loss: 0.3316 - accuracy: 0.9020 - val_loss: 0.4914 - val_accuracy: 0.8496
Epoch 21/30
68/68 [=====] - 4s 57ms/step - loss: 0.3314 - accuracy: 0.9030 - val_loss: 0.5474 - val_accuracy: 0.8398
Epoch 22/30
```

```
68/68 [=====] - 4s 57ms/step - loss: 0.3110 - accuracy: 0.9062 - val_loss: 0.5276 - val_accuracy: 0.8457
Epoch 23/30
68/68 [=====] - 4s 57ms/step - loss: 0.2830 - accuracy: 0.9188 - val_loss: 0.5317 - val_accuracy: 0.8379
Epoch 24/30
68/68 [=====] - 4s 59ms/step - loss: 0.2648 - accuracy: 0.9150 - val_loss: 0.5991 - val_accuracy: 0.8262
Epoch 25/30
68/68 [=====] - 4s 63ms/step - loss: 0.3069 - accuracy: 0.9141 - val_loss: 0.4708 - val_accuracy: 0.8672
Epoch 26/30
68/68 [=====] - 4s 58ms/step - loss: 0.2885 - accuracy: 0.9155 - val_loss: 0.5943 - val_accuracy: 0.8262
Epoch 27/30
68/68 [=====] - 4s 56ms/step - loss: 0.2828 - accuracy: 0.9150 - val_loss: 0.4600 - val_accuracy: 0.8730
Epoch 28/30
68/68 [=====] - 4s 56ms/step - loss: 0.2683 - accuracy: 0.9220 - val_loss: 0.4327 - val_accuracy: 0.8828
Epoch 29/30
68/68 [=====] - 4s 57ms/step - loss: 0.2550 - accuracy: 0.9220 - val_loss: 0.3939 - val_accuracy: 0.8789
Epoch 30/30
68/68 [=====] - 4s 57ms/step - loss: 0.2229 - accuracy: 0.9443 - val_loss: 0.4481 - val_accuracy: 0.8906
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 13:18:05.328487: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - ETA: 0s - loss: 1.5576 - accuracy: 0.3626
2023-03-17 13:18:12.189483: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - 8s 105ms/step - loss: 1.5576 - accuracy: 0.3626 - val_loss: 1.2081 - val_accuracy: 0.5234
Epoch 2/30
68/68 [=====] - 5s 69ms/step - loss: 1.0882 - accuracy: 0.5664 - val_loss: 1.1126 - val_accuracy: 0.5547
Epoch 3/30
```

```
68/68 [=====] - 4s 61ms/step - loss: 0.9905 - accuracy: 0.6128 - val_loss: 0.9707 - val_accuracy: 0.6367
Epoch 4/30
68/68 [=====] - 4s 59ms/step - loss: 0.8591 - accuracy: 0.6718 - val_loss: 1.0676 - val_accuracy: 0.6230
Epoch 5/30
68/68 [=====] - 4s 60ms/step - loss: 0.7788 - accuracy: 0.7349 - val_loss: 0.9230 - val_accuracy: 0.6367
Epoch 6/30
68/68 [=====] - 4s 57ms/step - loss: 0.7051 - accuracy: 0.7470 - val_loss: 0.7312 - val_accuracy: 0.7656
Epoch 7/30
68/68 [=====] - 4s 58ms/step - loss: 0.6411 - accuracy: 0.7883 - val_loss: 0.7041 - val_accuracy: 0.7734
Epoch 8/30
68/68 [=====] - 4s 59ms/step - loss: 0.5854 - accuracy: 0.8115 - val_loss: 0.5721 - val_accuracy: 0.8105
Epoch 9/30
68/68 [=====] - 4s 58ms/step - loss: 0.5559 - accuracy: 0.8217 - val_loss: 0.7211 - val_accuracy: 0.7363
Epoch 10/30
68/68 [=====] - 4s 57ms/step - loss: 0.5344 - accuracy: 0.8217 - val_loss: 0.5851 - val_accuracy: 0.8203
Epoch 11/30
68/68 [=====] - 4s 58ms/step - loss: 0.5180 - accuracy: 0.8333 - val_loss: 0.5130 - val_accuracy: 0.8379
Epoch 12/30
68/68 [=====] - 4s 57ms/step - loss: 0.4924 - accuracy: 0.8333 - val_loss: 0.5480 - val_accuracy: 0.8301
Epoch 13/30
68/68 [=====] - 4s 57ms/step - loss: 0.4491 - accuracy: 0.8528 - val_loss: 0.5291 - val_accuracy: 0.8301
Epoch 14/30
68/68 [=====] - 4s 57ms/step - loss: 0.4240 - accuracy: 0.8621 - val_loss: 0.4135 - val_accuracy: 0.8672
Epoch 15/30
68/68 [=====] - 4s 58ms/step - loss: 0.3939 - accuracy: 0.8686 - val_loss: 0.5330 - val_accuracy: 0.8281
Epoch 16/30
68/68 [=====] - 4s 57ms/step - loss: 0.3717 - accuracy: 0.8872 - val_loss: 0.4804 - val_accuracy: 0.8496
Epoch 17/30
68/68 [=====] - 4s 57ms/step - loss: 0.3397 - accuracy: 0.8965 - val_loss: 0.4953 - val_accuracy: 0.8496
Epoch 18/30
68/68 [=====] - 4s 58ms/step - loss: 0.3883 - accuracy: 0.8705 - val_loss: 0.4241 - val_accuracy: 0.8750
Epoch 19/30
```

```
68/68 [=====] - 4s 59ms/step - loss: 0.3458 - accuracy: 0.8886 - val_loss: 0.4065 - val_accuracy: 0.8770
Epoch 20/30
68/68 [=====] - 4s 59ms/step - loss: 0.3226 - accuracy: 0.9044 - val_loss: 0.3901 - val_accuracy: 0.8848
Epoch 21/30
68/68 [=====] - 4s 58ms/step - loss: 0.3558 - accuracy: 0.8918 - val_loss: 0.3476 - val_accuracy: 0.8945
Epoch 22/30
68/68 [=====] - 4s 58ms/step - loss: 0.3104 - accuracy: 0.9062 - val_loss: 0.3947 - val_accuracy: 0.8867
Epoch 23/30
68/68 [=====] - 4s 60ms/step - loss: 0.2804 - accuracy: 0.9211 - val_loss: 0.3208 - val_accuracy: 0.9082
Epoch 24/30
68/68 [=====] - 4s 58ms/step - loss: 0.2776 - accuracy: 0.9201 - val_loss: 0.3927 - val_accuracy: 0.8926
Epoch 25/30
68/68 [=====] - 4s 58ms/step - loss: 0.2548 - accuracy: 0.9253 - val_loss: 0.4063 - val_accuracy: 0.8730
Epoch 26/30
68/68 [=====] - 4s 59ms/step - loss: 0.2839 - accuracy: 0.9183 - val_loss: 0.3035 - val_accuracy: 0.9238
Epoch 27/30
68/68 [=====] - 4s 59ms/step - loss: 0.2686 - accuracy: 0.9285 - val_loss: 0.3128 - val_accuracy: 0.9023
Epoch 28/30
68/68 [=====] - 4s 59ms/step - loss: 0.3236 - accuracy: 0.9109 - val_loss: 0.3981 - val_accuracy: 0.8867
Epoch 29/30
68/68 [=====] - 4s 58ms/step - loss: 0.2836 - accuracy: 0.9285 - val_loss: 0.5441 - val_accuracy: 0.8594
Epoch 30/30
68/68 [=====] - 4s 57ms/step - loss: 0.2382 - accuracy: 0.9438 - val_loss: 0.3389 - val_accuracy: 0.8984
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 13:20:11.760597: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - ETA: 0s - loss: 1.5422 - accuracy: 0.3900
```

```
2023-03-17 13:20:19.625076: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - 9s 108ms/step - loss: 1.5422 - accuracy: 0.3900 - val_loss: 1.3731 - val_accuracy: 0.4277
Epoch 2/30
68/68 [=====] - 4s 64ms/step - loss: 1.1195 - accuracy: 0.5659 - val_loss: 1.0899 - val_accuracy: 0.5586
Epoch 3/30
68/68 [=====] - 4s 60ms/step - loss: 0.9591 - accuracy: 0.6300 - val_loss: 1.0148 - val_accuracy: 0.6172
Epoch 4/30
68/68 [=====] - 4s 60ms/step - loss: 0.8798 - accuracy: 0.6541 - val_loss: 0.8831 - val_accuracy: 0.6465
Epoch 5/30
68/68 [=====] - 4s 59ms/step - loss: 0.8088 - accuracy: 0.6996 - val_loss: 0.8334 - val_accuracy: 0.6875
Epoch 6/30
68/68 [=====] - 4s 59ms/step - loss: 0.7509 - accuracy: 0.7317 - val_loss: 0.9710 - val_accuracy: 0.6289
Epoch 7/30
68/68 [=====] - 4s 61ms/step - loss: 0.6937 - accuracy: 0.7632 - val_loss: 0.7370 - val_accuracy: 0.7812
Epoch 8/30
68/68 [=====] - 4s 62ms/step - loss: 0.6093 - accuracy: 0.7920 - val_loss: 0.6668 - val_accuracy: 0.7930
Epoch 9/30
68/68 [=====] - 4s 60ms/step - loss: 0.5599 - accuracy: 0.8208 - val_loss: 0.6615 - val_accuracy: 0.7832
Epoch 10/30
68/68 [=====] - 4s 61ms/step - loss: 0.5056 - accuracy: 0.8315 - val_loss: 0.5786 - val_accuracy: 0.8125
Epoch 11/30
68/68 [=====] - 4s 63ms/step - loss: 0.5082 - accuracy: 0.8319 - val_loss: 0.6070 - val_accuracy: 0.7988
Epoch 12/30
68/68 [=====] - 4s 59ms/step - loss: 0.4848 - accuracy: 0.8431 - val_loss: 0.5751 - val_accuracy: 0.8027
Epoch 13/30
68/68 [=====] - 4s 59ms/step - loss: 0.4622 - accuracy: 0.8477 - val_loss: 0.5048 - val_accuracy: 0.8301
Epoch 14/30
68/68 [=====] - 4s 58ms/step - loss: 0.4785 - accuracy: 0.8435 - val_loss: 0.4771 - val_accuracy: 0.8398
Epoch 15/30
68/68 [=====] - 4s 58ms/step - loss: 0.4251 - accuracy: 0.8691 - val_loss: 0.4185 - val_accuracy: 0.8672
```

```
Epoch 16/30
68/68 [=====] - 4s 58ms/step - loss: 0.4025 - accuracy: 0.8784 - val_loss: 0.4725 - val_accuracy: 0.8457
Epoch 17/30
68/68 [=====] - 4s 59ms/step - loss: 0.3942 - accuracy: 0.8770 - val_loss: 0.3783 - val_accuracy: 0.8711
Epoch 18/30
68/68 [=====] - 4s 59ms/step - loss: 0.3634 - accuracy: 0.8890 - val_loss: 0.5642 - val_accuracy: 0.7949
Epoch 19/30
68/68 [=====] - 4s 61ms/step - loss: 0.3368 - accuracy: 0.9011 - val_loss: 0.4789 - val_accuracy: 0.8438
Epoch 20/30
68/68 [=====] - 4s 59ms/step - loss: 0.3535 - accuracy: 0.8932 - val_loss: 0.4670 - val_accuracy: 0.8594
Epoch 21/30
68/68 [=====] - 4s 60ms/step - loss: 0.3608 - accuracy: 0.8849 - val_loss: 0.4331 - val_accuracy: 0.8457
Epoch 22/30
68/68 [=====] - 4s 59ms/step - loss: 0.2923 - accuracy: 0.9155 - val_loss: 0.3634 - val_accuracy: 0.8789
Epoch 23/30
68/68 [=====] - 4s 59ms/step - loss: 0.3004 - accuracy: 0.9123 - val_loss: 0.5333 - val_accuracy: 0.8301
Epoch 24/30
68/68 [=====] - 4s 59ms/step - loss: 0.2888 - accuracy: 0.9164 - val_loss: 0.3487 - val_accuracy: 0.8867
Epoch 25/30
68/68 [=====] - 4s 58ms/step - loss: 0.3103 - accuracy: 0.9104 - val_loss: 0.3442 - val_accuracy: 0.9004
Epoch 26/30
68/68 [=====] - 4s 60ms/step - loss: 0.2549 - accuracy: 0.9290 - val_loss: 0.2910 - val_accuracy: 0.9121
Epoch 27/30
68/68 [=====] - 4s 58ms/step - loss: 0.2630 - accuracy: 0.9290 - val_loss: 0.3993 - val_accuracy: 0.8770
Epoch 28/30
68/68 [=====] - 4s 58ms/step - loss: 0.2250 - accuracy: 0.9369 - val_loss: 0.3814 - val_accuracy: 0.8887
Epoch 29/30
68/68 [=====] - 4s 58ms/step - loss: 0.2808 - accuracy: 0.9234 - val_loss: 0.3950 - val_accuracy: 0.8848
Epoch 30/30
68/68 [=====] - 4s 59ms/step - loss: 0.2355 - accuracy: 0.9341 - val_loss: 0.3485 - val_accuracy: 0.8848
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
```

Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30

2023-03-17 13:22:19.473464: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - ETA: 0s - loss: 1.5699 - accuracy:
0.3835

2023-03-17 13:22:27.357034: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

68/68 [=====] - 9s 106ms/step - loss: 1.5699 -
accuracy: 0.3835 - val_loss: 1.3674 - val_accuracy: 0.4199
Epoch 2/30

68/68 [=====] - 5s 67ms/step - loss: 1.0676 - accuracy:
0.5840 - val_loss: 1.3078 - val_accuracy: 0.4453
Epoch 3/30

68/68 [=====] - 4s 62ms/step - loss: 0.9769 - accuracy:
0.6249 - val_loss: 1.0104 - val_accuracy: 0.6543
Epoch 4/30

68/68 [=====] - 4s 63ms/step - loss: 0.8636 - accuracy:
0.6871 - val_loss: 0.9076 - val_accuracy: 0.7051
Epoch 5/30

68/68 [=====] - 4s 59ms/step - loss: 0.7777 - accuracy:
0.7224 - val_loss: 0.7849 - val_accuracy: 0.6973
Epoch 6/30

68/68 [=====] - 4s 59ms/step - loss: 0.7026 - accuracy:
0.7651 - val_loss: 0.8152 - val_accuracy: 0.7402
Epoch 7/30

68/68 [=====] - 4s 60ms/step - loss: 0.6276 - accuracy:
0.7925 - val_loss: 0.6811 - val_accuracy: 0.7949
Epoch 8/30

68/68 [=====] - 4s 59ms/step - loss: 0.6448 - accuracy:
0.7883 - val_loss: 0.6875 - val_accuracy: 0.7695
Epoch 9/30

68/68 [=====] - 4s 59ms/step - loss: 0.5758 - accuracy:
0.8097 - val_loss: 0.8034 - val_accuracy: 0.7246
Epoch 10/30

68/68 [=====] - 4s 59ms/step - loss: 0.5868 - accuracy:
0.7934 - val_loss: 0.6003 - val_accuracy: 0.8125
Epoch 11/30

68/68 [=====] - 4s 62ms/step - loss: 0.5187 - accuracy:
0.8361 - val_loss: 0.7759 - val_accuracy: 0.7637
Epoch 12/30

68/68 [=====] - 4s 61ms/step - loss: 0.5311 - accuracy:
0.8180 - val_loss: 0.6263 - val_accuracy: 0.7891

```
Epoch 13/30
68/68 [=====] - 4s 59ms/step - loss: 0.4569 - accuracy: 0.8584 - val_loss: 0.6117 - val_accuracy: 0.7949
Epoch 14/30
68/68 [=====] - 4s 60ms/step - loss: 0.4495 - accuracy: 0.8542 - val_loss: 0.4943 - val_accuracy: 0.8438
Epoch 15/30
68/68 [=====] - 4s 59ms/step - loss: 0.4417 - accuracy: 0.8593 - val_loss: 0.4785 - val_accuracy: 0.8477
Epoch 16/30
68/68 [=====] - 4s 61ms/step - loss: 0.4185 - accuracy: 0.8723 - val_loss: 0.4768 - val_accuracy: 0.8574
Epoch 17/30
68/68 [=====] - 4s 61ms/step - loss: 0.3807 - accuracy: 0.8821 - val_loss: 0.5839 - val_accuracy: 0.8086
Epoch 18/30
68/68 [=====] - 4s 60ms/step - loss: 0.3962 - accuracy: 0.8742 - val_loss: 0.5986 - val_accuracy: 0.8125
Epoch 19/30
68/68 [=====] - 4s 60ms/step - loss: 0.3646 - accuracy: 0.8830 - val_loss: 0.5293 - val_accuracy: 0.8184
Epoch 20/30
68/68 [=====] - 4s 59ms/step - loss: 0.3697 - accuracy: 0.8812 - val_loss: 0.4856 - val_accuracy: 0.8379
Epoch 21/30
68/68 [=====] - 4s 60ms/step - loss: 0.3169 - accuracy: 0.9095 - val_loss: 0.4451 - val_accuracy: 0.8555
Epoch 22/30
68/68 [=====] - 4s 59ms/step - loss: 0.3168 - accuracy: 0.9044 - val_loss: 0.4048 - val_accuracy: 0.8906
Epoch 23/30
68/68 [=====] - 4s 59ms/step - loss: 0.3170 - accuracy: 0.9002 - val_loss: 0.3823 - val_accuracy: 0.9023
Epoch 24/30
68/68 [=====] - 4s 60ms/step - loss: 0.2701 - accuracy: 0.9280 - val_loss: 0.3163 - val_accuracy: 0.9238
Epoch 25/30
68/68 [=====] - 4s 60ms/step - loss: 0.2831 - accuracy: 0.9211 - val_loss: 0.4774 - val_accuracy: 0.8516
Epoch 26/30
68/68 [=====] - 4s 59ms/step - loss: 0.2632 - accuracy: 0.9229 - val_loss: 0.3577 - val_accuracy: 0.8984
Epoch 27/30
68/68 [=====] - 4s 59ms/step - loss: 0.2839 - accuracy: 0.9197 - val_loss: 0.5375 - val_accuracy: 0.8340
Epoch 28/30
68/68 [=====] - 4s 59ms/step - loss: 0.3239 - accuracy: 0.9095 - val_loss: 0.4109 - val_accuracy: 0.8945
```

```
Epoch 29/30
68/68 [=====] - 4s 59ms/step - loss: 0.2484 - accuracy: 0.9369 - val_loss: 0.4573 - val_accuracy: 0.8730
Epoch 30/30
68/68 [=====] - 4s 59ms/step - loss: 0.2326 - accuracy: 0.9420 - val_loss: 0.4695 - val_accuracy: 0.8672
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Epoch 1/30
2023-03-17 13:24:28.273044: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - ETA: 0s - loss: 1.5086 - accuracy: 0.3788
2023-03-17 13:24:35.454565: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
68/68 [=====] - 8s 97ms/step - loss: 1.5086 - accuracy: 0.3788 - val_loss: 1.2633 - val_accuracy: 0.4785
Epoch 2/30
68/68 [=====] - 4s 64ms/step - loss: 1.1055 - accuracy: 0.5534 - val_loss: 1.2033 - val_accuracy: 0.4805
Epoch 3/30
68/68 [=====] - 5s 68ms/step - loss: 0.9638 - accuracy: 0.6198 - val_loss: 0.9063 - val_accuracy: 0.6191
Epoch 4/30
68/68 [=====] - 4s 62ms/step - loss: 0.8744 - accuracy: 0.6727 - val_loss: 1.0363 - val_accuracy: 0.5840
Epoch 5/30
68/68 [=====] - 4s 63ms/step - loss: 0.7692 - accuracy: 0.7340 - val_loss: 0.7476 - val_accuracy: 0.7266
Epoch 6/30
68/68 [=====] - 4s 60ms/step - loss: 0.7089 - accuracy: 0.7591 - val_loss: 0.7185 - val_accuracy: 0.7617
Epoch 7/30
68/68 [=====] - 4s 61ms/step - loss: 0.6800 - accuracy: 0.7688 - val_loss: 0.8569 - val_accuracy: 0.6953
Epoch 8/30
68/68 [=====] - 4s 62ms/step - loss: 0.6368 - accuracy: 0.7767 - val_loss: 0.5639 - val_accuracy: 0.8066
Epoch 9/30
68/68 [=====] - 4s 60ms/step - loss: 0.5624 - accuracy: 0.8124 - val_loss: 0.5502 - val_accuracy: 0.8105
```

```
Epoch 10/30
68/68 [=====] - 4s 63ms/step - loss: 0.5099 - accuracy: 0.8357 - val_loss: 0.5703 - val_accuracy: 0.8047
Epoch 11/30
68/68 [=====] - 4s 62ms/step - loss: 0.4938 - accuracy: 0.8315 - val_loss: 0.4924 - val_accuracy: 0.8398
Epoch 12/30
68/68 [=====] - 4s 60ms/step - loss: 0.4866 - accuracy: 0.8357 - val_loss: 0.6282 - val_accuracy: 0.7891
Epoch 13/30
68/68 [=====] - 4s 60ms/step - loss: 0.4712 - accuracy: 0.8454 - val_loss: 0.4462 - val_accuracy: 0.8594
Epoch 14/30
68/68 [=====] - 4s 60ms/step - loss: 0.4519 - accuracy: 0.8468 - val_loss: 0.5364 - val_accuracy: 0.8301
Epoch 15/30
68/68 [=====] - 4s 60ms/step - loss: 0.4071 - accuracy: 0.8672 - val_loss: 0.4958 - val_accuracy: 0.8516
Epoch 16/30
68/68 [=====] - 4s 62ms/step - loss: 0.3997 - accuracy: 0.8663 - val_loss: 0.4683 - val_accuracy: 0.8516
Epoch 17/30
68/68 [=====] - 4s 62ms/step - loss: 0.3722 - accuracy: 0.8863 - val_loss: 0.5340 - val_accuracy: 0.8301
Epoch 18/30
68/68 [=====] - 4s 61ms/step - loss: 0.3492 - accuracy: 0.8909 - val_loss: 0.4158 - val_accuracy: 0.8770
Epoch 19/30
68/68 [=====] - 4s 61ms/step - loss: 0.3417 - accuracy: 0.8923 - val_loss: 0.4544 - val_accuracy: 0.8574
Epoch 20/30
68/68 [=====] - 4s 60ms/step - loss: 0.3642 - accuracy: 0.8895 - val_loss: 0.4595 - val_accuracy: 0.8633
Epoch 21/30
68/68 [=====] - 4s 62ms/step - loss: 0.3205 - accuracy: 0.9048 - val_loss: 0.4325 - val_accuracy: 0.8652
Epoch 22/30
68/68 [=====] - 4s 60ms/step - loss: 0.2965 - accuracy: 0.9109 - val_loss: 0.4514 - val_accuracy: 0.8555
Epoch 23/30
68/68 [=====] - 4s 60ms/step - loss: 0.3000 - accuracy: 0.9123 - val_loss: 0.3728 - val_accuracy: 0.9062
Epoch 24/30
68/68 [=====] - 4s 60ms/step - loss: 0.3118 - accuracy: 0.9076 - val_loss: 0.5239 - val_accuracy: 0.8516
Epoch 25/30
68/68 [=====] - 4s 60ms/step - loss: 0.2723 - accuracy: 0.9266 - val_loss: 0.4666 - val_accuracy: 0.8691
```

```
Epoch 26/30
68/68 [=====] - 4s 60ms/step - loss: 0.2535 - accuracy: 0.9364 - val_loss: 0.3698 - val_accuracy: 0.9023
Epoch 27/30
68/68 [=====] - 4s 60ms/step - loss: 0.2476 - accuracy: 0.9364 - val_loss: 0.3700 - val_accuracy: 0.9023
Epoch 28/30
68/68 [=====] - 4s 60ms/step - loss: 0.2470 - accuracy: 0.9355 - val_loss: 0.2867 - val_accuracy: 0.9258
Epoch 29/30
68/68 [=====] - 4s 60ms/step - loss: 0.2776 - accuracy: 0.9243 - val_loss: 0.3258 - val_accuracy: 0.9082
Epoch 30/30
68/68 [=====] - 4s 60ms/step - loss: 0.2670 - accuracy: 0.9262 - val_loss: 0.5536 - val_accuracy: 0.8379
```

The GridSearchCV prints all the scores as well as the best params which is very convenient:

```
[30]: grid_search_cv.summary()
```

```
-----
-----
Best score: [0.87890625 0.8984375 0.8886719 0.87890625 0.8574219 ] using
params: {'augment': False, 'l2_param': 0.0001}
-----
```

```
Idx: 0 - Score: [0.7597656 0.796875 0.71484375 0.7714844 0.83203125] with
param: {'augment': False, 'l2_param': 0.01}
Idx: 1 - Score: [0.8769531 0.89453125 0.82421875 0.8671875 0.828125 ] with
param: {'augment': False, 'l2_param': 0.001}
Idx: 2 - Score: [0.87890625 0.8984375 0.8886719 0.87890625 0.8574219 ] with
param: {'augment': False, 'l2_param': 0.0001}
-----
```

```
[30]: """
-----\nBest score: [0.87890625 0.8984375
0.8886719 0.87890625 0.8574219 ] using params: {'augment': False, 'l2_param':
0.0001}\n-----\n-----Idx: 0 - Score: [0.7597656
0.796875 0.71484375 0.7714844 0.83203125] with param: {'augment': False,
'l2_param': 0.01}Idx: 1 - Score: [0.8769531 0.89453125 0.82421875 0.8671875
0.828125 ] with param: {'augment': False, 'l2_param': 0.001}Idx: 2 - Score:
[0.87890625 0.8984375 0.8886719 0.87890625 0.8574219 ] with param: {'augment':
False, 'l2_param': 0.0001}-----\n"
```

5.6 Training the Model

We are now training the model directly with data augmentation to make predictions on the test data and to interpret its performance.

```
[39]: model_optimized = create_model_optimized(augment=True, weight_factor=2)
model_optimized.summary()
```

```
Found 2666 files belonging to 5 classes.
Using 2133 files for training.
Found 2666 files belonging to 5 classes.
Using 533 files for validation.
Found 2499 files belonging to 5 classes.
Model: "sequential_34"
```

Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 150, 150, 3)	0
batch_normalization_31 (BatchNormalization)	(None, 150, 150, 3)	12
conv2d_165 (Conv2D)	(None, 150, 150, 64)	1792
max_pooling2d_165 (MaxPooling2D)	(None, 75, 75, 64)	0
dropout_155 (Dropout)	(None, 75, 75, 64)	0
conv2d_166 (Conv2D)	(None, 75, 75, 128)	73856
max_pooling2d_166 (MaxPooling2D)	(None, 37, 37, 128)	0
dropout_156 (Dropout)	(None, 37, 37, 128)	0
conv2d_167 (Conv2D)	(None, 37, 37, 256)	295168
max_pooling2d_167 (MaxPooling2D)	(None, 18, 18, 256)	0
dropout_157 (Dropout)	(None, 18, 18, 256)	0
conv2d_168 (Conv2D)	(None, 18, 18, 512)	1180160
max_pooling2d_168 (MaxPooling2D)	(None, 9, 9, 512)	0

dropout_158 (Dropout)	(None, 9, 9, 512)	0
conv2d_169 (Conv2D)	(None, 9, 9, 1024)	4719616
max_pooling2d_169 (MaxPooling2D)	(None, 4, 4, 1024)	0
dropout_159 (Dropout)	(None, 4, 4, 1024)	0
flatten_33 (Flatten)	(None, 16384)	0
dense_66 (Dense)	(None, 256)	4194560
dense_67 (Dense)	(None, 5)	1285
<hr/>		
Total params:	10,466,449	
Trainable params:	10,466,443	
Non-trainable params:	6	

```
[40]: history = model_optimized.fit(train_ds, validation_data=val_ds, epochs=50)
```

Epoch 1/50

2023-03-17 13:38:01.730401: I
 tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
 Plugin optimizer for device_type GPU is enabled.
 134/134 [=====] - ETA: 0s - loss: 2.3965 - accuracy: 0.3919

2023-03-17 13:38:22.712039: I
 tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
 Plugin optimizer for device_type GPU is enabled.

134/134 [=====] - 23s 152ms/step - loss: 2.3965 -
 accuracy: 0.3919 - val_loss: 1.6807 - val_accuracy: 0.5497

Epoch 2/50

134/134 [=====] - 16s 122ms/step - loss: 1.5764 -
 accuracy: 0.5373 - val_loss: 1.3333 - val_accuracy: 0.5929

Epoch 3/50

134/134 [=====] - 16s 121ms/step - loss: 1.3037 -
 accuracy: 0.5823 - val_loss: 1.1786 - val_accuracy: 0.6248

Epoch 4/50

134/134 [=====] - 16s 123ms/step - loss: 1.1532 -
 accuracy: 0.6160 - val_loss: 1.0931 - val_accuracy: 0.6623

Epoch 5/50

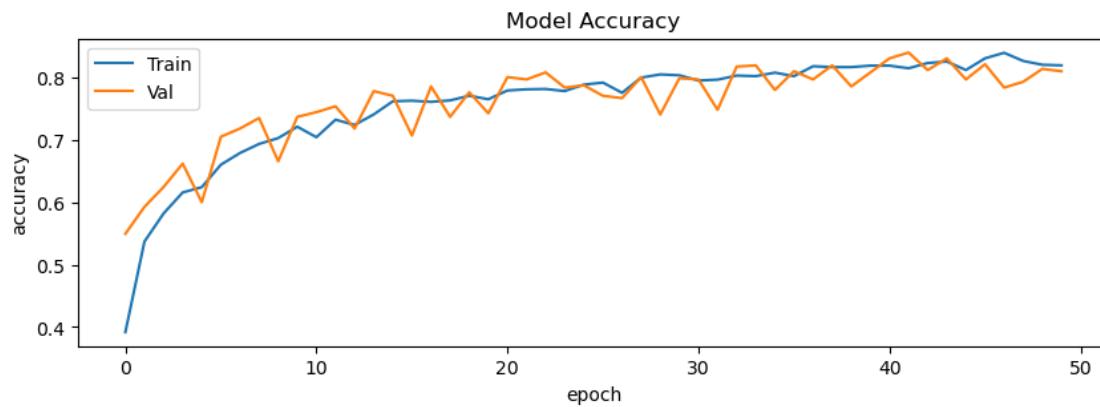
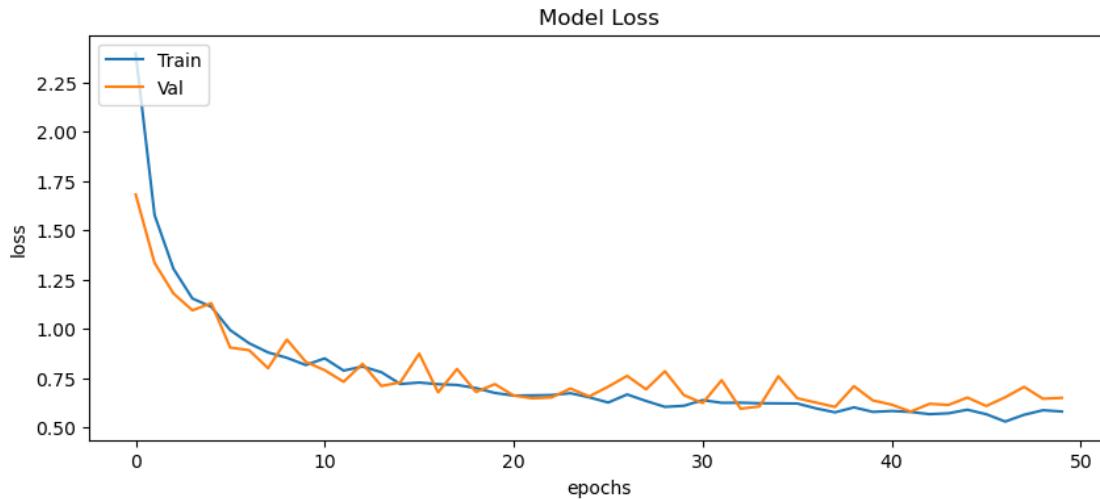
134/134 [=====] - 16s 120ms/step - loss: 1.1109 -
 accuracy: 0.6245 - val_loss: 1.1293 - val_accuracy: 0.6004

Epoch 6/50
134/134 [=====] - 17s 128ms/step - loss: 0.9926 -
accuracy: 0.6606 - val_loss: 0.9041 - val_accuracy: 0.7054
Epoch 7/50
134/134 [=====] - 17s 126ms/step - loss: 0.9260 -
accuracy: 0.6793 - val_loss: 0.8908 - val_accuracy: 0.7186
Epoch 8/50
134/134 [=====] - 16s 120ms/step - loss: 0.8795 -
accuracy: 0.6939 - val_loss: 0.7994 - val_accuracy: 0.7355
Epoch 9/50
134/134 [=====] - 16s 118ms/step - loss: 0.8520 -
accuracy: 0.7032 - val_loss: 0.9447 - val_accuracy: 0.6660
Epoch 10/50
134/134 [=====] - 16s 121ms/step - loss: 0.8156 -
accuracy: 0.7215 - val_loss: 0.8318 - val_accuracy: 0.7373
Epoch 11/50
134/134 [=====] - 16s 120ms/step - loss: 0.8492 -
accuracy: 0.7046 - val_loss: 0.7892 - val_accuracy: 0.7448
Epoch 12/50
134/134 [=====] - 16s 118ms/step - loss: 0.7872 -
accuracy: 0.7328 - val_loss: 0.7305 - val_accuracy: 0.7542
Epoch 13/50
134/134 [=====] - 16s 121ms/step - loss: 0.8078 -
accuracy: 0.7243 - val_loss: 0.8215 - val_accuracy: 0.7186
Epoch 14/50
134/134 [=====] - 16s 120ms/step - loss: 0.7789 -
accuracy: 0.7412 - val_loss: 0.7092 - val_accuracy: 0.7786
Epoch 15/50
134/134 [=====] - 16s 120ms/step - loss: 0.7192 -
accuracy: 0.7623 - val_loss: 0.7277 - val_accuracy: 0.7711
Epoch 16/50
134/134 [=====] - 16s 117ms/step - loss: 0.7265 -
accuracy: 0.7632 - val_loss: 0.8733 - val_accuracy: 0.7073
Epoch 17/50
134/134 [=====] - 16s 119ms/step - loss: 0.7182 -
accuracy: 0.7614 - val_loss: 0.6767 - val_accuracy: 0.7861
Epoch 18/50
134/134 [=====] - 16s 120ms/step - loss: 0.7143 -
accuracy: 0.7637 - val_loss: 0.7955 - val_accuracy: 0.7373
Epoch 19/50
134/134 [=====] - 16s 122ms/step - loss: 0.6985 -
accuracy: 0.7712 - val_loss: 0.6787 - val_accuracy: 0.7767
Epoch 20/50
134/134 [=====] - 16s 121ms/step - loss: 0.6740 -
accuracy: 0.7656 - val_loss: 0.7186 - val_accuracy: 0.7430
Epoch 21/50
134/134 [=====] - 16s 121ms/step - loss: 0.6604 -
accuracy: 0.7797 - val_loss: 0.6605 - val_accuracy: 0.8011

Epoch 22/50
134/134 [=====] - 16s 119ms/step - loss: 0.6614 -
accuracy: 0.7815 - val_loss: 0.6464 - val_accuracy: 0.7974
Epoch 23/50
134/134 [=====] - 16s 119ms/step - loss: 0.6629 -
accuracy: 0.7820 - val_loss: 0.6510 - val_accuracy: 0.8086
Epoch 24/50
134/134 [=====] - 16s 119ms/step - loss: 0.6726 -
accuracy: 0.7787 - val_loss: 0.6963 - val_accuracy: 0.7842
Epoch 25/50
134/134 [=====] - 16s 119ms/step - loss: 0.6514 -
accuracy: 0.7890 - val_loss: 0.6562 - val_accuracy: 0.7880
Epoch 26/50
134/134 [=====] - 16s 119ms/step - loss: 0.6252 -
accuracy: 0.7923 - val_loss: 0.7060 - val_accuracy: 0.7711
Epoch 27/50
134/134 [=====] - 16s 121ms/step - loss: 0.6662 -
accuracy: 0.7759 - val_loss: 0.7603 - val_accuracy: 0.7674
Epoch 28/50
134/134 [=====] - 16s 118ms/step - loss: 0.6331 -
accuracy: 0.8003 - val_loss: 0.6923 - val_accuracy: 0.8011
Epoch 29/50
134/134 [=====] - 16s 119ms/step - loss: 0.6032 -
accuracy: 0.8054 - val_loss: 0.7842 - val_accuracy: 0.7411
Epoch 30/50
134/134 [=====] - 16s 119ms/step - loss: 0.6086 -
accuracy: 0.8040 - val_loss: 0.6629 - val_accuracy: 0.7992
Epoch 31/50
134/134 [=====] - 16s 119ms/step - loss: 0.6375 -
accuracy: 0.7956 - val_loss: 0.6225 - val_accuracy: 0.7974
Epoch 32/50
134/134 [=====] - 16s 118ms/step - loss: 0.6237 -
accuracy: 0.7970 - val_loss: 0.7387 - val_accuracy: 0.7486
Epoch 33/50
134/134 [=====] - 16s 118ms/step - loss: 0.6242 -
accuracy: 0.8036 - val_loss: 0.5937 - val_accuracy: 0.8180
Epoch 34/50
134/134 [=====] - 16s 119ms/step - loss: 0.6217 -
accuracy: 0.8026 - val_loss: 0.6055 - val_accuracy: 0.8199
Epoch 35/50
134/134 [=====] - 16s 120ms/step - loss: 0.6210 -
accuracy: 0.8083 - val_loss: 0.7583 - val_accuracy: 0.7805
Epoch 36/50
134/134 [=====] - 16s 119ms/step - loss: 0.6202 -
accuracy: 0.8026 - val_loss: 0.6469 - val_accuracy: 0.8105
Epoch 37/50
134/134 [=====] - 16s 119ms/step - loss: 0.5954 -
accuracy: 0.8186 - val_loss: 0.6248 - val_accuracy: 0.7974

```
Epoch 38/50
134/134 [=====] - 16s 118ms/step - loss: 0.5753 -
accuracy: 0.8172 - val_loss: 0.6029 - val_accuracy: 0.8199
Epoch 39/50
134/134 [=====] - 16s 119ms/step - loss: 0.6007 -
accuracy: 0.8172 - val_loss: 0.7081 - val_accuracy: 0.7861
Epoch 40/50
134/134 [=====] - 16s 118ms/step - loss: 0.5777 -
accuracy: 0.8195 - val_loss: 0.6357 - val_accuracy: 0.8086
Epoch 41/50
134/134 [=====] - 16s 119ms/step - loss: 0.5822 -
accuracy: 0.8195 - val_loss: 0.6144 - val_accuracy: 0.8311
Epoch 42/50
134/134 [=====] - 16s 118ms/step - loss: 0.5776 -
accuracy: 0.8153 - val_loss: 0.5792 - val_accuracy: 0.8405
Epoch 43/50
134/134 [=====] - 16s 117ms/step - loss: 0.5661 -
accuracy: 0.8237 - val_loss: 0.6186 - val_accuracy: 0.8124
Epoch 44/50
134/134 [=====] - 16s 119ms/step - loss: 0.5706 -
accuracy: 0.8261 - val_loss: 0.6132 - val_accuracy: 0.8311
Epoch 45/50
134/134 [=====] - 16s 119ms/step - loss: 0.5887 -
accuracy: 0.8125 - val_loss: 0.6502 - val_accuracy: 0.7974
Epoch 46/50
134/134 [=====] - 16s 118ms/step - loss: 0.5658 -
accuracy: 0.8312 - val_loss: 0.6077 - val_accuracy: 0.8218
Epoch 47/50
134/134 [=====] - 16s 119ms/step - loss: 0.5286 -
accuracy: 0.8401 - val_loss: 0.6519 - val_accuracy: 0.7842
Epoch 48/50
134/134 [=====] - 16s 118ms/step - loss: 0.5631 -
accuracy: 0.8270 - val_loss: 0.7047 - val_accuracy: 0.7936
Epoch 49/50
134/134 [=====] - 16s 118ms/step - loss: 0.5861 -
accuracy: 0.8209 - val_loss: 0.6443 - val_accuracy: 0.8143
Epoch 50/50
134/134 [=====] - 16s 117ms/step - loss: 0.5796 -
accuracy: 0.8200 - val_loss: 0.6489 - val_accuracy: 0.8105
```

```
[41]: visualize_results(history)
```



The diagram may look very volatile, but keep in mind the y-axis is scaled.

```
[46]: model_optimized.evaluate(test_ds)
```

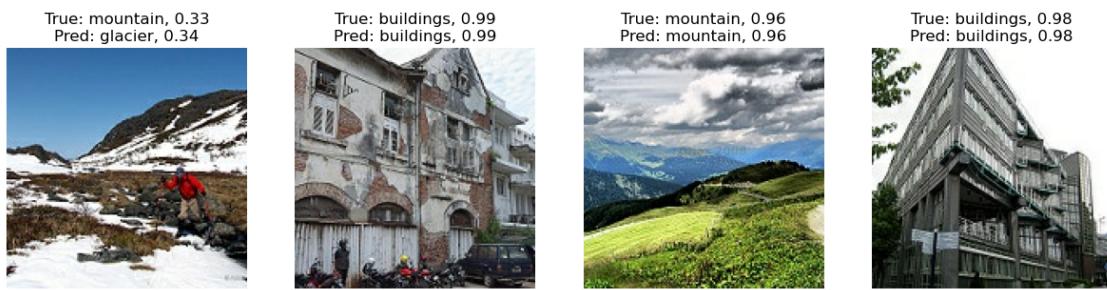
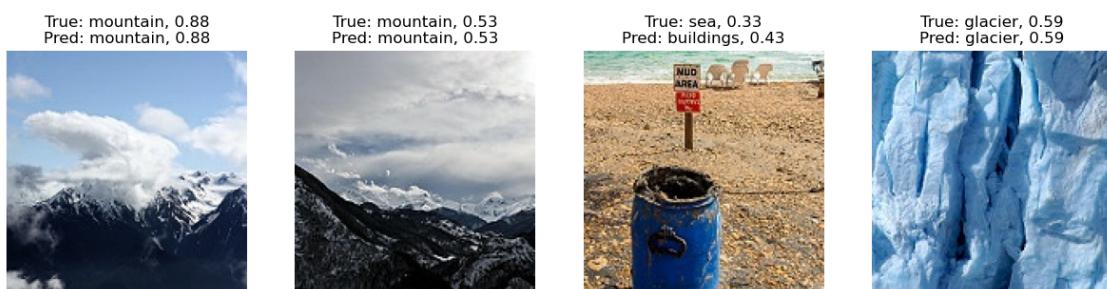
```
157/157 [=====] - 4s 22ms/step - loss: 0.6600 -
accuracy: 0.8055
```

```
[46]: [0.6600058674812317, 0.8055222034454346]
```

```
[47]: show_images_with_predictions(model_optimized)
```

```
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 69ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
```

```
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
```



5.7 Confusion Matrix

```
[44]: true_labels, pred_labels = show_confusion_matrix(model_optimized)
```

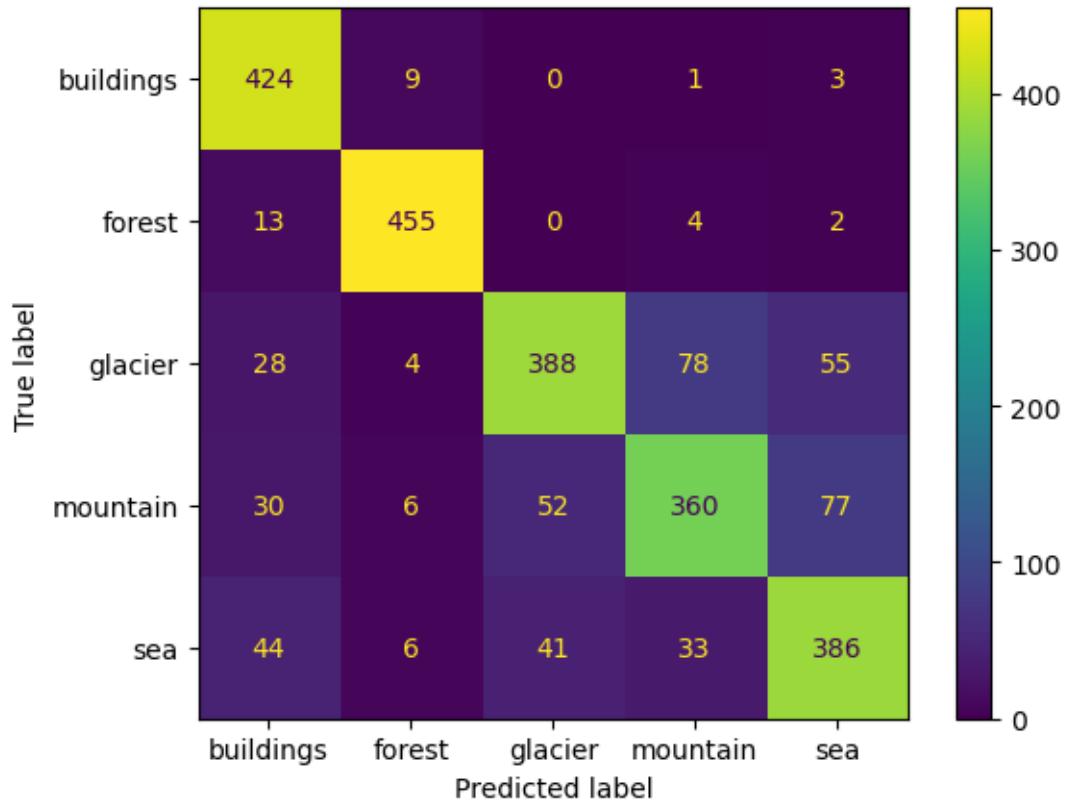
```
1/1 [=====] - 0s 278ms/step
1/1 [=====] - 0s 62ms/step
1/1 [=====] - 0s 12ms/step
```



```

1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 11ms/step
1/1 [=====] - 0s 10ms/step
1/1 [=====] - 0s 217ms/step

```



As from the above Confusion Matrix, the predictions for the class forest and buildings are very accurate while predictions for glacier and mountain get confused the most often. A possible reason for this is that some images containing snow and are labeled as mountain while others are being labeled as glacier. In one example, the predictions for mountain and glacier are really close together where an ensemble method of multiple models with other weaknesses could help to improve the predictions:

True: mountain, 0.33
Pred: glacier, 0.34



5.8 Classification Report

```
[45]: print('Classification Report:\n', metrics.  
        classification_report(y_true=true_labels, y_pred=pred_labels))
```

	precision	recall	f1-score	support
0	0.79	0.97	0.87	437
1	0.95	0.96	0.95	474
2	0.81	0.70	0.75	553
3	0.76	0.69	0.72	525
4	0.74	0.76	0.75	510
accuracy			0.81	2499
macro avg	0.81	0.81	0.81	2499
weighted avg	0.81	0.81	0.80	2499

6 Conclusion

This project was exciting as it was challenging. It made us really try out numerous ways on how to tweak both performance in the sense of speed and of model accuracy. It is interesting how a small

change can have a huge impact on the model's learning behavior. We've also extensively dealt with the over- and underfitting effects which consolidated our understanding on how and when to use regularization and generalization optimizations.

In summary, creating a model with tensorflow is very easy and convenient but creating a performant model can be an overwhelming and time consuming process that can be frustrating at some point. But seeing the results of a successful improvement, on the other hand, is very satisfactory. The biggest difficulty in our opinion is to not loose the overview of the data handling.

7 Outlook

As a wrap-up we'd like to give a short outlook on what could be further steps to improve the model or to take the model into production.

7.1 Export and Reuse

Tensorflow allows to export the weights of a trained model. In this way a new model can be initialized using the saved weights and can be used directly to make predictions on unseen data. This would be the method we'd use if we'd take this model to some productive use.

7.2 Deployment of the model

Furthermore, it would be possible to implement the model on a cloud-based platform or server for productive purposes. This would require making an API that can handle requests and return predictions, while also constructing the essential infrastructure required to handle the computational power.

7.3 Further Optimizations

The parameter optimizations we've made with Cross Validation are valid but would need some improvements for productive models. For example it would be necessary to optimize with multiple parameter changes e.g. simultaneously train for the optimal parameter of batch size, l2 parameter, dropout values, etc. We've also noticed that the results are not reproducible which implies that we would need to train more folds than 5 to obtain more stable results.

Even though it wasn't a competition to create the most accurate model, we tried to find optimal parameters to increase the accuracy and minimize the loss. We were stuck at 80% to 85% accuracy and we've learned that the last 15% were the most difficult to achieve. In the beginning it was way easier to get from 60% to 70% by tuning a single parameter or adding a new layer to the model. But for the last 15% there are so many factors that can be improved in order to obtain an "optimal" model.

To overcome this difficulty of creating the "optimal" model (which we know isn't the goal), we could train multiple models with subsets of the dataset and then apply some ensemble algorithm to find the correct predictions. Furthermore, it is crucial that we don't have wrong labeled data in the dataset since it would confuse the model a lot so manual data preprocessing would be an important step for productive use-cases.

7.4 General Outlook

We are curious to see how machine learning will change the way everyone works and what impact it might have on our lives. In this example, how classifying images may help computervision thrive. We are delighted to be able to take an AI class and therefore learn more about AI as well as create and train our own models during such exciting times.